

CprE 381: Computer Organization and Assembly Level Programming

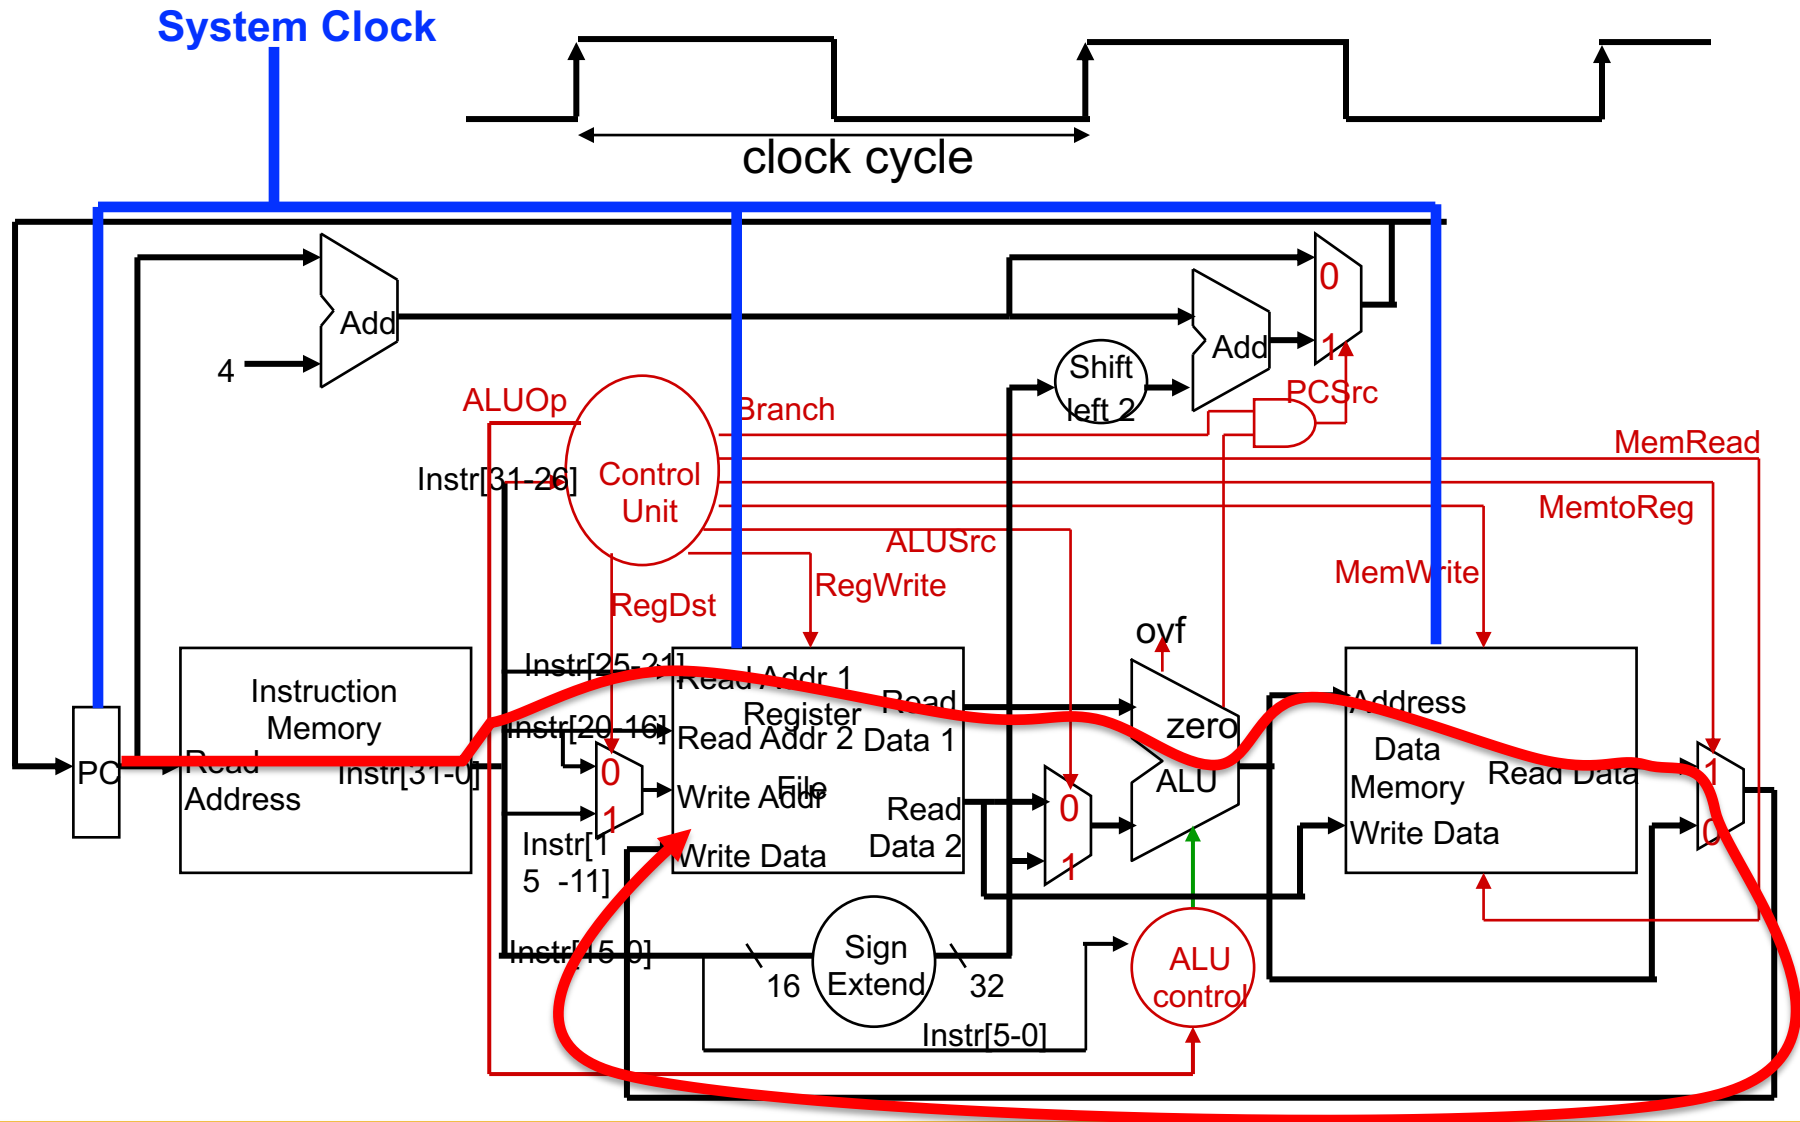
Control Hazards

Henry Duwe
Electrical and Computer Engineering
Iowa State University

Administrative

- Term Proj 2 due this week in lab
 - For keeps
- HW7
 - Due Mar 27 11:59pm
 - Last HW for exam 2 coverage
 - Question 3 best exam prep if done well
- Exam 2
 - When: April 1
 - Where: **TBD**
 - What: MIPS Arithmetic through pipelining (including hazards)

Review: Worst Case Timing (Load Instruction)




Review: Execution Time

- Drawing on the previous equation:

$$\textit{Execution Time} = \# \textit{ Instructions} \times \frac{\textit{Cycles}}{\textit{Instruction}} \times \frac{\textit{Seconds}}{\textit{Cycle}}$$


- To improve performance (i.e., reduce execution time)
 - Increase clock rate (decrease clock cycle time) OR
 - Decrease CPI OR
 - Reduce the number of instructions
- Designers balance cycle time against the number of cycles required
 - Improving one factor may make the other one worse...

Review: Multicycle Processor



Review: A Simple MIPS Pipeline

- All control signals can be determined during Decode
 - And held in the **state registers** between pipeline stages



Review: Oh, the *Hazards* I've Seen

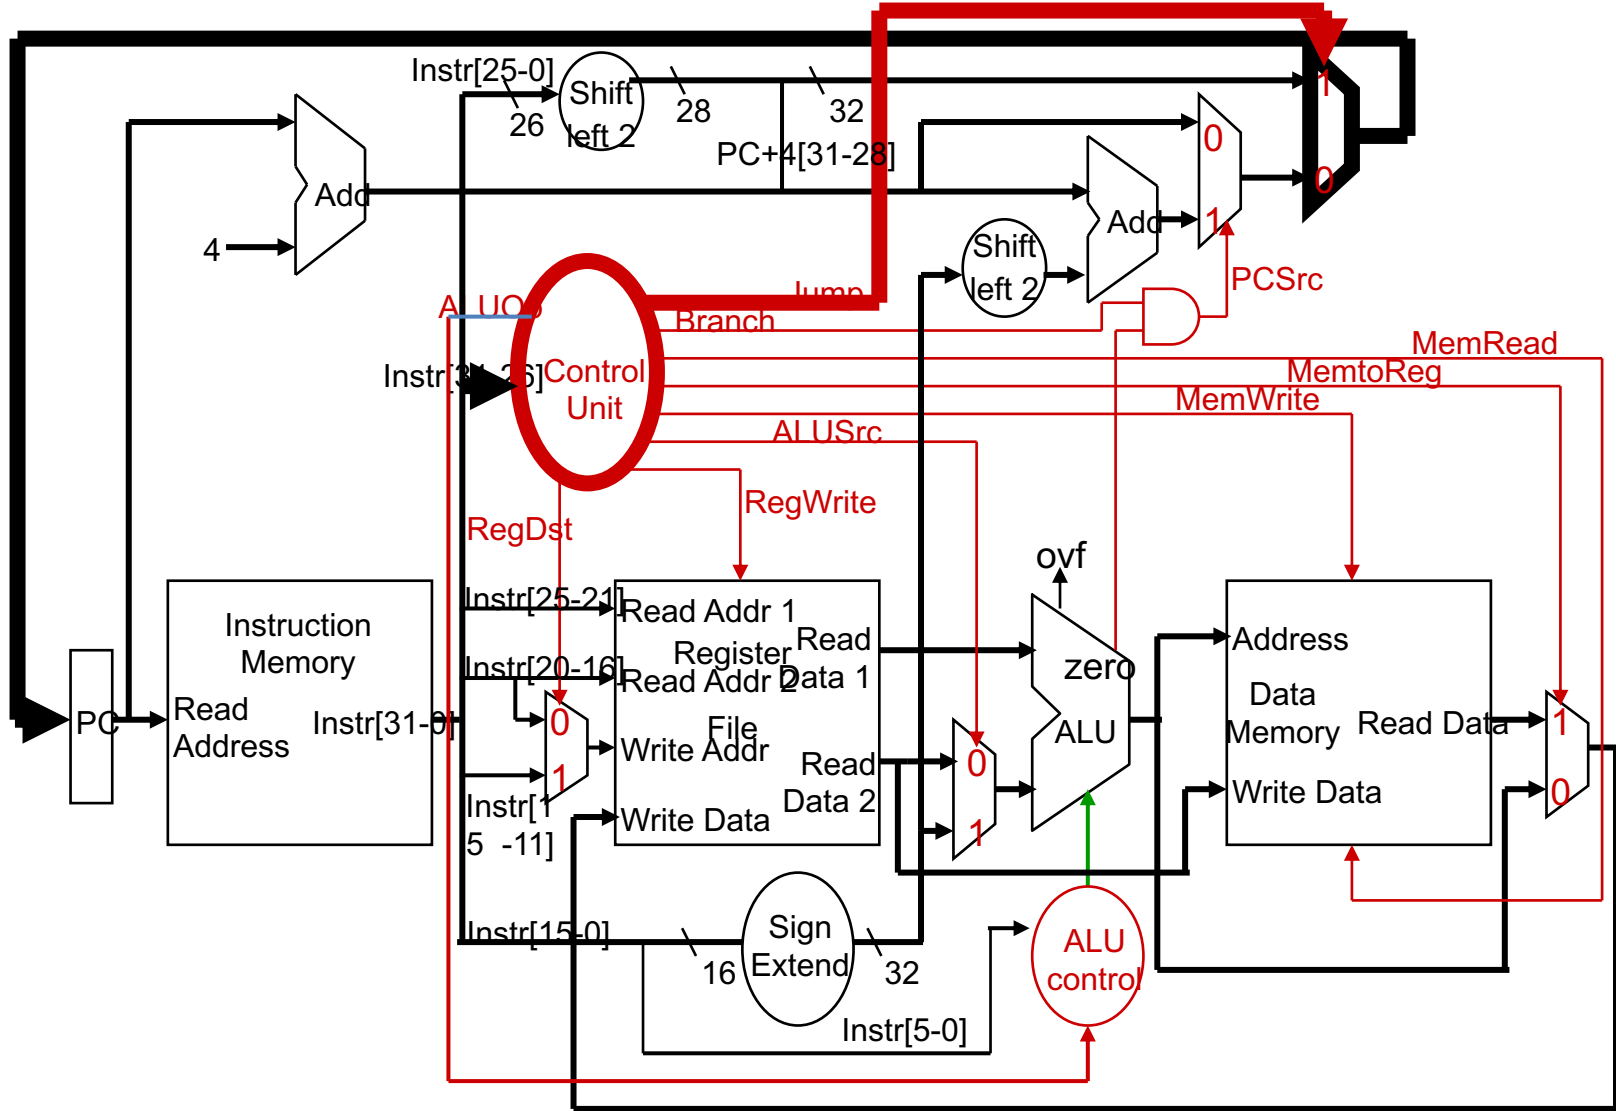
- Pipeline Hazards
 - **Structural hazards**: attempt to use the same resource by two different instructions at the same time
 - **Data hazards**: attempt to use data before it is ready
 - An instruction's source operand(s) are produced by a prior instruction still in the pipeline
 - **Control hazards**: attempt to make a decision about program control flow before the condition has been evaluated and the new PC target address calculated
 - Branch and jump instructions, exceptions
- Can always resolve hazards by **waiting**
 - Pipeline control must **detect** the hazard
 - And take action to **resolve** hazards



Control Hazards

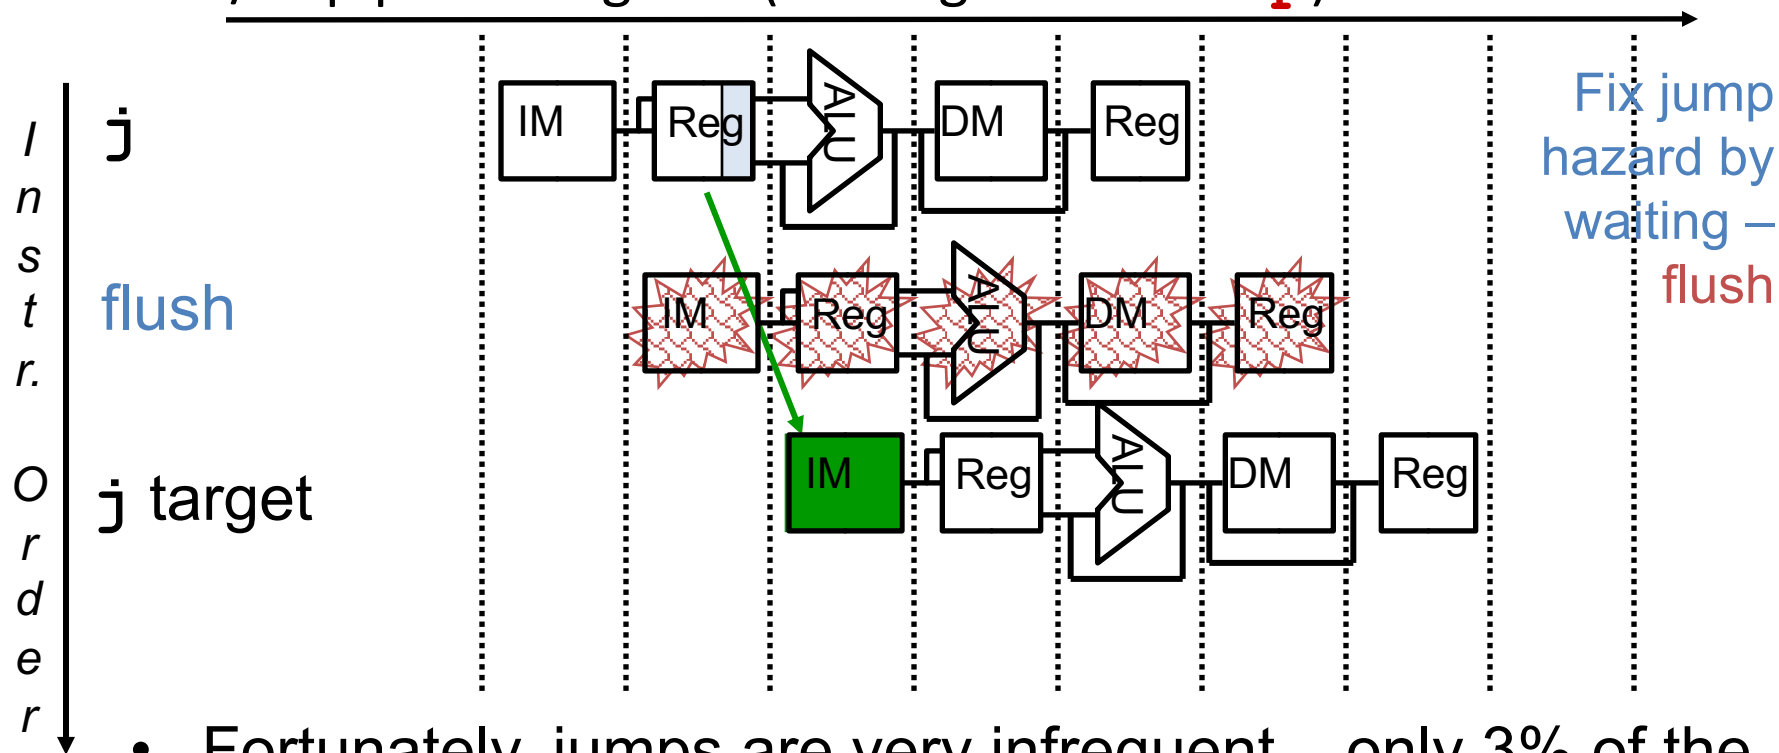
- When the flow of instruction addresses is not sequential (i.e., $PC = PC + 4$)
 - Conditional branches (`beq`, `bne`)
 - Unconditional branches (`j`, `jal`, `jr`)
 - Exceptions
- Possible “solutions”
 - Stall (impacts performance)
 - Move branch decision point as early in the pipeline as possible, thereby reducing the number of stall cycles
 - Delay decision (requires compiler support)
 - Predict and hope for the best!
- Control hazards occur less frequently than data hazards, but there is *nothing* as effective against control hazards as forwarding is for data hazards

Remember: Adding the Jump Instruction



Jumps Incur One Stall

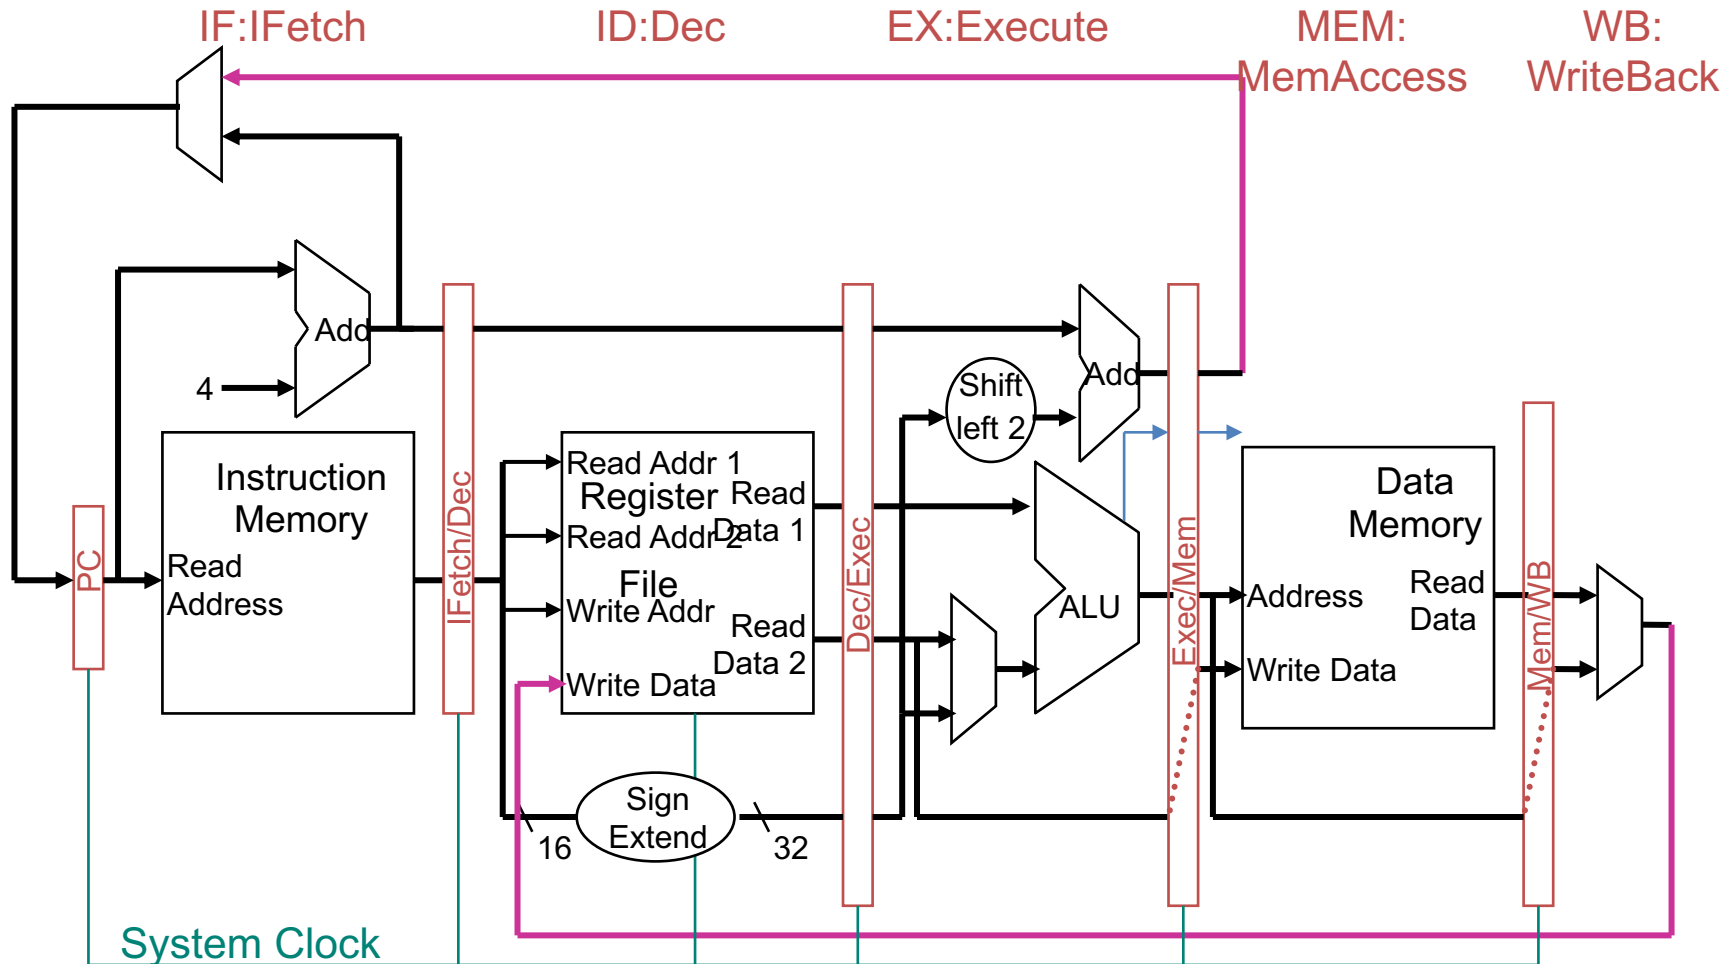
- Jumps not decoded until ID, so one **flush** is needed
 - To flush, set `IF.Flush` to zero the instruction field of the IF/ID pipeline register (turning it into a **nop**)



- Fortunately, jumps are very infrequent – only 3% of the SPECint instruction mix

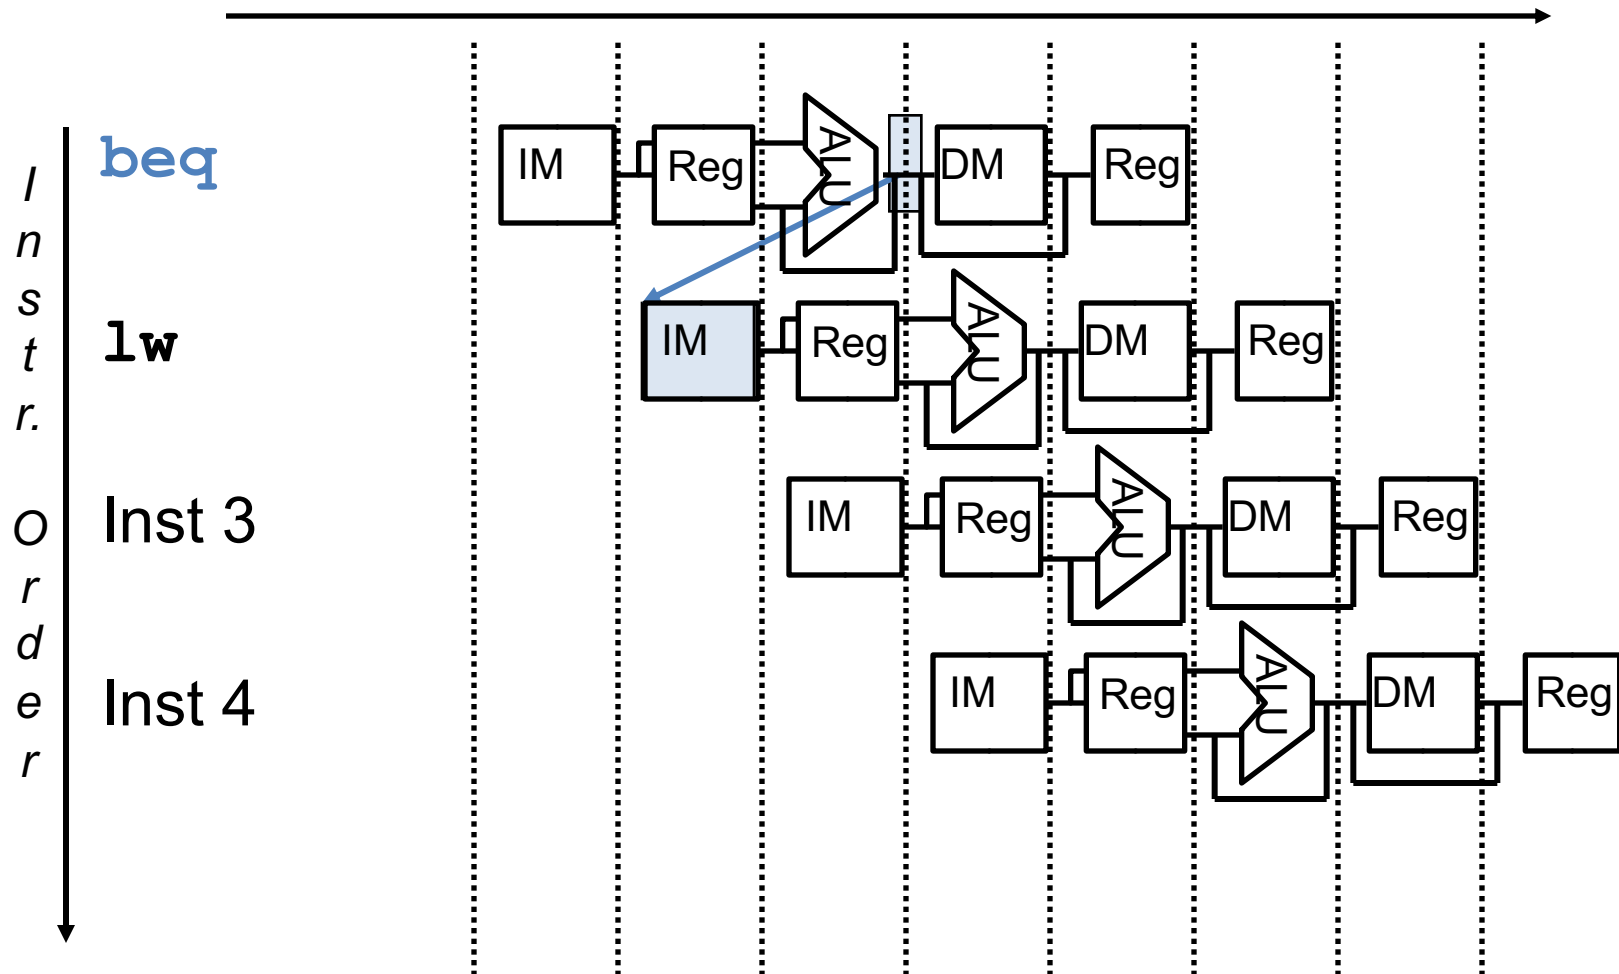
Remember: MIPS Pipeline Datapath Modifications

- What do we need to add/modify in our MIPS datapath?
 - State registers between each pipeline stage to **isolate** them

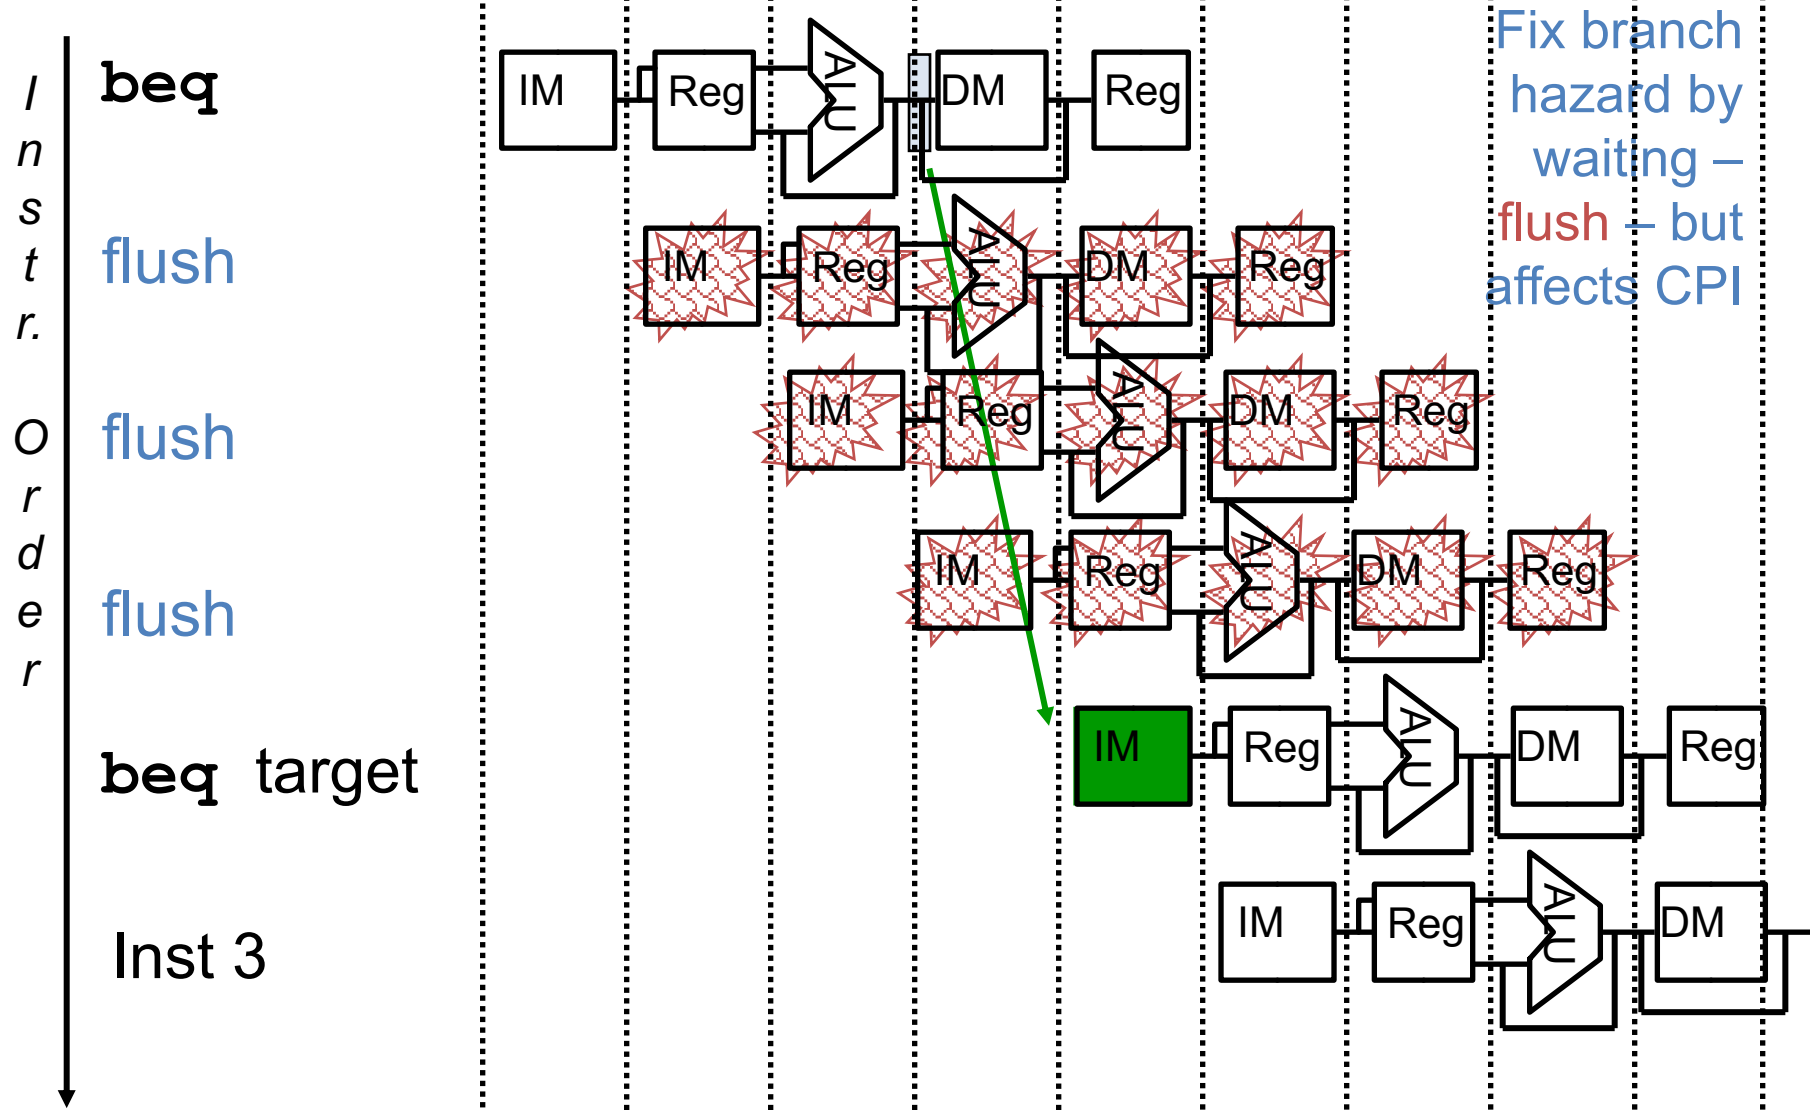


Branches Cause Control Hazards

- Dependencies backward in time cause **hazards**

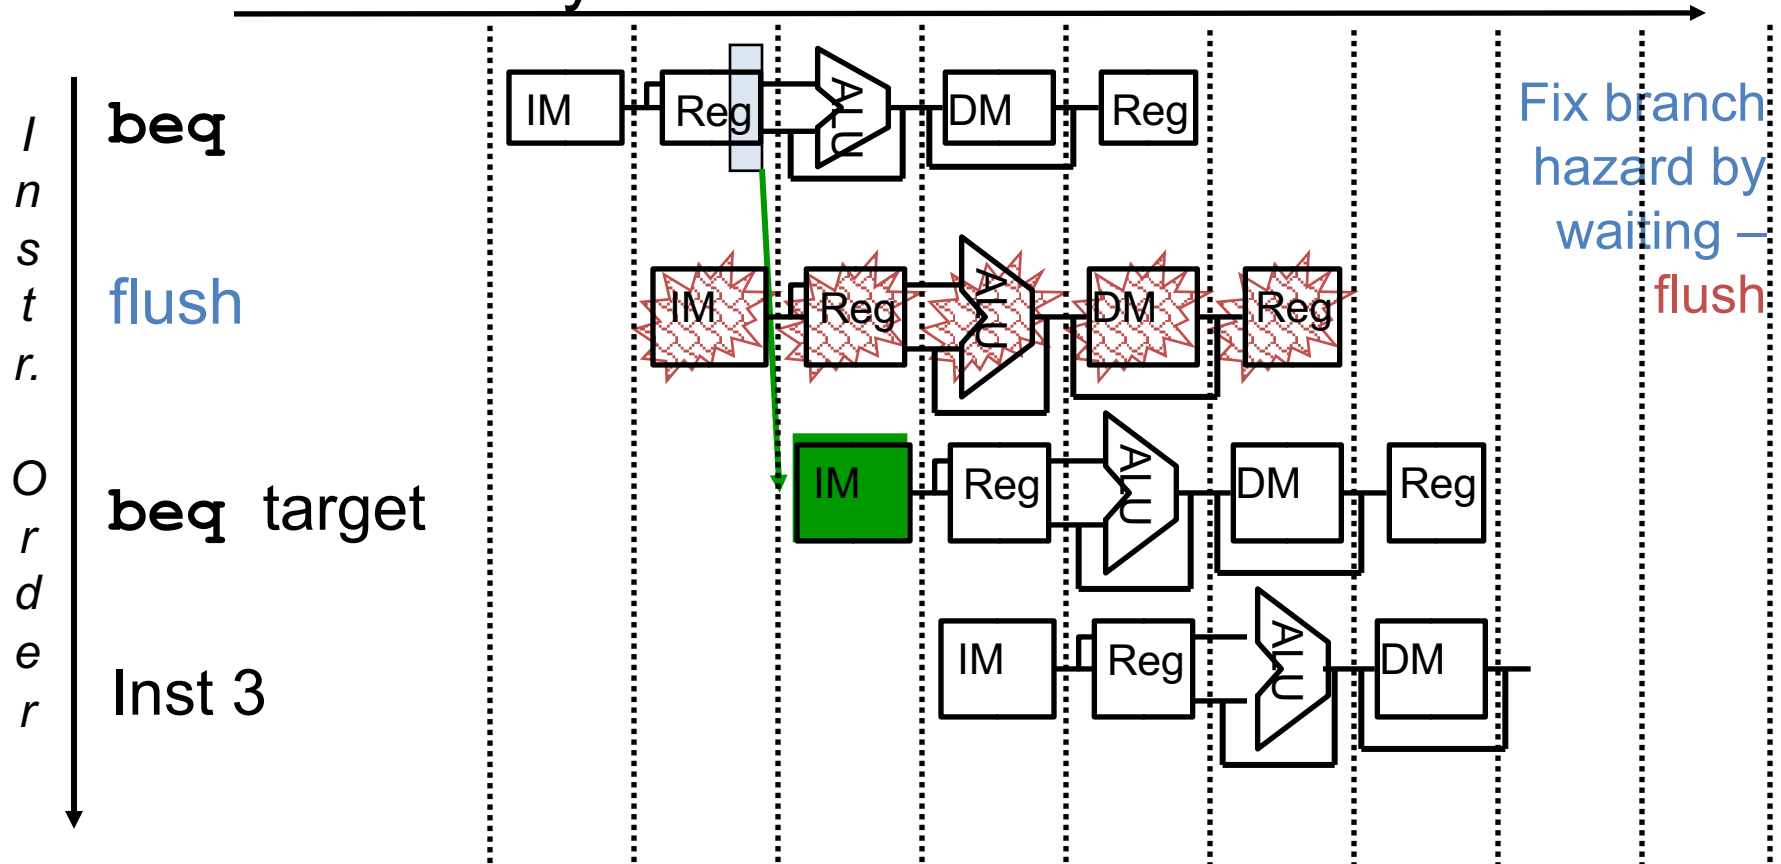


One Way to “Fix” a Branch Control Hazard

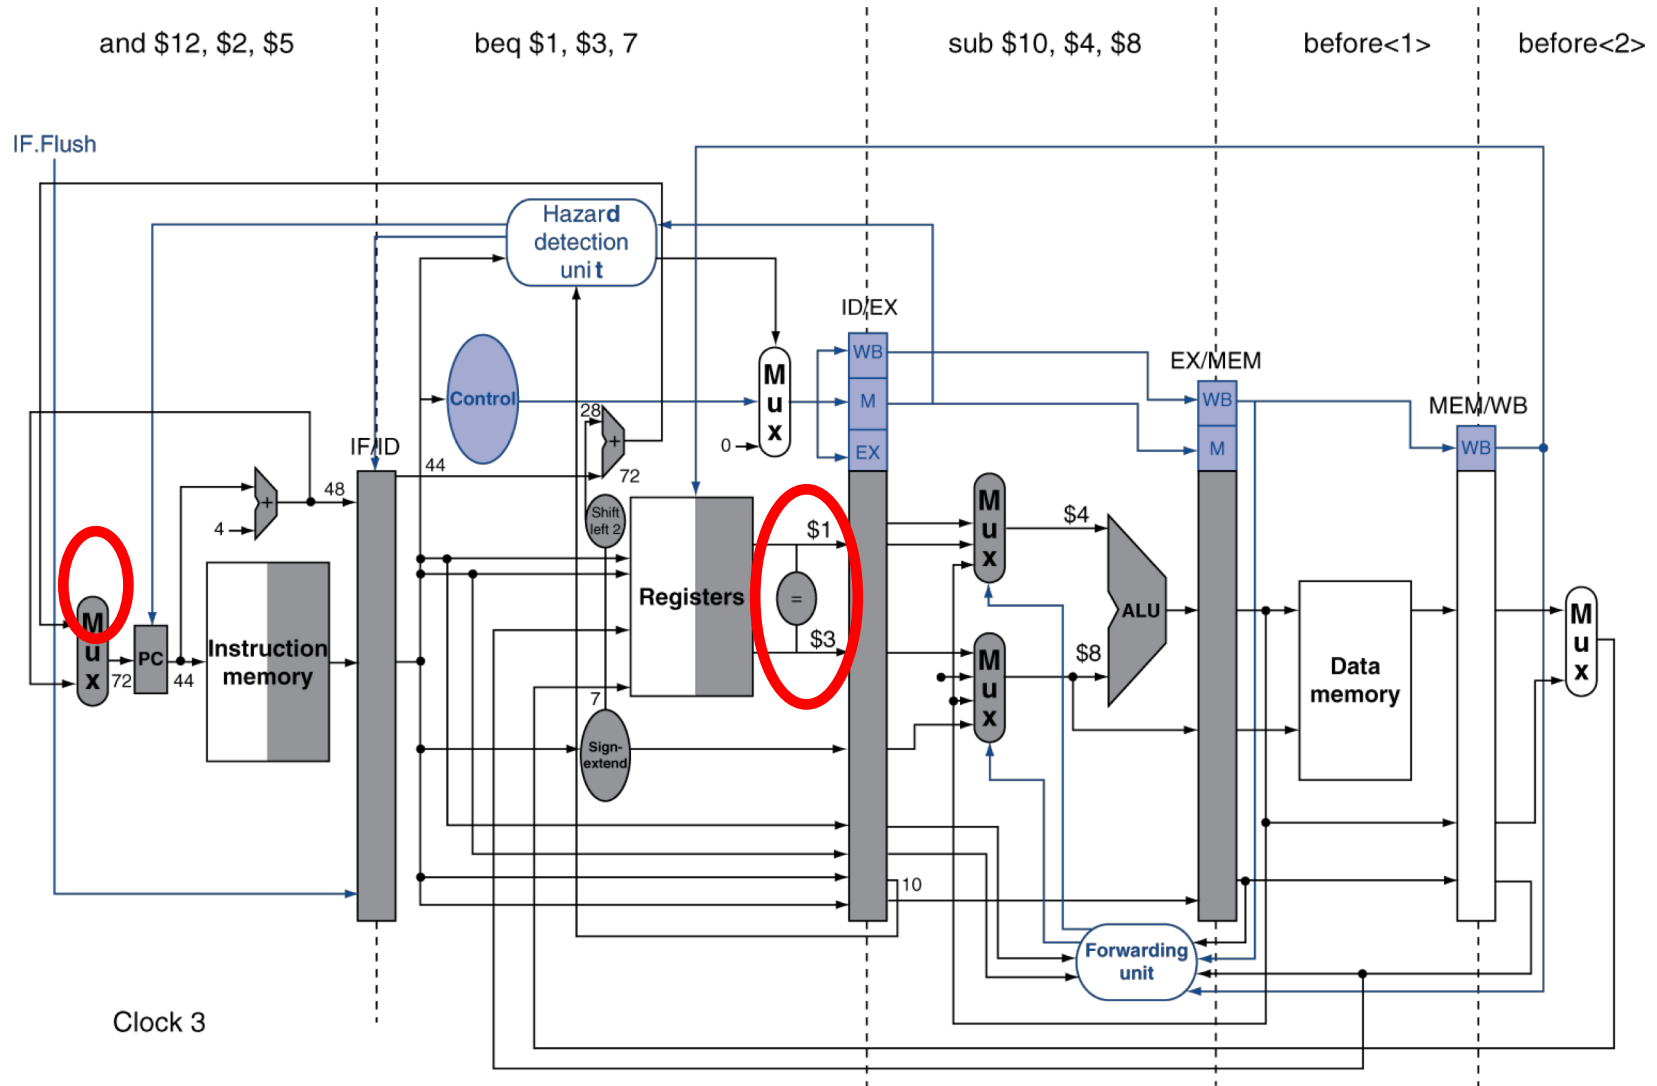


Another Method

- Move branch decision hardware back to as **early** in the pipeline as possible – i.e., during the decode cycle



Branch Comparison Hardware



Branch Stalls and Performance

- Need to stall for one cycle on every branch!
 - ID control point
- Consider the following case
 - The ideal CPI of the machine is 1 (applies to all instr not branches)
 - The branch causes a stall
- What is the new effective CPI after moving the branch decision if 15% of the instructions are branches?

Branch Stalls and Performance

- Need to stall for one cycle on every branch!
 - ID control point

In-class Assessment!

Access Code: QED

Note: sharing access code to those outside of classroom or using access while outside of classroom is considered cheating

- What is the new effective CPI after moving the branch decision if 15% of the instructions are branches?

Branch Stalls and Performance

- Need to stall for one cycle on every branch!
 - ID control point
- Consider the following case
 - The ideal CPI of the machine is 1 (applies to all instr not branches)
 - The branch causes a stall
- What is the new effective CPI after moving the branch decision if 15% of the instructions are branches?
- The new effective CPI is $1 + 1 \times 0.15 = 1.15$
- The old effective CPI was $1 + 3 \times 0.15 = 1.45$

Yet Another Solution: Delayed Branches

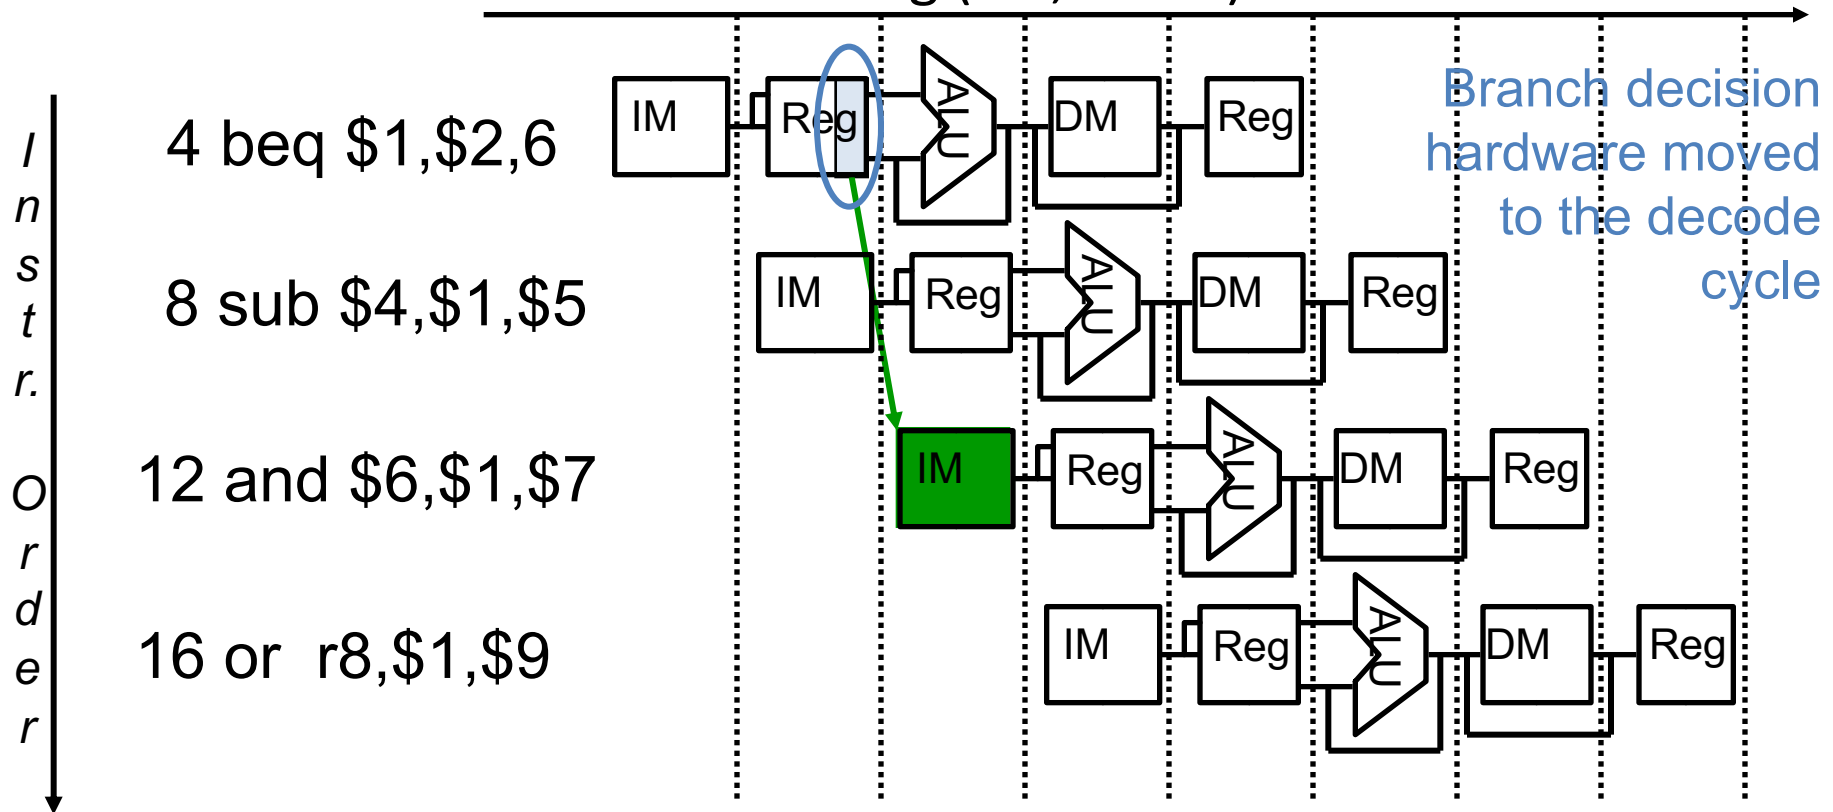
- Solution used in early MIPS machines
 - Had a branch delay of one
 - Branch does not take effect until the cycle after its execution
- Example:

```
beq r1, r2, L      # Branch instruction
sub r4, r1, r3     # instr ALWAYS is executed
and r6, r2, r7     # instr executes if not taken
```

- Worked well initially, but now is a pain
 - Compiler can fill one slot 50% of the time
 - Machines have many branch delay slots
 - Issue more than one instruction per cycle

Still More Fixes

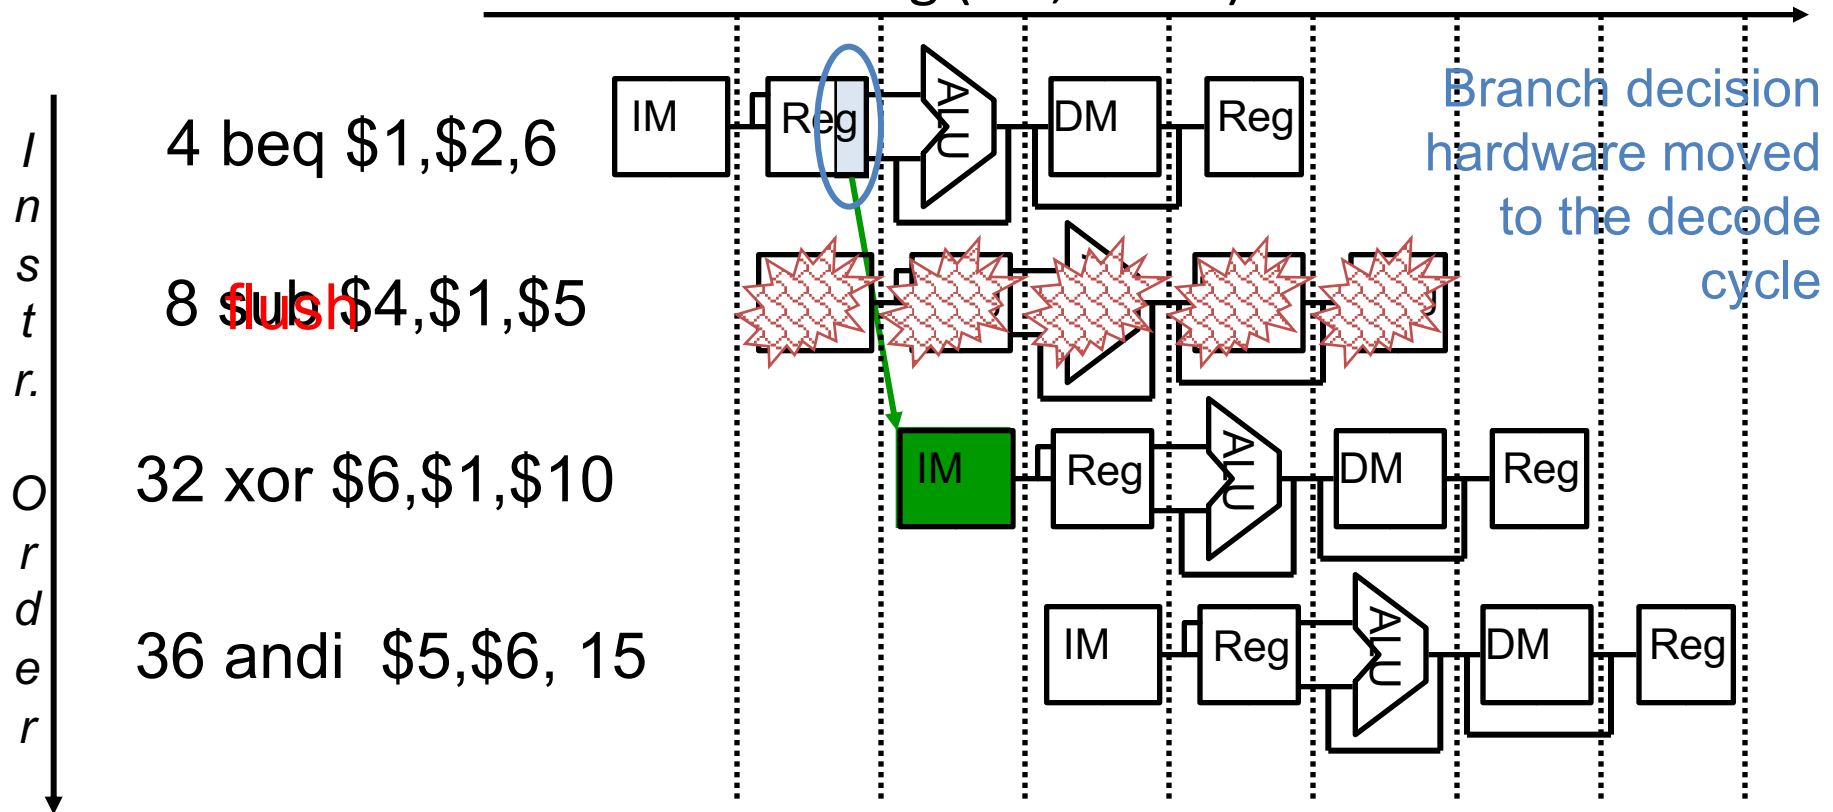
- Predict branches are always not taken – and take corrective action when wrong (i.e., taken)



- To flush, set **IF.Flush** to zero the instruction field of the IF/ID pipeline register (turning it into a **nop**)

Still More Fixes

- Predict branches are always not taken – and take corrective action when wrong (i.e., taken)




- To flush, set **IF.Flush** to zero the instruction field of the IF/ID pipeline register (turning it into a **nop**)

Dynamic Branch Prediction

- In deeper and superscalar pipelines, branch penalty is more significant
- Use dynamic prediction
 - Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
 - To execute a branch
 - Check table, expect the same outcome
 - Start fetching from fall-through or target
 - If wrong, flush pipeline and flip prediction

1-Bit Predictor: Shortcoming

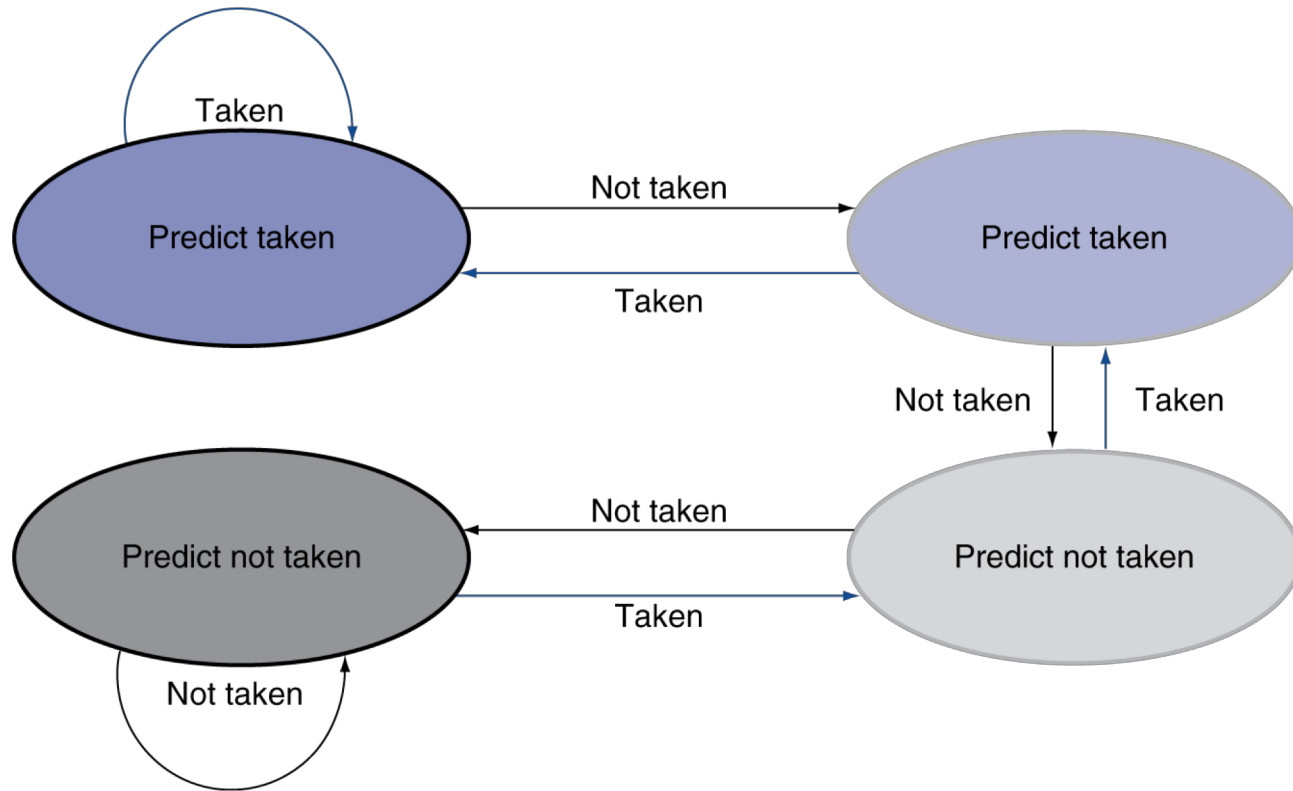
- Inner loop branches mispredicted twice!



- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around

2-Bit Predictor

- Only change prediction on two successive mispredictions



Preview: Calculating the Branch Target

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch

Acknowledgments

- These slides contain material developed and copyright by:
 - Joe Zambreno (Iowa State)
 - David Patterson (UC Berkeley)
 - Mary Jane Irwin (Penn State)
 - Christos Kozyrakis (Stanford)
 - Onur Mutlu (Carnegie Mellon)
 - Krste Asanović (UC Berkeley)