

Name Sean Gordon

ISU ID # 495 762 295

Lab Section (circle one—your exam will be given a 0 if you do this incorrectly):

A (W 10-Noon), B (R 10-Noon), C (F 10-Noon), D (W 4-6)

CprE 381
Computer Organization and Assembly Level
Programming

Exam #1
2/18/2019 9:00-9:50AM

Directions: There are 5 questions in this exam (on 4 pages). Each question is worth the points indicated. You should roughly spend 1 minute for every two points—plan accordingly. If a problem appears to be hard, move on and come back. Please read the questions carefully. Show your work, including any assumptions you need to use to solve the problems.

Calculators should NOT be used.

Problem	Score
1	<u>18</u> / 25 points
2	<u>5</u> / 25 points
3	<u>30</u> / 30 points
4	<u>0</u> / 5 points
5	<u>15</u> / 15 points
Total	<u>68</u> / 100 points

+18

1. MIPS Assembly and SIMD (25 points). Consider the two following MIPS code segments.

```
char *a;                                     // assume a in register $a0
...
for (int i=0; i<N; i++) {                   // assume N in register $a1
    sum += a[i];
}
```

Implementation 1 (no SIMD instructions):

```
1  add $t0, $zero, $zero
2  add $t2, $zero, $zero
3  j    cond
4  loop:
5  addu $t1, $a0, $t0
6  lb   $t1, 0($t1)
7  addu $t2, $t2, $t1
8  addiu $t0, $t0, 1
9  cond:
10 slt  $t1, $t0, $a1
11 bne  $t1, $zero, loop
12 exit:
13
```

Implementation 2 (SIMD):

```
add $t0, $zero, $zero
add $t2, $zero, $zero
j    cond
loop:
addu $t1, $a0, $t0
lw   $t1, 0($t1)
addu $t1, $t1, $t1
addu $t2, $t2, $t1
addu $t0, $t0, 4
cond:
slt  $t1, $t0, $a1
bne  $t1, $zero, loop
exit:
```

(a) (10 points) What is the minimum size of each binary in bytes? Show your work.

Each instruction is 32 bits long / 4 bytes long.

Implementation 1 : 9 lines

Implementation 2 : 10 lines

$$9 \cdot 4 = 36 \text{ bytes}$$

$$10 \cdot 4 = 40 \text{ bytes}$$

(b) (10 points) How many instructions get dynamically executed by each implementation in terms of N? You do not need to simplify, but SHOW YOUR WORK.

$$\begin{aligned} &3 + \\ &2 + \text{Initial condition} \\ &6 \cdot N \\ &= \end{aligned}$$

$$5 + 6 \cdot N$$

$$\begin{aligned} &3 + \\ &2 + \text{Initial condition} \\ &7 \cdot N \\ &= \end{aligned}$$

$$5 + 7 \cdot N$$

(c) (5 points) Describe one programming optimization you could perform to implementation 1 to save an additional instruction per iteration.

~~You could combine implementation 2 lines 7 + 8~~

-5

how?

2. ⁺⁵ Pseudo-Instructions (25 points). Consider the case of the MIPS `ulw` instruction that ^{loads} stores a word to an unaligned memory address (e.g., a word stored to the address `0x10010003`). `ulw` is not supported in hardware, rather it is a pseudo-instruction.

- (a) (15 points) Support `ulw` in software. First, finish labelling the memory diagram provided and shade in the bytes you are intending to load (you may assume either endianness). Second, using only the basic arithmetic, logical, and shift instructions, along with `lw`, provide a sequence of instructions that corresponds to the following `ulw` instruction instance.

`ulw $a0, 3($t0)`

same operands as `lw` expects
you may assume `$t0`'s value is
`0x10010000`

Word Address	Byte 0	Byte 1	Byte 2	Byte 3
...				
				05
0x10010000				
	04	05	06	
...				

~~`srl $t1, $a0, 12` # Get most significant nibble~~

~~`lw $t1, 0($t0)` # load word 03~~

~~`sll $t1, $a0, 4` # Get 3 lower nibbles~~

~~`lw $t1, 1($t0)` # load word 04-06~~

ok... I see the way you labelled

still unaligned addresses
values never used

-10

- (b) (10 points) Why does the MIPS ISA not implement a `ulw` instruction? Provide a technical reason. [Hint: think about how you'd have to modify your Lab 4 datapath to support this.]

~~Loading on an unaligned memory address would not be particularly useful, as the MIPS infrastructure is not equipped to handle something that is not aligned, and as such the stored data would be difficult to use~~

-10

↑ `lb, lh` already load such data

3. **MIPS Formats (30 points)**. MIPS has three base instruction formats (see the R-, I-, and J-formats on the reference sheet attached to your exam). However, domain-specific experts with tunnel-vision always want more. Consider the four newly proposed formats below.

Format C (Compact):

Opcode		rs		rt		rd	
15	9	8	6	5	3	2	0

Imm1		rt		Imm2		Opcode	
31	21	20	16	15	6	5	0

Format F (Partial):			
Opcode		rs	Imm
31	26	25 21	20

Format I (Extended operands):							
Opcode		rs		rt		rd	Funct
31	26	25	21	20	16	15	10
Imm							0
63							

Format C + E are immediately incompatible, as each instruction is classically 32 bits in length. Format C as well can only access regs 0-7.

(b) Of the remaining formats, identify any additional hardware complexities that arise from the format. [Hint: think about how you'd have to modify your Lab 4 datapath to support such operations.]

In the case of \mathcal{P} , this would necessitate an additional Max for the new imm with its larger bundle size

4. ⁺⁰ ISA vs ABI Twitter War (5 Points). In 140 chars or less (total): what is the primary purpose of an ISA? of an ABI?

The primary purpose ~~of~~ an ISA is to provide a medium to translate high level code into machine code.

An ABI does what an ISA does, but with hardware designed for that ~~that~~ specific job.

5. ⁺¹⁵ Processor Design Preparedness (15 Points).

(a) (10 points) Dissassemble the following MIPS machine instruction (i.e., list the instructions, registers, and the types and decimal values of the operand fields). Show your work!

0x00000000

R-Type:

OP	Rs	Rr	Rd	Shift	Func.
000000	00000	00000	00000	00000	000000

Sll \$0, ~~\$0~~, 0

- (b) (5 points) True or False: I am ready to design a processor. If False, why?

[✓] True (although not especially well). ^{Oh whoops}

I have learned the building blocks of a processor, and using machine code, am able to appropriately direct data from memory or immediate.

(END OF EXAM)

381 Notes Sheet


Sean Gordon
Sgordon u
C/10:00

R: ~~addu \$t0, \$s1, \$s2~~



Used for
Register/ALU
Ops.

I: lw \$t2, 0(\$t1)



Load/Store
ALU ; imm.
Conditional Branch

J: j li Used if Li is very far, too big for a branch or something



JUMP
Jump & Link

PC Relative Addressing

When branching to an address is actually
PC + 4 + imm

	...		PC	PC + 4 + imm
-1.4	0x1000	bne \$t2, \$s3, Exit	0x1000	Exit == 2 because
0.4	0x1004	add ...	↓	PC + 4 + 4 * imm = 100C =>
1.4	0x1008	lw ...		100C - PC - 4 = 4 * imm =>
2.4	0x100C	Exit: addui \$s2, \$s1, 10	0x100C	100C - 1000 - 4 = 4 * imm

J-type example: What if "Li" is very far away, too far for 16 bits?

bne \$s0, \$s1, L1

- can be rewritten as -

bne \$s0, \$s1, L2 # Inverted

j L1

Unconditional, can use 16 bits

L2:

\$1 or \$a1 is reserved for the compiler

blt \$t0, \$t1, L1 # Go to L1 if \$t0 < \$t1

↓

slt \$a1, \$t0, \$t1 # set \$a1 = 1 if \$t0 < \$t1

bne \$a1, \$0, L1 # Go to L1 if \$a1 != 0

Name	Reg #	Usage
\$0, \$zero	0	Constant 0
\$1, \$at	1	Reserved for pseudo instructions
\$v0 - \$v2	2-3	Values for results
\$a0 - \$a3	4-7	arguments/parameters
\$t0 - \$t7	8-15	temps.
\$s0 - \$s7	16-23	saved regs
\$t8 - \$t9	24-25	temps 2
\$gp	28	Global pointer
\$sp	29	stack pointer
\$fp	30	Frame pointer
\$ra	31	return address

\$s0 - \$s7 guaranteed to be preserved over a function call (convention)
 \$t0 - \$t9 not guaranteed

\$gp points to the center of the 6th block of memory in static data
 \$sp points to the last location on stack
 \$fp points to beginning of stack, not really