

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

Lecture 11: Inter-Process Communication (IPC) I



Agenda

- **Recap**
- **Inter-Process Communication I**
 - **Basic Concepts**
 - **Race Condition & Critical Region**
 - **Solutions of Mutual Exclusion**

Recap

- Scheduling Algorithms
 - First-Come, First-Served (FCFS)
 - Shortest-Job-First (SJF)
 - Shortest Remaining Time Next
 - Round Robin (RR)
 - Priority Scheduling

Recap

- Scheduling Algorithms


- Consider 5 processes arriving in the order of P1-P5; all are ready at time 0
- Draw Gantt Chart for each algorithm and calculate average waiting time

Process	Burst Time
P_1	10
P_2	29
P_3	3
P_4	7
P_5	12


- FCFS: 28



- SJF: 13



- RR (quantum = 10): 23



Agenda

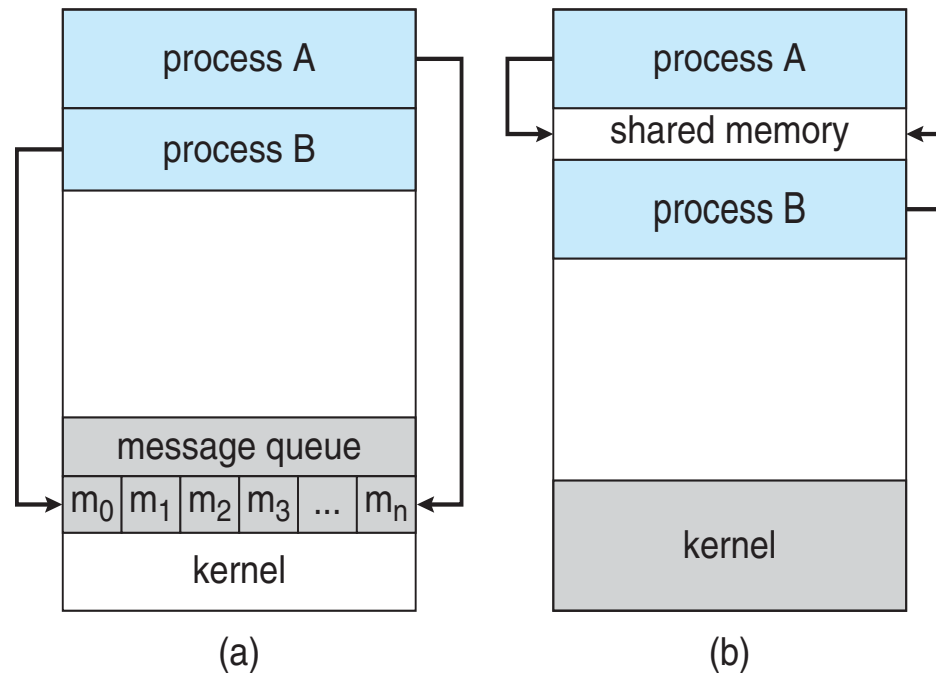
- ~~Recap~~
- **Inter-Process Communication I**
 - **Basic Concepts**
 - **Race Condition & Critical Region**
 - **Solutions of Mutual Exclusion**

Basic Concepts

- Processes within a system may be independent or cooperating
 - Example reasons for cooperating:
 - Information sharing
 - Computation speedup
 - ...
 - Cooperating processes need **inter-process communication (IPC)** mechanism:
 - Communicate with each other
 - Ensure
 - (1) do not get in each other's way
 - (2) proper ordering when dependencies are present
 - Violating the two may lead to “concurrency bugs”

Basic Concepts

- Two basic methods of IPC
 - Message passing
 - Shared memory

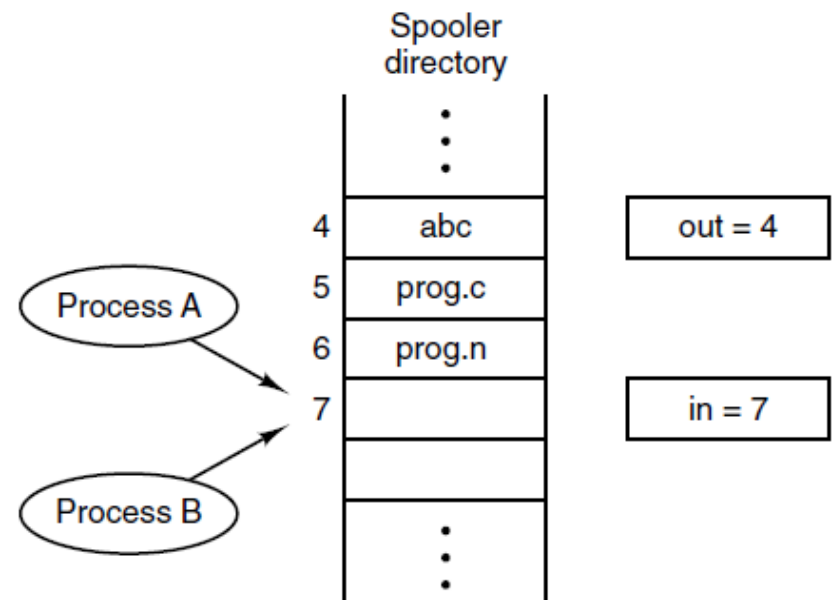


Agenda

- ~~Recap~~
- **Inter-Process Communication I**
 - ~~Basic Concepts~~
 - **Race Condition & Critical Region**
 - **Solutions of Mutual Exclusion**

Race Condition

- two or more processes are reading or writing some shared data and the final result depends on who runs precisely when
- E.g., a printer with two processes
 - Two shared variables
 - out: next file to be printed
 - in: next free slot



Race Condition

- E.g., two threads perform “counter = counter + 1”
 - “counter” is a shared variable
 - Initially counter = 50
 - Possible result?

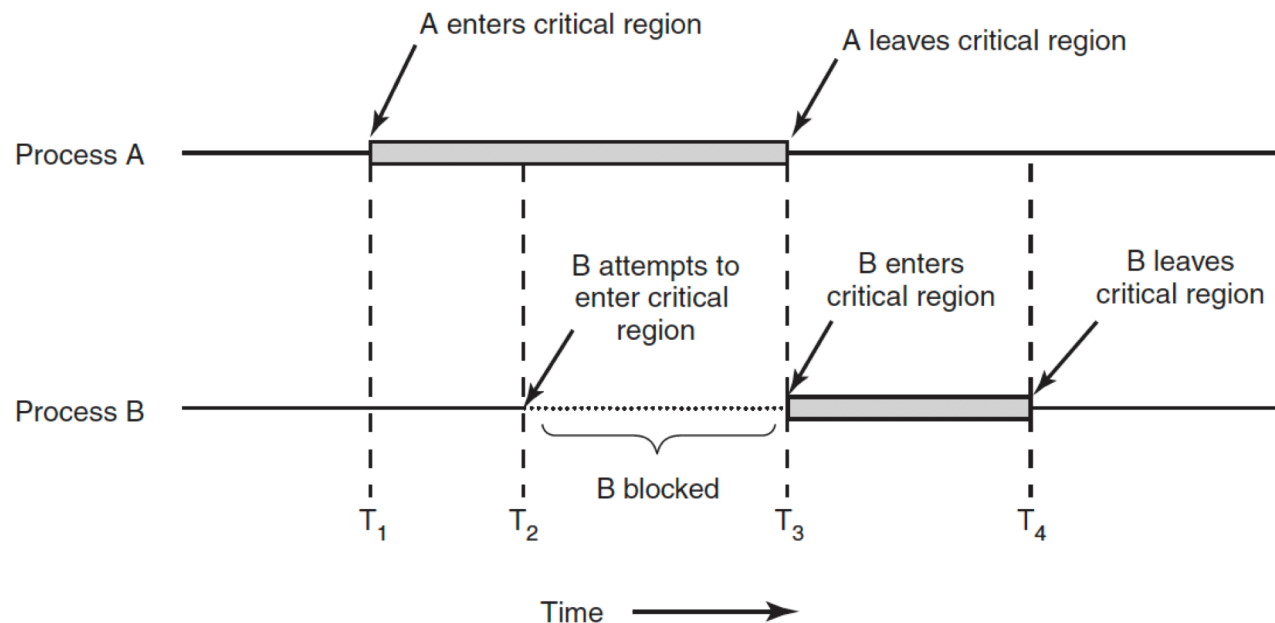
Race Condition

- E.g., two threads perform “counter = counter + 1”
 - “counter” is a shared variable
 - Initially counter = 50

OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
			100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
save T1's state					
restore T2's state			100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
interrupt					
save T2's state					
restore T1's state			108	51	50
	mov %eax, 0x8049a1c		113	51	51

Critical Region

- A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread
 - Multiple threads executing critical section can result in a race condition.
 - Need to support **mutual exclusion** or critical sections



Critical Region

- Four conditions for a good solution
 - No two processes may be simultaneously inside their critical regions
 - No assumptions may be made about speeds or the number of CPUs
 - No process running outside its critical region may block any process
 - No process should have to wait forever to enter its critical region

Agenda

- ~~Recap~~
- **Inter-Process Communication I**
 - ~~Basic Concepts~~
 - ~~Race Condition & Critical Region~~
 - **Solutions of Mutual Exclusion**

Solutions of Mutual Exclusion

- Disabling Interrupts
 - Disable interrupts when a process is in critical region
 - Cons
 - Not safe
 - What if a process does not turn them on?
 - More suitable for kernel instead of user processes
 - Not applicable for multiple processors
 - processes on other processors may still access shared variables
 - Less used today

Solutions of Mutual Exclusion

- Strict Alternation
 - Use an variable to keep track of whose turn it is to enter the critical region

```
while (TRUE) {  
    while (turn != 0)      /* loop */ ;  
    critical_region();  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)      /* loop */ ;  
    critical_region();  
    turn = 0;  
    noncritical_region();  
}
```

(b)

Solutions of Mutual Exclusion

- Strict Alternation
 - Use an variable to keep track of whose turn it is to enter the critical region

```
while (TRUE) {  
    while (turn != 0)      /* loop */ ;  
    critical_region();  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)      /* loop */ ;  
    critical_region();  
    turn = 0;  
    noncritical_region();  
}
```

(b)

- Cons
 - **Busy waiting** (continuously testing a variable until some value appears) wastes CPU cycles
 - A lock that uses busy waiting is called a **spin lock**

Agenda

- **Recap**

Questions?

- **Inter-Process Communication I**

- **Basic Concepts**

- **Race Condition & Critical Region**

- **Solutions of Mutual Exclusion**



*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.