

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

Lecture 09: CPU Scheduling I

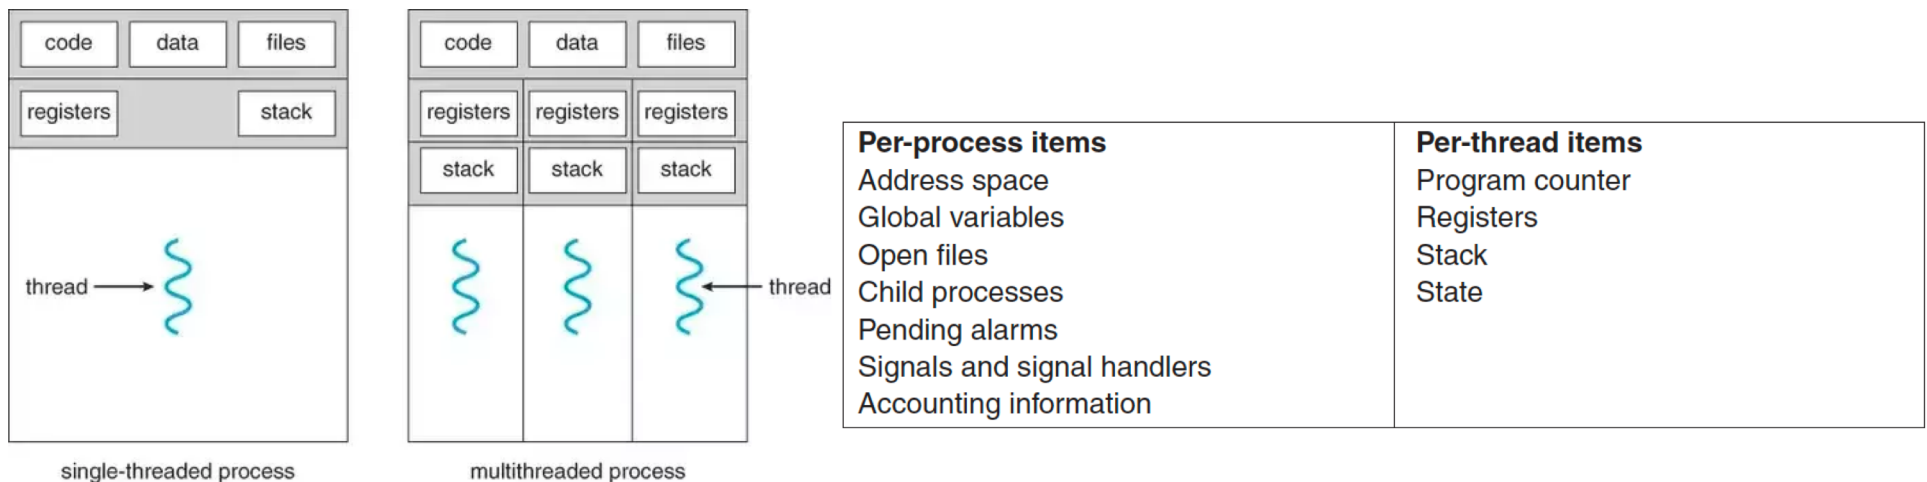


Agenda

- **Recap**
- **CPU Scheduling I**
 - **Scheduling Concept**
 - **POSIX Threads**

Recap

- Thread Concept
 - Multiple threads of control *within* a process
 - All threads within a process share the same
 - text, heap, static data segments, open files ...
 - each thread has its own
 - State, PC, registers, stack



Recap

- POSIX Threads
 - Thread creation

```
#include <pthread.h>

Int pthread_create(  pthread_t*      thread,
                    const pthread_attr_t* attr,
                    void*             (*start_routine) (void*),
                    void*             arg);
```

- Thread join

```
int pthread_join(pthread_t thread, void **value_ptr);
```

Recap

- Example

```
1  typedef struct __myarg_t {
2      int a;
3      int b;
4  } myarg_t;
5
6  typedef struct __myret_t {
7      int x;
8      int y;
9  } myret_t;
10
11 void *mythread(void *arg) {
12     myarg_t *m = (myarg_t *) arg;
13     printf("%d %d\n", m->a, m->b);
14     myret_t *r = malloc(sizeof(myret_t));
15     r->x = 1;
16     r->y = 2;
17     return (void *) r;
18 }
```

Recap

- Example (cont')

```
19. int main(int argc, char *argv[]) {
20.     int rc;
21.     pthread_t p;
22.     myret_t *m;
23.
24.     myarg_t args;
25.     args.a = 10;
26.     args.b = 20;
27.     pthread_create(&p, NULL, mythread, &args);
28.     pthread_join(p, (void **) &m);
29.     printf("returned %d %d\n", m->x, m->y);
30.     free(m);
31.     return 0;
32. }
```

Recap

- Be careful with how values are returned from a thread

```
1  void *mythread(void *arg) {  
2      myarg_t *m = (myarg_t *) arg;  
3      printf("%d %d\n", m->a, m->b);  
4      myret_t r; // ALLOCATED ON STACK!  
5      r.x = 1;  
6      r.y = 2;  
7      return (void *) &r;  
8  }
```


- When the function returns, `r` is automatically deallocated (i.e., the stack frame is destroyed)
 - pointer to `r` will point to invalid data

Agenda

- **Recap**
- **CPU Scheduling I**
 - **Scheduling Concept**
 - **Scheduling Algorithms**


Scheduling Concept

- Process state & multiprogramming (revisit)




1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

One program counter




(a)

Four program counters



(b)




(c)

Scheduling Concept

- Process scheduling
 - When to run a different process?
 - New process is created
 - Process terminates
 - A process blocks on I/O etc.
 - interrupt occurs
 - Which one to run?
 - Depends on scheduling algorithms

Scheduling Concept

- Process behavior
 - CPU-bound (compute-bound)
 - I/O-bound



Scheduling Concept

- Non-preemptive V.S. preemptive scheduling
 - Non-preemptive
 - Processes run until they are blocked (for I/O) or voluntarily releases the CPU (e.g., terminates)
 - Preemptive
 - Can forcibly suspend a process and switch to another

Scheduling Concept

- Goals of scheduling algorithms
 - General goals
 - Fairness
 - giving each process a fair share of the CPU
 - Policy enforcement
 - ensuring that desired policy is carried out
 - Balance
 - keeping all parts of the system busy
 - Different systems may have different specific goals

Scheduling Concept

- Goals of scheduling algorithms
 - Specific goals for batch systems
 - Throughput
 - maximize jobs per hour
 - Turnaround time
 - minimize time between submission and termination
 - CPU utilization
 - keep the CPU busy all the time

Scheduling Concept

- Goals of scheduling algorithms
 - Specific goals for batch systems
 - Throughput
 - maximize jobs per hour
 - Turnaround time
 - minimize time between submission and termination
 - CPU utilization
 - keep the CPU busy all the time

Scheduling Concept

- Goals of scheduling algorithms
 - Specific goals for interactive systems
 - Response time
 - respond to user requests quickly
 - Proportionality
 - meet users' expectations

Scheduling Concept

- Goals of scheduling algorithms
 - Specific goals for real-time systems
 - Meeting deadlines
 - hard real time
 - there are absolute deadlines that must be met
 - soft real time
 - missing an occasional deadline is undesirable, but nevertheless tolerable

Agenda

- **Recap**
- **CPU Scheduling I**
 - **Scheduling Concept**
 - **Scheduling Algorithms**

Scheduling Algorithms

- First-Come, First-Served (FCFS)
 - Scheduling based on the arrival order of processes

Scheduling Algorithms

- First-Come, First-Served (FCFS)
 - Scheduling based on the arrival order of processes

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Assume the processes arrive in the order: P_1 , P_2 , P_3
- The **Gantt Chart** for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

Scheduling Algorithms

- First-Come, First-Served (FCFS)
 - Scheduling based on the arrival order of processes

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Assume the processes arrive in the order: P_2 , P_3 , P_1
- The Gantt Chart for the schedule is:
- Waiting time:
- Average waiting time:

Scheduling Algorithms

- First-Come, First-Served (FCFS)
 - Scheduling based on the arrival order of processes

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Assume the processes arrive in the order: P_2 , P_3 , P_1
- The Gantt Chart for the schedule is:




- Waiting time: $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$

Scheduling Algorithms

- Shortest-Job-First (SJF)
 - Scheduling based on (predicted) CPU burst time

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

- Assume all processes are ready at time 0
- Gantt Chart:



- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$
- Turnaround time: P1 = 9; P2 = 24; P3 = 16; P4 = 3

Agenda

- **Recap**
- **CPU Scheduling I**
 - **Scheduling Concept**
 - **Scheduling Algorithms**

Questions?



*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.