

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

Lecture 12: Inter-Process Communication (IPC) II

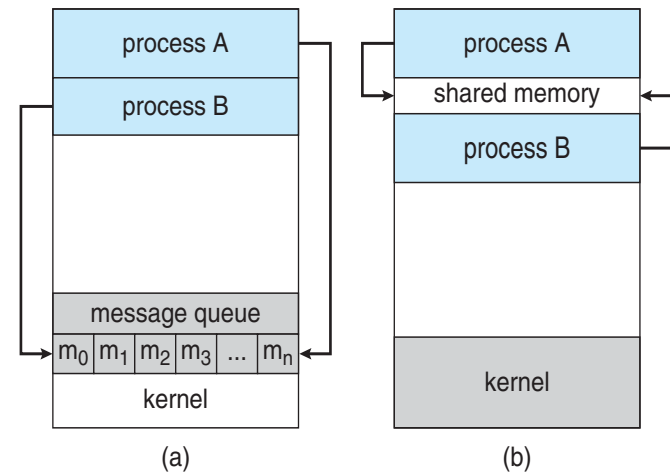


Agenda

- **Recap**
- **Inter-Process Communication II**
 - **Solutions of Mutual Exclusion (cont.)**

Recap

- IPC Concepts
 - Cooperating processes need **inter-process communication (IPC)** mechanism:
 - Communicate with each other
 - Ensure
 - (1) do not get in each other's way
 - (2) proper ordering when dependencies are present
 - Two basic methods of IPC
 - Message passing
 - Shared memory



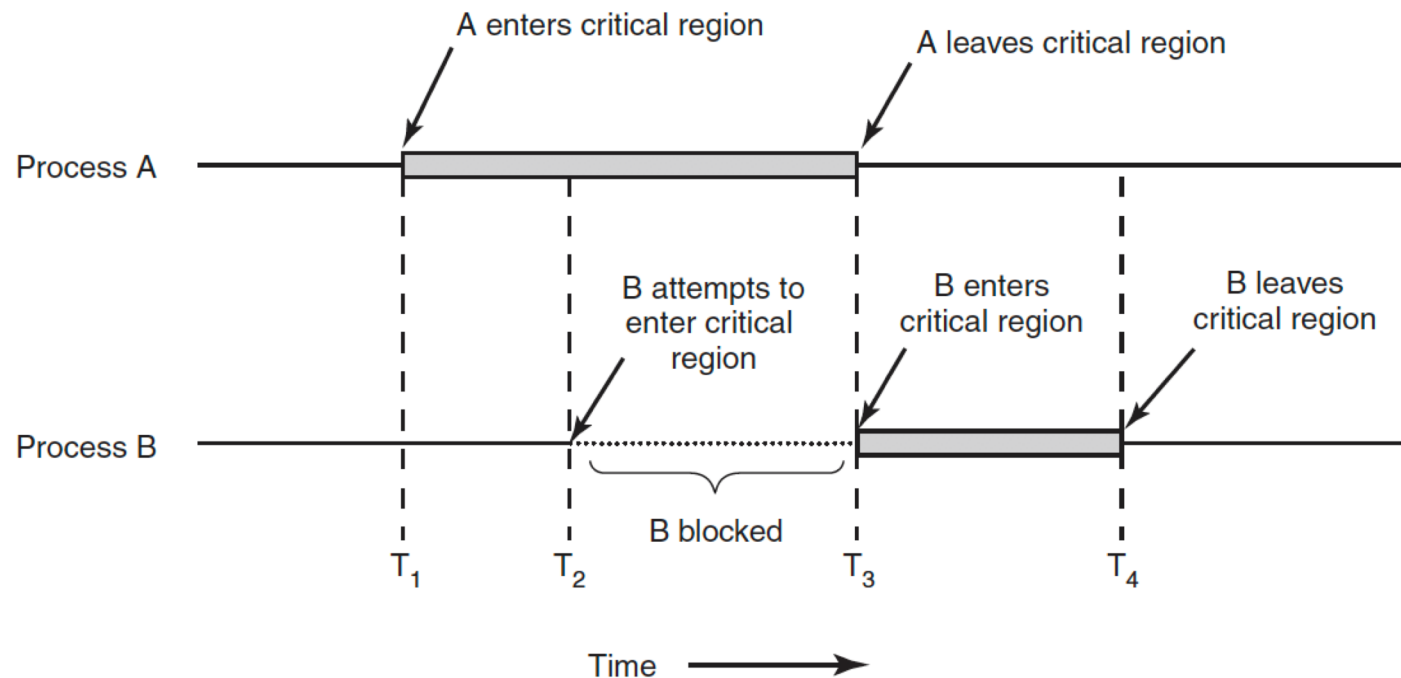
Recap

- Race Condition
 - E.g., two threads perform “counter = counter + 1”
 - “counter” is a shared variable

| OS | Thread1 | Thread2 | (after instruction) | | |
|--------------------|---------------------|---------------------|---------------------|------|---------|
| | | | PC | %eax | counter |
| | | | 100 | 0 | 50 |
| | mov 0x8049a1c, %eax | | 105 | 50 | 50 |
| | add \$0x1, %eax | | 108 | 51 | 50 |
| interrupt | | | | | |
| save T1's state | | | | | |
| restore T2's state | | | 100 | 0 | 50 |
| | | mov 0x8049a1c, %eax | 105 | 50 | 50 |
| | | add \$0x1, %eax | 108 | 51 | 50 |
| | | mov %eax, 0x8049a1c | 113 | 51 | 51 |
| interrupt | | | | | |
| save T2's state | | | | | |
| restore T1's state | | | 108 | 51 | 50 |
| | mov %eax, 0x8049a1c | | 113 | 51 | 51 |

Recap

- Critical Region (Critical Section)
 - A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread
 - Need to support **mutual exclusion**



Recap

- Solutions of Mutual Exclusion
 - Disabling Interrupts
 - Disable interrupts when a process is in critical region
 - Cons
 - Not safe
 - Not applicable for multiple processors
 - Strict Alternation
 - Use an variable to keep track of whose turn it is to enter the critical region
 - Cons
 - **Busy waiting** (continuously testing a variable until some value appears) wastes CPU cycles
 - requires two processes strictly alternate in entering their critical regions

Agenda

- ~~Recap~~
- **Inter-Process Communication II**
 - **Solutions of Mutual Exclusion (cont.)**

Solutions of Mutual Exclusion

- Peterson's Solution (1981)
 - General Structure

```
enter_region();  
critical_region();  
leave_region();
```


Solutions of Mutual Exclusion

- Peterson's Solution (1981)

```
#define FALSE 0
#define TRUE  1
#define N      2                /* number of processes */

int turn;                       /* whose turn is it? */
int interested[N];              /* all values initially 0 (FALSE) */

void enter_region(int process); /* process is 0 or 1 */
{
    int other;                  /* number of the other process */

    other = 1 - process;        /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;             /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process) /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

Solutions of Mutual Exclusion

- Hardware Support
 - special instruction: TSL (Test and Set Lock)
 - TSL R, Lock
 - TSL = “Test and Set Lock”
 - R = register,
 - Lock = memory location
 - reads the contents of the memory word Lock into register R, and then stores a nonzero value at the memory address Lock.
 - Atomic (“all or nothing”/indivisible)
 - reading the word and storing into it are guaranteed to be indivisible—no other processor can access the memory word until the instruction is finished
 - The CPU executing the TSL instruction locks the memory bus to prohibit other CPUs from accessing memory until it is done
 - different from disabling interrupt

Solutions of Mutual Exclusion

- Hardware Support
 - special instruction: TSL (Test and Set Lock)
 - Example

enter_region:

```
TSL REGISTER,LOCK
CMP REGISTER,#0
JNE enter_region
RET
```

```
| copy lock to register and set lock to 1
| was lock zero?
| if it was not zero, lock was set, so loop
| return to caller; critical region entered
```

leave_region:

```
MOVE LOCK,#0
RET
```

```
| store a 0 in lock
| return to caller
```

Solutions of Mutual Exclusion

- Hardware Support
 - another instruction: XCHG (exchange)
 - exchanges the contents of a register and a memory location atomically

enter_region:

```
MOVE REGISTER,#1
XCHG REGISTER,LOCK
CMP REGISTER,#0
JNE enter_region
RET
```

```
| put a 1 in the register
| swap the contents of the register and lock variable
| was lock zero?
| if it was non zero, lock was set, so loop
| return to caller; critical region entered
```

leave_region:

```
MOVE LOCK,#0
RET
```

```
| store a 0 in lock
| return to caller
```

Summary of Solutions So Far

- Software solution
 - Disabling interrupts
 - Single processor only
 - Use for kernel
 - Strict alternation
 - Strict ordering
 - Busy waiting
 - Peterson's solution
 - Busy waiting
- Hardware solution
 - TSL/XCHG
 - Work on multiprocessors
 - Busy waiting

Better solution?

Better solution?

- Sleep and Wakeup
 - Yield to other thread/process if unable to lock
 - If cannot enter critical region, calls *sleep()* to give up CPU
 - wakes up another thread using *wakeup()*

Pthread Locks

- Provide mutual exclusion to a critical section
 - Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Usage (*w/o lock initialization and error check*)

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- No other thread holds the lock → the thread will acquire the lock and enter the critical section.
- If another thread hold the lock → the thread will not return from the call until it has acquired the lock.

Pthread Locks

- All locks must be properly initialized.
 - One way: using `PTHREAD_MUTEX_INITIALIZER`

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- The dynamic way: using `pthread_mutex_init()`

```
int rc = pthread_mutex_init(&lock, NULL);  
assert(rc == 0); // always check success!
```

Pthread Locks

- These two calls are also used in **lock acquisition**

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);  
int pthread_mutex_timelock(pthread_mutex_t *mutex,  
                           struct timespec *abs_timeout);
```

- `trylock`: return failure if the lock is already held
- `timelock`: return after a timeout or after acquiring the lock

Agenda

- **Recap**

- **Inter-Process Communication II**

- **Solutions of Mutual Exclusion (cont.)**

Questions?



*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.