

# IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

## Lecture 13: Inter-Process Communication (IPC) III



# Agenda

- **Recap**
- **Inter-Process Communication III**
  - **Semaphore**
  - **Conditional Variables**

# Recap

- Solutions of Mutual Exclusion
  - Peterson's Solution

```
#define FALSE 0
#define TRUE 1
#define N      2                /* number of processes */

int turn;                       /* whose turn is it? */
int interested[N];              /* all values initially 0 (FALSE) */

void enter_region(int process);  /* process is 0 or 1 */
{
    int other;                  /* number of the other process */

    other = 1 - process;        /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;             /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process)   /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

# Recap

- Solutions of Mutual Exclusion
  - Hardware Support
    - special instruction: TSL (Test and Set Lock)

enter\_region:

```
TSL REGISTER,LOCK  
CMP REGISTER,#0  
JNE enter_region  
RET
```

```
| copy lock to register and set lock to 1  
| was lock zero?  
| if it was not zero, lock was set, so loop  
| return to caller; critical region entered
```

leave\_region:

```
MOVE LOCK,#0  
RET
```

```
| store a 0 in lock  
| return to caller
```

# Recap

- Solutions of Mutual Exclusion
  - Hardware Support
    - special instruction: XCHG (exchange)

enter\_region:

```
MOVE REGISTER,#1
XCHG REGISTER,LOCK
CMP REGISTER,#0
JNE enter_region
RET
```

```
| put a 1 in the register
| swap the contents of the register and lock variable
| was lock zero?
| if it was non zero, lock was set, so loop
| return to caller; critical region entered
```

leave\_region:

```
MOVE LOCK,#0
RET
```

```
| store a 0 in lock
| return to caller
```

# Recap

- Pthread Locks
  - Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Lock Initialization

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

```
int rc = pthread_mutex_init(&lock, NULL);  
assert(rc == 0); // always check success!
```

- alternatives

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);  
int pthread_mutex_timelock(pthread_mutex_t *mutex,  
                           struct timespec *abs_timeout);
```

# Agenda

- ~~Recap~~
- **Inter-Process Communication III**
  - Semaphore
  - Conditional Variables

# Semaphore

- Synchronization method that provides more sophisticated ways (than mutex locks) for process to synchronize their activities
- Semaphore **S** – integer variable
  - Can only be accessed via two indivisible (atomic) operations
    - **down()** and **up()**
      - originally called **P()** and **V()**
      - also called **wait()** and **signal()**



# Semaphore

- Semaphore **S**
  - Definition of the **down()** operation

```
down(S) {  
    while (S <= 0)  
        ; // busy waiting  
    S--;  
}
```
  - Definition of the **up()** operation

```
up(S) {  
    S++;  
}
```


# Semaphore

- **Counting semaphore**
    - integer value can range over an unrestricted domain
  - **Binary semaphore**
    - integer value can range only between 0 and 1
    - same as a **mutex lock**
  - Can solve various synchronization problems
    - e.g., consider P1 and P2 that require S1 to happen before S2
      - Create a semaphore “synch” initialized to 0
- ```
P1:  
    S1; up(synch);  
P2:  
    down(synch); S2;
```


# Producer-Consumer Problem

- Shared bounded buffer
  - A “Producer” process inserts item
  - A “Consumer” process removes item

Producer



Consumer



# Producer-Consumer Problem

- Solution with semaphore

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

```
/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */
```

# Producer-Consumer Problem

- Solution with semaphore

```
void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}
```

/\* TRUE is the constant 1 \*/  
/\* generate something to put in buffer \*/  
/\* decrement empty count \*/  
/\* enter critical region \*/  
/\* put new item in buffer \*/  
/\* leave critical region \*/  
/\* increment count of full slots \*/

# Producer-Consumer Problem

- Solution with semaphore

```
void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item( );
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```

/\* infinite loop \*/  
/\* decrement full count \*/  
/\* enter critical region \*/  
/\* take item from buffer \*/  
/\* leave critical region \*/  
/\* increment count of empty slots \*/  
/\* do something with the item \*/

# Use Semaphore with Caution

- **Deadlock**

- two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Eg., Let **S** and **Q** be two semaphores initialized to 1

| $P_0$    | $P_1$    |
|----------|----------|
| down(S); | down(Q); |
| down(Q); | down(S); |
| ...      | ...      |
| up(S);   | up(Q);   |
| up(Q);   | up(S);   |

# Use Semaphore with Caution

- **Starvation** – indefinite blocking
  - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process



# Conditional Variables

- Allows a thread to wait till a condition is satisfied
- Testing the condition must be done within a mutex
- A mutex is associated with every condition variable

| Thread call            | Description                                  |
|------------------------|----------------------------------------------|
| Pthread_cond_init      | Create a condition variable                  |
| Pthread_cond_destroy   | Destroy a condition variable                 |
| Pthread_cond_wait      | Block waiting for a signal                   |
| Pthread_cond_signal    | Signal another thread and wake it up         |
| Pthread_cond_broadcast | Signal multiple threads and wake all of them |

# Conditional Variables

- Comparison with semaphore
  - If a signal is sent to a conditional variable on which no thread is waiting, the signal is lost
  - Semaphore will accumulate 'signals' by up()
- Comparison with mutex lock
  - Mutex lock is good to guarantee mutual exclusion
    - Allow and block access to the critical regions
  - Conditional variables block threads due to some condition not met

# Conditional Variables

- Typical Usage

```
pthread_cond_t condition_variable;  
pthread_mutex_t mutex;
```

Waiting Thread:

```
pthread_mutex_lock(&mutex);  
  
while (condition not satisfied) {  
    pthread_cond_wait(  
        &condition_variable,  
        &mutex);  
    }  
  
pthread_mutex_unlock(&mutex);
```

Signaling Thread:

```
pthread_mutex_lock(&mutex);  
  
/* change variable value */  
  
if (condition satisfied) {  
    pthread_cond_signal(  
        &condition_variable);  
    }  
  
pthread_mutex_unlock(&mutex);
```

# Conditional Variables

- Example
  - Write a program using two threads
    - Thread 1 prints “hello”
    - Thread 2 prints “world”
    - Thread 2 should wait till thread 1 finishes before printing
  - Use condition variables

# Conditional Variables

- Example

```
int  thread1_done = 0;

pthread_cond_t  cv;
pthread_mutex_t mutex;
```

Thread 1:

```
printf("hello ");
```

```
pthread_mutex_lock(&mutex);
thread1_done = 1;
pthread_cond_signal(&cv);
pthread_mutex_unlock(&mutex);
```

Thread 2:

```
pthread_mutex_lock(&mutex);

pthread_cond_wait(&cv, &mutex);

printf(" world\n");
pthread_mutex_unlock(&mutex);
```

# Conditional Variables

- Example

```
int  thread1_done = 0;

pthread_cond_t  cv;
pthread_mutex_t mutex;
```

Thread 1:

```
printf("hello ");
```

```
pthread_mutex_lock(&mutex);
thread1_done = 1;
pthread_cond_signal(&cv);
pthread_mutex_unlock(&mutex);
```

Thread 2:

```
pthread_mutex_lock(&mutex);
```

```
while (thread1_done == 0) {
    pthread_cond_wait(&cv,
    &mutex);
}
```

```
printf(" world\n");
pthread_mutex_unlock(&mutex);
```

# Agenda

- **Recap**

- **Inter-Process Communication III**

- **Semaphore**

- **Conditional Variables**

Questions?



\*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.