

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

Lecture 15: Deadlocks



Agenda

- **Recap**
- **Deadlocks**
 - **Deadlock Concepts**
 - **Deadlock Solutions**

Recap

- Conditional Variables
 - Allows a thread to wait till a condition is satisfied
 - Testing the condition must be done within a mutex
 - A mutex is associated with every condition variable
 - A mutex is passed into wait:
`pthread_cond_wait(cond_var, mutex)`
 - **Mutex is unlocked before the thread sleeps**
 - **Mutex is locked again before pthread_cond_wait() returns**
 - Safe to use pthread_cond_wait() in a while loop and check condition again before proceeding

Recap

- Example

```
int  thread1_done = 0;

pthread_cond_t  cv;
pthread_mutex_t mutex;
```

Thread 1:

```
printf("hello ");
```

```
pthread_mutex_lock(&mutex);
thread1_done = 1;
pthread_cond_signal(&cv);
pthread_mutex_unlock(&mutex);
```

Thread 2:


```
pthread_mutex_lock(&mutex);
```

```
while (thread1_done == 0) {
    pthread_cond_wait(&cv,
                     &mutex);
}
```

```
printf(" world\n");
pthread_mutex_unlock(&mutex);
```

Recap

- Dining Philosophers Problem
 - Philosopher
 - eat, think
 - eat, think
 -
 - Philosopher = Process
 - Eating needs two resources
 - Issues with solutions
 - Deadlock
 - Starvation
 - Poor performance



Recap

- Readers-Writers Problem
 - Multiple threads reading/writing
 - Many threads can read simultaneously
 - Only one can be writing at any time
 - When a writer is executing, nobody else can read or write
 - One solution idea
 - Readers:
 - First reader locks the database
 - If a reader inside, other readers enter without locking again
 - Checking for readers occurs within a mutex
 - Writer:
 - Always lock database before entering
 - May run into starvation

Recap

- Example solution

READER:

```
While (1) {  
    down(protector);  
    rc++;  
    if (rc == 1) //first reader  
        down(database);  
    up(protector);  
  
    read();  
  
    down(protector);  
    rc--;  
    If (rc == 0) then // last one  
        up(database);  
    up(protector);  
    ....  
}
```

WRITER:

```
While (1) {  
    generate_data();  
    down(database);  
    write();  
    up(database);  
}
```

Two semaphores:

database

protector

Initial: protector=1, database =1

rc =0

Agenda

- ~~Recap~~
- **Deadlocks**
 - **Deadlock Concepts**
 - **Deadlock Solutions**

Deadlock Concepts

- What types of bugs exist?
 - A survey on four major open-source applications
 - MySQL, Apache, Mozilla, OpenOffice.

Application	What it does	Non-Deadlock	Deadlock
MySQL	Database Server	14	9
Apache	Web Server	13	4
Mozilla	Web Browser	41	16
Open Office	Office Suite	6	2
Total		74	31

- Non-deadlock bugs:
 - Atomicity violation, ordering violation, ...

Deadlock Concepts

- What's deadlock?
 - “When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone.”
 - A Treasury of Railroad Folklore, B.A. Botkin & A.F. Harlow, p. 381

Deadlock Concepts

- A more formal definition
 - A set of processes/threads is deadlocked if each process/thread in the set is waiting for an event that only another process/thread in the set can cause
- Usually the event is release of a currently held resource
- None of the processes/threads can ...
 - run
 - release resources

Deadlock Concepts

- Why do deadlocks occur?

Deadlock Concepts

- Why do deadlocks occur?
 - Reason 1:
 - In large code bases, **complex dependencies** arise between components.
 - Reason 2:
 - Due to the nature of **encapsulation**
 - Hide details of implementations and make software easier to build in a modular way.
 - Such **modularity** *does not mesh* well with locking.

Deadlock Concepts

- Four conditions need to hold for a deadlock to occur

Condition	Description
Mutual Exclusion	Threads claim exclusive control of resources that they require.
Hold-and-wait	Threads hold resources allocated to them while waiting for additional resources
No preemption	Resources cannot be forcibly removed from threads that are holding them.
Circular wait	There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain




- If any of these four conditions are not met, **deadlock cannot occur.**

Agenda

- ~~Recap~~
- Deadlocks
 - ~~Deadlock Concepts~~
 - Deadlock Solutions

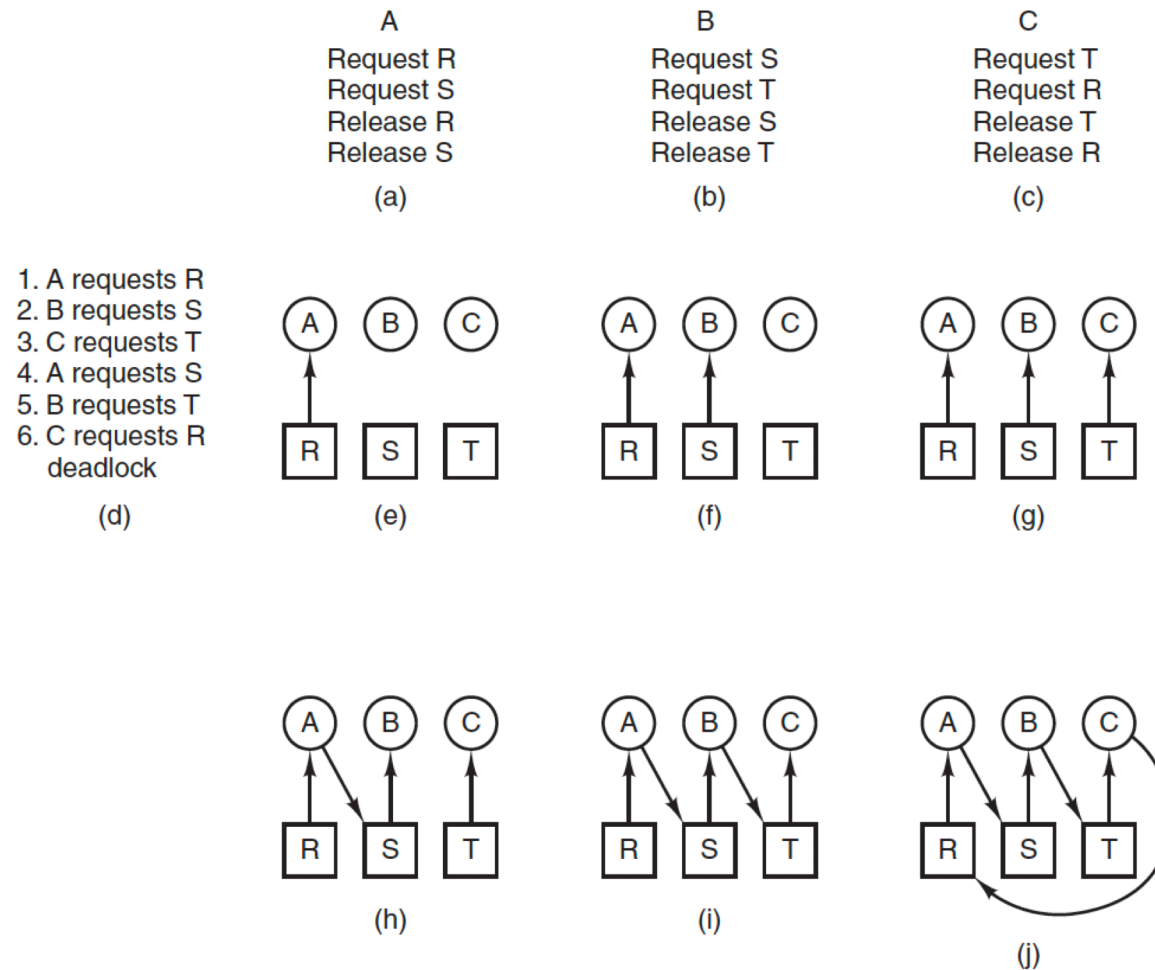
Deadlock Solutions

- Deadlock modeling
 - **Resource allocation graph**
 - resource R assigned to process A
 - process B is requesting/waiting for resource S
 - process C and D are in deadlock over resources T and U



Deadlock Solutions

- Deadlock modeling



Deadlock Solutions

- Deadlock detection & recovery
 - **Allow deadlock** to occasionally occur and then *take some action*.
 - E.g.: if an OS froze, you would reboot it.
 - Many database systems employ *deadlock detection and recovery technique*
 - A deadlock detector **runs periodically**.
 - Building a **resource allocation graph** and checking it for cycles.
 - In deadlock, the **system performs recovery**
 - Preemption, rollback, killing process, ...

Deadlock Solutions


- Deadlock avoidance via scheduling
 - Do not allow deadlock to occur
 - Global knowledge is required:
 - Which locks various threads might grab during their execution.
 - Subsequently schedules said threads in a way as to guarantee no deadlock can occur.

Deadlock Solutions

- Deadlock avoidance via scheduling
 - E.g., we have two processors and four threads.
 - Lock acquisition demands of the threads:

	T1	T2	T3	T4
L1	yes	yes	no	no
L2	yes	yes	yes	no

- A smart scheduler could compute that as long as T1 and T2 are not run at the same time, **no deadlock** could ever arise.




Deadlock Solutions

- Deadlock avoidance via scheduling
 - E.g., we have two processors and four threads.
 - Lock acquisition demands of the threads:
 - More contention for the same resources

	T1	T2	T3	T4
L1	yes	yes	yes	no
L2	yes	yes	yes	no

- A possible schedule that guarantees that *no deadlock* could ever occur.



Deadlock Solutions

- Deadlock prevention
 - Lock ordering
 - Provide **a total ordering** on lock acquisition
 - require *careful design* of global locking strategies.
 - E.g.: there are two locks in the system (L1 and L2); we can prevent deadlock by always acquiring L1 before L2
 - Trylock()
 - Grab the lock (if it is available).
 - Or, return -1: you should try again later.

```
1  top:
2      lock(L1);
3      if( tryLock(L2) == -1 ){
4          unlock(L1);
5          goto top;
6      }
```

Agenda

- **Recap**
- **Deadlocks**
 - **Deadlock Concepts**
 - **Deadlock Solutions**

Questions?



*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.