

# IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

## Lecture 22: TLB & Multi-Level Page Tables



# Agenda

- **Recap**
- **Translation Lookaside Buffer (TLB)**
- **Multi-Level Page Tables**

# Recap

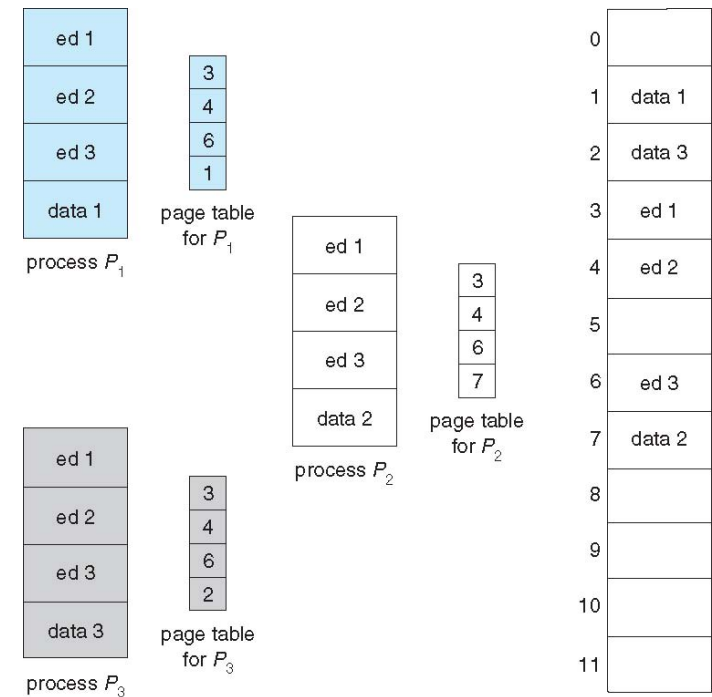

- Where is page table
  - In physical memory
  - accessible by kernel
  - Page-table base register (PTBR) points to the page table
  - Page-table length register (PTLR) indicates size of the page table

# Recap

- What is in the page table
  - Page table entries (PTEs)
    - base address of each page in the physical memory
    - may also contain additional bits for other purposes (e.g., protection)
      - **Valid Bit**: Indicating whether the particular translation is valid
      - **Protection Bit**: Indicating whether the page could be read from, written to, or executed from
      - **Present Bit**: Indicating whether this page is in physical memory or on disk(swapped out)
      - **Dirty Bit**: Indicating whether the page has been modified since it was brought into memory
      - **Reference Bit(Accessed Bit)**: Indicating that a page has been accessed

# Recap

- Address translation using a page table
  - Enable *shared pages*



# Recap

- Page Fault

- Requested page not in the physical memory
  - MMU reports a page fault to CPU
  - CPU gives control to the OS (page fault handler routine)
  - OS fetches a page from the disk
    - May need to evict an existing page from memory
  - Instruction is restarted

- Potential Issues of Paging

- Time
  - for every memory access, one additional access is needed
- Space
  - A page table can get awfully large
  - Each process needs to have a page table

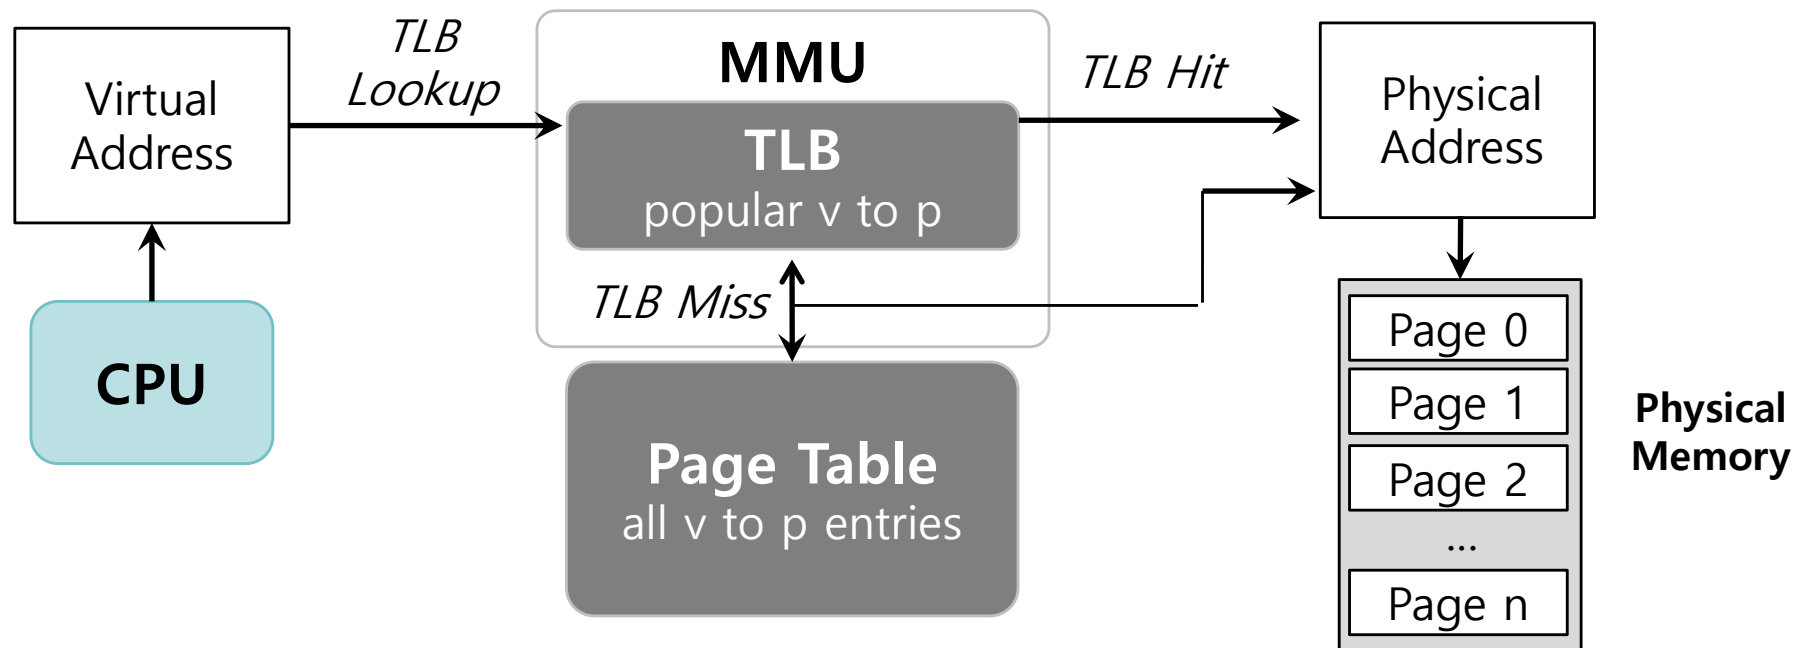
# Agenda

- ~~Recap~~

- Translation Lookaside Buffer (TLB)
- Multi-Level Page Tables

# Translation Lookaside Buffer (TLB)

- A hardware cache of popular virtual-to-physical address translation
  - Purpose: to speedup paging
  - Location: part of the memory-management unit (MMU)





# Translation Lookaside Buffer (TLB)


- Structure
  - Associate registers – parallel search

Page #	Frame #

- Address translation (Virtual Page#, Physical Frame#)
  - If virtual page# is in associative register
    - get the corresponding physical frame# from TLB
  - Otherwise
    - Walk through page table to get the frame#
- Each entry may also contain other bits
  - E.g., ASID: address space ID

# Translation Lookaside Buffer (TLB)

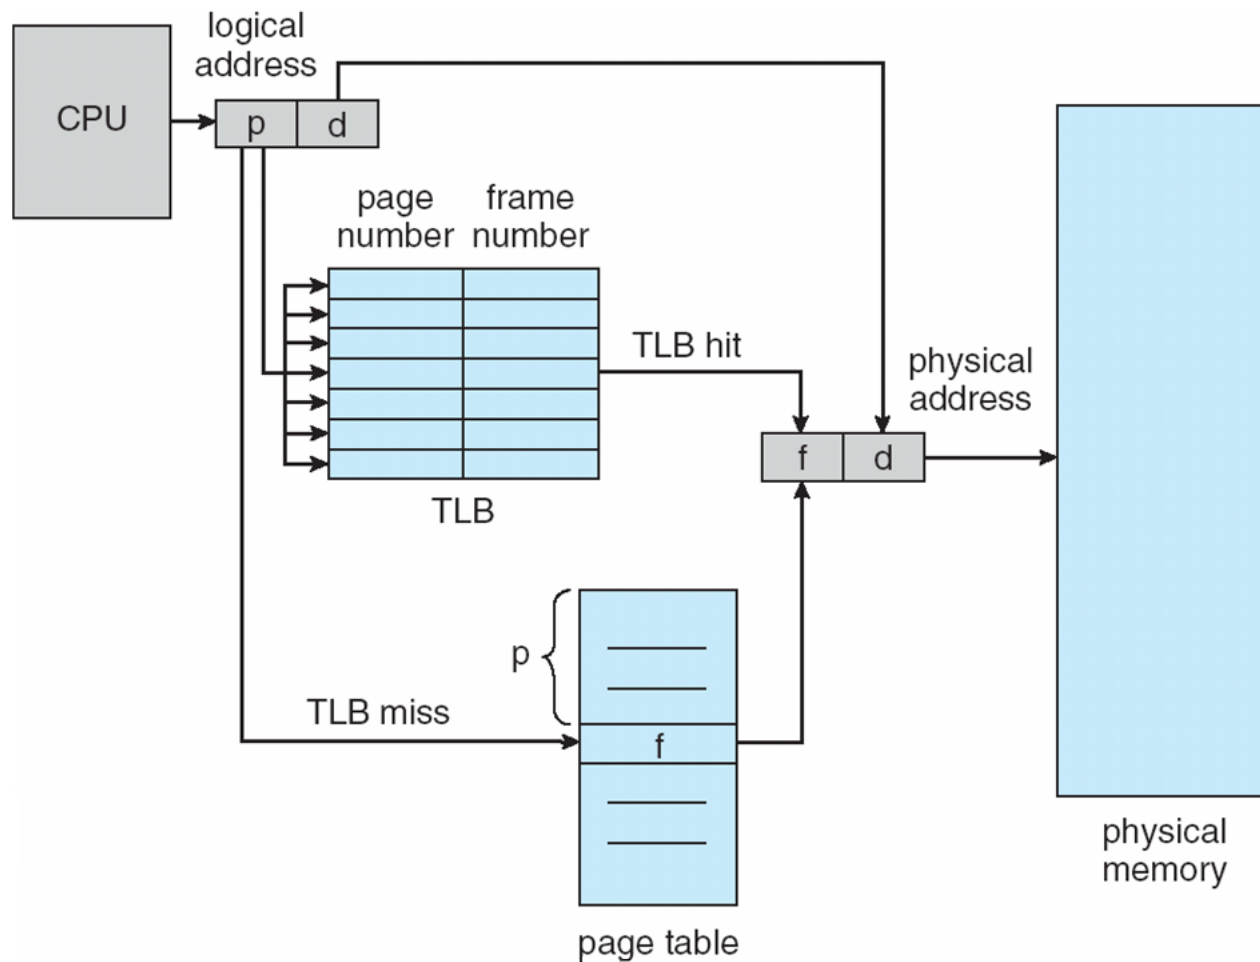
- A Real TLB Entry (MIPS R4000)
  - 64 bits



Flag	Content
<b>19-bit VPN</b>	The rest reserved for the kernel.
<b>24-bit PFN</b>	<b>Systems can support with up to 64GB of main memory( <math>2^{24} * 4KB</math> pages ).</b>
Global bit(G)	Used for pages that are globally-shared among processes.
<b>ASID</b>	<b>OS can use to distinguish between address spaces.</b>
Coherence bit(C)	determine how a page is cached by the hardware.
Dirty bit(D)	marking when the page has been written.
Valid bit(V)	tells the hardware if there is a valid translation present in the entry.

# Translation Lookaside Buffer (TLB)

- Address translation with TLB & page table



# Translation Lookaside Buffer (TLB)

- Effective Access Time (EAT)
  - effective time for accessing memory with TLB
  - TLB Hit ratio =  $\alpha$ 
    - percentage of times that a page number is found in the TLB
    - TLB hit: one memory access
    - TLB miss: two memory accesses
  - Consider  $\alpha = 80\%$ , 100ns for each memory access
    - $EAT = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
  - Consider a more realistic hit ratio  $\alpha = 99\%$ ; still 100ns for each memory access
    - $EAT = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$

# How can TLB improve performance

- Example: accessing an array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++){
2:                  sum+=a[i];
3:      }
```

# How can TLB improve performance

- Example: accessing an array

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++){
2:                  sum+=a[i];
3:      }
```

**The TLB improves performance  
due to **spatial locality****


3 misses and 7 hits.  
Thus **TLB hit rate** is 70%.

# How can TLB improve performance

- Locality


- Temporal Locality

- An instruction or data item that has been recently accessed will likely be re-accessed soon in the future.



- Spatial Locality

- If a program accesses memory at address x, it will likely soon access memory near x.



# Agenda

- ~~Recap~~
- ~~Translation Lookaside Buffer (TLB)~~
- Multi-Level Page Tables




# Space Overhead of Page Tables

- A (linear) page table can be large
  - Example (revisit):
    - 32-bit address space (4GB) with 4KB pages
    - 12 bits for offset within a page ( $4K=2^{12}$ )
    - 20 bits for VPN ( $32 - 12 = 20$ )
    - $4MB = 2^{20} \text{ entries} * 4 \text{ Bytes per page table entry}$
- Each process needs to have a page table
  - 100 processes needs  $4MB * 100 = 400MB$

# Observation

- Many pages are unused, many PTEs are empty

**A 16KB Address Space  
with 1KB Pages**



PFN	valid	prot	present	dirty
10	1	r-x	1	0
-	0	-	-	-
-	0	-	-	-
-	0	-	-	-
15	1	rw-	1	1
...	...	...	...	...
-	0	-	-	-
3	1	rw-	1	1
23	1	rw-	1	1

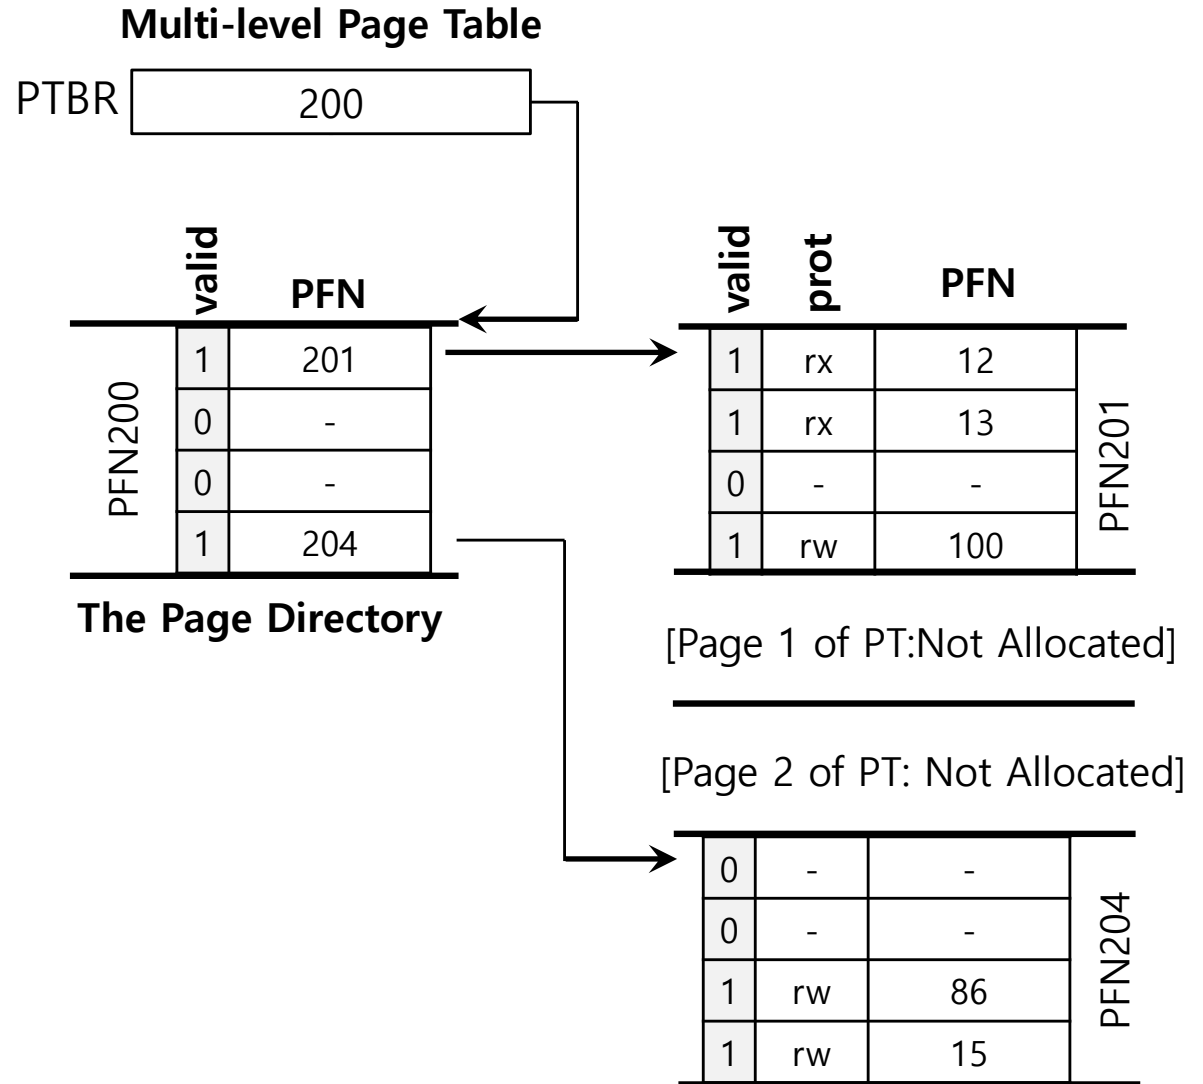

**A Page Table For 16KB Address Space**

Physical Memory

# Multi-Level Page Tables

- Paging the (linear) page table
  - Chop up the page table into page-sized units
  - If an entire page of page-table entries is invalid, don't allocate that page of the page table at all
  - To track whether a page of the page table is valid, use a new structure, called **page directory**
    - It consists of a number of page directory entries (PDE)
      - one PDE per page of the page table
      - PDE has a valid bit and page frame number (PFN)

# Multi-Level Page Tables



# Agenda

- ~~Recap~~

- ~~Translation Lookaside Buffer (TLB)~~

- ~~Multi-Level Page Tables~~

## Questions?



\*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.