

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

Lecture 24: Page Replacement Algorithms

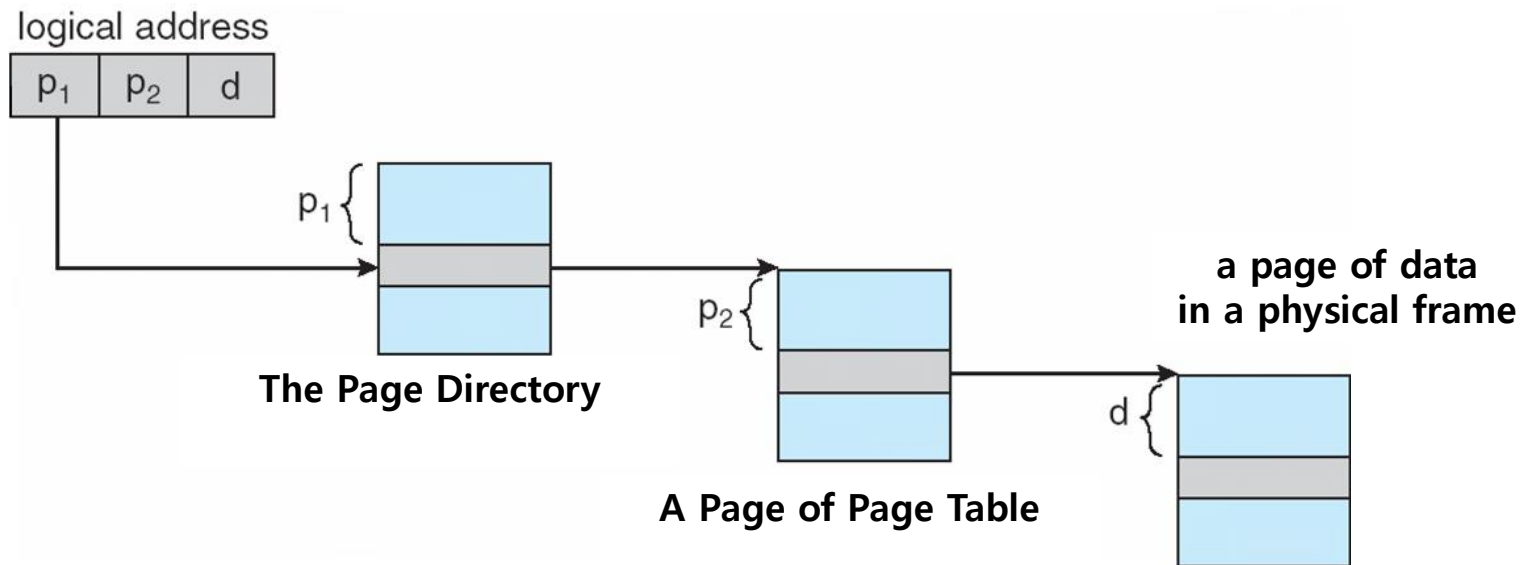


Agenda

- **Recap**
- **Page Replacement Algorithms**
 - **FIFO (First-In-First-Out)**
 - **LRU (Least-Recently-Used)**
 - **Clock**


Recap

- Multi-Level Page Tables
 - Paging the page table itself
 - e.g., a two-level page table
 - the page directory index (p_1) is used to identify a page directory entry (PDE) in the page directory
 - the page table index (p_2) is used to identify a page table entry



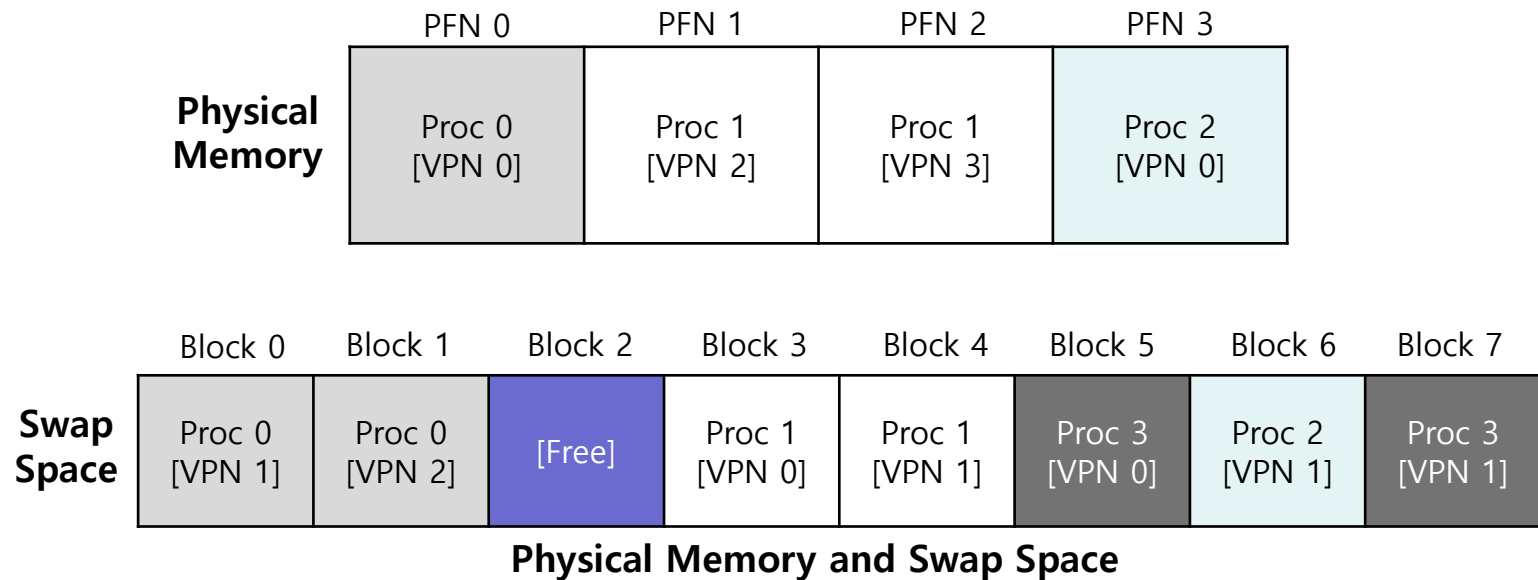
Recap

- Intel IA-32 Paging Architecture



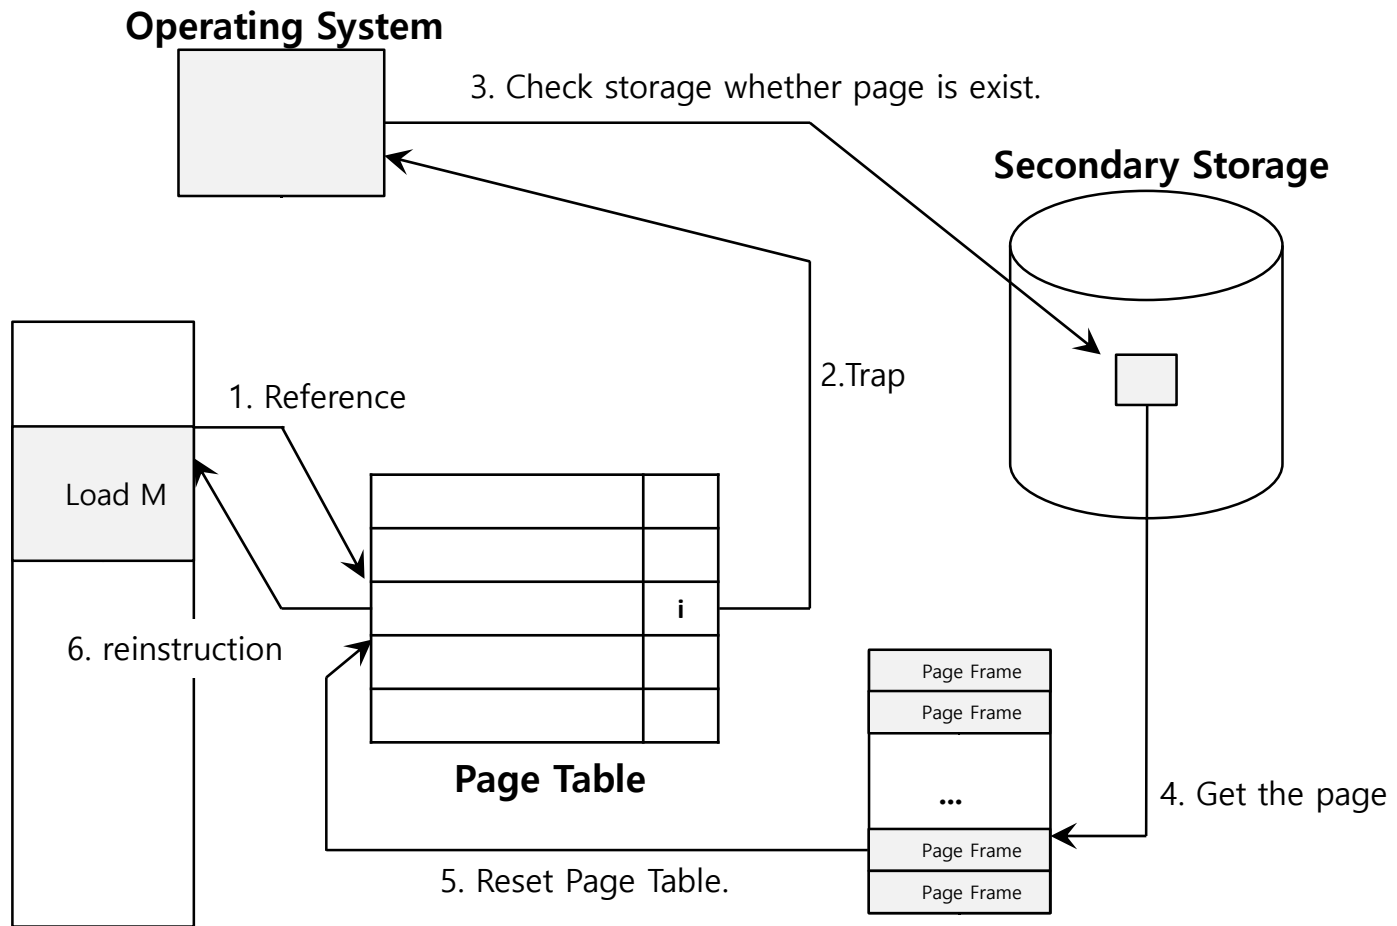
Recap

- Swapping
 - stash away portions of address space that currently aren't in great demand
 - the unpopular pages are placed in the **swap space** at the next layer of memory/storage hierarchy



Recap

- Swap in a page from disk at page fault



Agenda

- ~~Recap~~
- Page Replacement Algorithms
 - FIFO (First-In-First-Out)
 - LRU (Least-Recently-Used)
 - Clock

Page Replacement Algorithms

- When free space in physical memory is low, OS starts evicting pages to disks
- Page replacement algorithms decide which page to evict
 - the physical memory serves as a cache of the swap space
 - a cache miss leads to a page fault
 - the goal is to minimize the number of cache misses/page faults
 - i.e., minimize the accesses to the swap space

The Optimal Algorithm

- Lead to the fewest number of misses overall
 - Replace the page that will be accessed furthest in the future
 - Resulting in the fewest-possible cache misses
 - or the highest **hit rate**
- Need future access information
 - usually unknown in practice
 - serve only as a comparison point, to see how close we are to the (theoretical) best case scenario

The Optimal Algorithm

- Example

References

0 1 2 0 1 3 0 3 1 2 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	2	0,1,3
0	Hit		0,1,3
3	Hit		0,1,3
1	Hit		0,1,3
2	Miss	3	0,1,2
1	Hit		0,1,2

The Optimal Algorithm

- Example

References										
0	1	2	0	1	3	0	3	1	2	1


Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	2	0,1,3
0	Hit		0,1,3
3	Hit		0,1,3
1	Hit		0,1,3
2	Miss	3	0,1,2
1	Hit		0,1,2

Hit rate:

$$\frac{\text{Hits}}{\text{Hits} + \text{Misses}} = 54.6\%$$

The Optimal Algorithm

- Another Example
 - 20 references in sequence
 - 3 page frames
 - # of page fault: 9



First In First Out (FIFO) Algorithm

- Pages were recorded in a queue when they enter the system
 - the page on the head of the queue (the “**First-in**” pages) is evicted first
- simple to implement
- irrelevant to the popularity/importance of pages

First In First Out (FIFO) Algorithm

- Example

Reference Row

0 1 2 0 1 3 0 3 1 2 1

Hit rate:

$$\frac{\text{Hits}}{\text{Hits} + \text{Misses}} = 36.4\%$$


Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	0	1,2,3
0	Miss	1	2,3,0
3	Hit		2,3,0
1	Miss		3,0,1
2	Miss	3	0,1,2
1	Hit		0,1,2

First In First Out (FIFO) Algorithm

- Another Example
 - 20 references in sequence
 - 3 page frames
 - # of page fault: 15

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

Using History

- Lean on the past and use history.
 - Two type of historical information.

Historical Information	Meaning	Algorithms
recency	The more recently a page has been accessed, the more likely it will be accessed again	LRU (least-recently-used)
frequency	If a page has been accessed many times, It should not be replaced as it clearly has some value	LFU (least-frequently-used)

LRU Algorithm

- Replace the least-recently-used page

Reference Row

0 1 2 0 1 3 0 3 1 2 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		1,2,0
1	Hit		2,0,1
3	Miss	2	0,1,3
0	Hit		1,3,0
3	Hit		1,0,3
1	Hit		0,3,1
2	Miss	0	3,1,2
1	Hit		3,2,1

LRU Algorithm

- Another Example
 - 20 references in sequence
 - 3 page frames
 - # of page fault: 12

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		


page frames

LRU Algorithm

- How to determine “least-recently-used”?
 - need to do some extra accounting work on every memory reference
 - e.g., hardware counters
 - may be expensive
- Approximating LRU with simple hardware support
 - use a **reference bit** in page table entry (PTE)
 - “accessed bit”, “use bit”
 - Whenever a page is referenced, the reference bit is set to 1 by the hardware
 - OS may clear it

Clock Algorithm

- All pages arrange in a circular list
 - A clock hand points to some particular page to begin with; the algorithm continues until it finds a reference bit that is 0 (i.e., not used)



reference bit	action
0	Evict the page
1	Clear reference bit and advance hand

Considering Dirty Pages

- Page table entry includes a modified bit (a.k.a dirty bit)
 - Page has been modified (i.e., dirty bit is 1)
 - it must be written back to disk to evict it
 - Page has not been modified (i.e., dirty bit is 0)
 - the eviction is free, i.e., don't need to write to disk

Agenda

- ~~Recap~~

- ~~Page Replacement Algorithms~~

- ~~FIFO (First-In-First-Out)~~

- ~~LRU (Least-Recently-Used)~~

- ~~Clock~~

Questions?



*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.