

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

Lecture 25: Memory APIs




Agenda

- **Recap**
- **Prefetching & Thrashing**
- **Memory APIs**
 - **malloc() / free()**
 - **calloc()**
 - **realloc()**

Recap

- Page Replacement Algorithms
 - The Optimal Algorithm
 - FIFO Algorithm
 - LRU Algorithm
 - Clock Algorithm





Agenda

- ~~Recap~~
- Prefetching & Thrashing
- Memory APIs
 - malloc() / free()
 - calloc()
 - realloc()

Prefetching


- The OS guess that a page is about to be used, and thus bring it in ahead of time.



Page 2 likely **soon to be accessed**, so bring it into memory together with (requested) page 1

Thrashing

- The OS keeps swapping pages in and out for processes
 - Usually happen when the memory is **oversubscribed**
 - the memory demands of the set of running processes **exceeds** the available physical memory significantly
 - Low CPU utilization



Agenda

- ~~Recap~~
- ~~Prefetching & Thrashing~~
- Memory APIs
 - malloc() / free()
 - calloc()
 - realloc()

Memory API: malloc()

```
#include <stdlib.h>

void* malloc(size_t size)
```

- Allocate a memory region on the heap
 - Argument
 - `size_t size` : size of the memory block(in bytes)
 - `size_t` is an unsigned integer type.
 - Return
 - Success : a void type pointer to the memory block allocated by `malloc`
 - Fail : a null pointer

Memory API: malloc()

- Instead of typing in a number directly for `size` in `malloc`, use `sizeof` to ensure the requested number of bytes is accurate
- e.g., `sizeof()`

```
int *x = malloc(10 * sizeof(int));  
printf("%d\n", sizeof(x));
```

4

```
int x[10];  
printf("%d\n", sizeof(x));
```

40


Memory API: free()

```
#include <stdlib.h>

void free(void* ptr)
```


- Free a memory region allocated by a call to `malloc`.
 - Argument
 - `void *ptr` : a pointer to a memory block allocated with `malloc`
 - Return
 - none

Memory Allocation




```
int *pi; // local variable
```

Memory Allocation




```
int *pi; // local variable
```




-----> pointer

```
pi = (int *)malloc(sizeof(int) * 4);
```

Memory Freeing

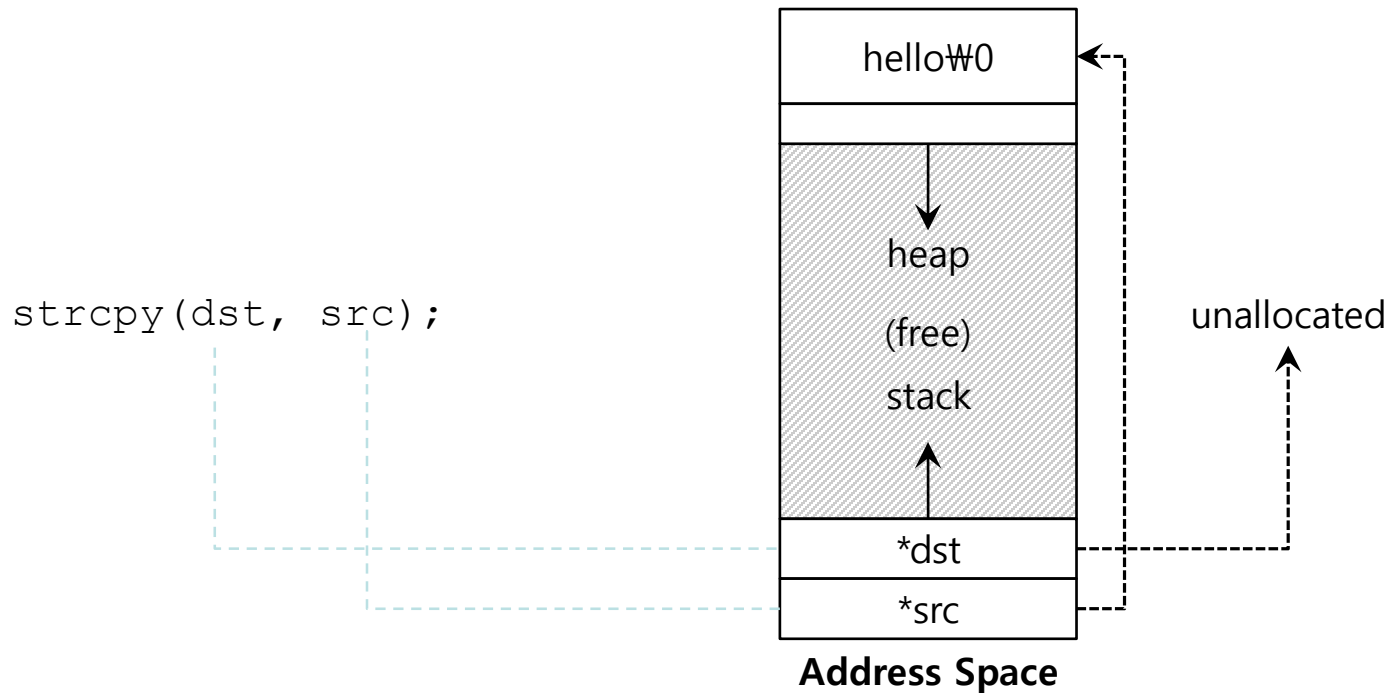


```
free(pi);
```




Forget To Allocate Memory

```
char *src = "hello"; //character string constant  
char *dst;           //unallocated  
strcpy(dst, src);    //segfault and die
```




Forget To Allocate Memory (cont')

```
char *src = "hello";    //character string constant
char *dst (char *)malloc(strlen(src) + 1 ); // allocated
strcpy(dst, src);       //work properly
```




Forget To Initialize Memory

```
int *x = (int *)malloc(sizeof(int)); // allocated
printf("*x = %d\n", *x); // uninitialized memory access
```



Address Space




Address Space


Memory Leak

- Keep allocating memory but forget to deallocate (free)


unused : unused, but not freed



Address Space






Address Space



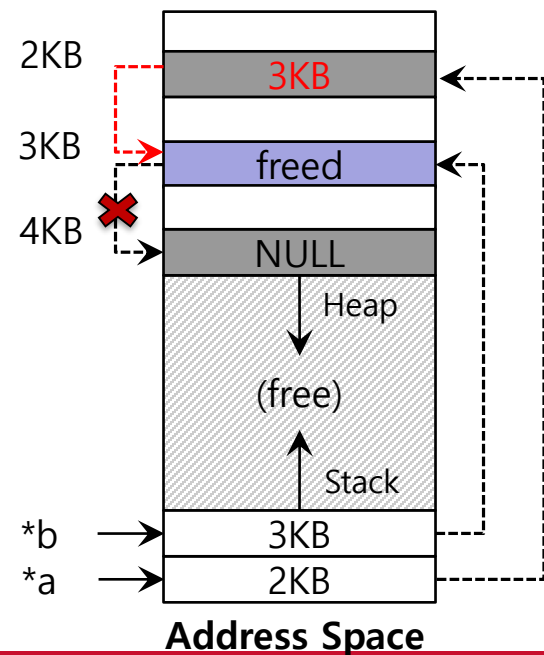
Address Space

Dangling Pointer

- Free memory before it is finished using
 - a process accesses to memory with an invalid pointer




`free(b)`



Double Free

- Free memory that had been freed already

```
int *x = (int *)malloc(sizeof(int)); // allocated
free(x); // free memory
free(x); // free repeatedly
```



Other Memory APIs: calloc()

```
#include <stdlib.h>

void *calloc(size_t num, size_t size)
```

- Allocate memory on the heap and zeroes it before returning.
 - Argument
 - `size_t num` : number of blocks to allocate
 - `size_t size` : size of each block(in bytes)
 - Return
 - Success : a void type pointer to the memory block allocated by `calloc`
 - Fail : a null pointer

Other Memory APIs: realloc()

```
#include <stdlib.h>

void *realloc(void *ptr, size_t size)
```

- Change the size of memory block.
 - A pointer returned by `realloc` may be either the same as `ptr` or a new.
 - Argument
 - `void *ptr`: Pointer to memory block allocated with `malloc`, `calloc` or `realloc`
 - `size_t size`: New size for the memory block(in bytes)
 - Return
 - Success: Void type pointer to the memory block
 - Fail : Null pointer

Related System Calls

- Internally, `malloc` library call invokes `brk/sbrk` system calls
 - expand the program's *break*.
 - *break*: The location of **the end of the heap** in address space
 - programmers should never directly call either `brk` or `sbrk`.

```
#include <unistd.h>

int brk(void *addr)
void *sbrk(intptr_t increment);
```

Agenda

- ~~Recap~~
- ~~Prefetching & Thrashing~~
- ~~Memory APIs~~
 - ~~malloc() / free()~~
 - ~~calloc()~~
 - ~~realloc()~~

Questions?



*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.