

IOWA STATE UNIVERSITY

Department of Electrical and Computer Engineering

Lecture 32: File System Implementation II

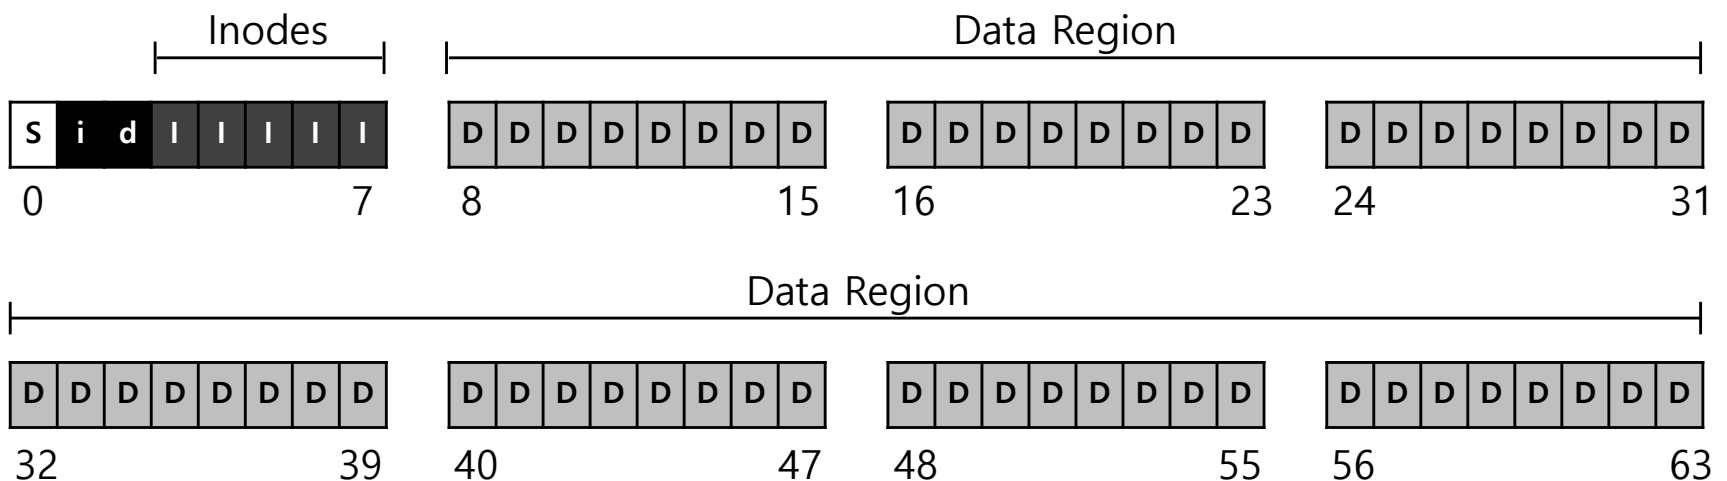


Agenda

- **Recap**
- **File System Implementation II**
 - **Access Paths**
 - **Caching & Buffering**
 - **Fast File System (FFS)**

Recap

- Two important perspectives of an FS
 - Data structures
 - Access methods
- Basic design
 - data region: user data
 - metadata region: inodes, bitmaps, superblock



Recap

- Each user file is represented by one `inode`
 - `inode` contains all information about the file
 - e.g., type, size, permissions, timestamps, pointers to data blocks
 - data blocks may be indexed by multi-level indirect pointers
- Each `inode` is referred to by an inode number
 - FS calculates where the inode is on the disk based on the inode number

The Inode table

				iblock 0				iblock 1				iblock 2				iblock 3				iblock 4								
Super	i-bmap				d-bmap				0	1	2	3	16	17	18	19	32	33	34	35	48	49	50	51	64	65	66	67
									4	5	6	7	20	21	22	23	36	37	38	39	52	53	54	55	68	69	70	71
									8	9	10	11	24	25	26	27	40	41	42	43	56	57	58	59	72	73	74	75
									12	13	14	15	28	29	30	31	44	45	46	47	60	61	62	63	76	77	78	79
0KB	4KB	8KB	12KB	16KB				20KB				24KB				28KB				32KB								

Agenda

- ~~Recap~~
- **File System Implementation II**
 - **Access Paths**
 - **Caching & Buffering**
 - **Fast File System (FFS)**

Access Paths

- Reading a file from disk
 - Issue `open ("/foo/bar", O_RDONLY)`
 - Traverse the pathname and locate the desired inode
 - Begin at the root of the file system (/)
 - In many Unix file systems, the root inode number is 2
 - based on the inode number, FS finds the block in the inode table that contains inode #2
 - Read inode #2 to find pointers to its data blocks (contents of the root directory)
 - read the data blocks and find an entry containing "foo" and its inode#
 - Traverse recursively the pathname until finding the inode for "bar"
 - Check permissions, allocate a file descriptor for this process and returns file descriptor to user.

Access Paths

- Reading a file from disk (cont')
 - Issue `read()` to read from the file
 - Already found the inode for the open file (via `open`)
 - read the inode to find the location of its data blocks
 - read in the first block of the file
 - Update the inode with a new last accessed time
 - Update the file offset in the open file table for the process
 - When the file is closed:
 - File descriptor should be deallocated
 - that is all the FS really needs to do; no disk I/Os take place

Access Paths

- Reading a file from disk (cont')
 - Timeline (Time Increasing downward)

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
open(bar)			read	read	read	read	read			
read()					read			read		
read()					read				read	
read()					read					read

Access Paths

- Writing to a file on disk
 - Issue `write()` to update the file with new contents.
 - may allocate a new data block (unless a block is being overwritten).
 - Need to update data block, data bitmap
 - generates five I/Os (at least):
 - one to read the data bitmap
 - one to write the bitmap (to reflect its new state to disk)
 - two more to read and then write the inode
 - one to write the actual block itself
 - To create file, FS also needs to update the directory, causing additional I/O traffic

Access Paths

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
create (/foo/bar)		read write	read	read write	read write	read	read write			
write()	read write				read write			write		
write()	read write				read write				write	
write()	read write				read write					write

Agenda

- ~~Recap~~
- ~~File System Implementation II~~
 - ~~Access Paths~~
 - Caching & Buffering
 - Fast File System (FFS)

Caching and Buffering

- Reading and writing files are expensive, incurring many I/Os
 - E.g., long pathname(/1/2/3/.../100/file.txt)
 - One to read the inode of the directory and at least one read its data
 - Literally perform hundreds of reads just to open the file
- To reduce I/O traffic, FSes use system memory (DRAM) to cache reads
 - Read I/O can be avoided by large cache
 - **page cache** in Linux

Caching and Buffering


- Write traffic has to go to disk for persistence
 - caching is less useful for writes
- But buffering writes in memory may still help improve performance
 - FS can optimize the writes in memory, e.g.:
 - batch some updates into a smaller set of I/Os
 - avoiding unnecessary I/O (e.g., overwritten in memory)
- Applications may force flush dirty data to disk by calling `fsync()`

Agenda

- ~~Recap~~
- File System Implementation II
 - ~~Access Paths~~
 - ~~Caching & Buffering~~
 - Fast File System (FFS)

The 1st Unix File System (~1974)

- Similar to the basic FS we discussed
 - a simplified view



- The Good Thing
 - Simple and supports the basic abstractions
 - Easy to implement
- The Problem
 - Terrible performance

The 1st Unix File System (~1974)

- Why the performance is terrible
 - major issue: treated the disk as a **random-access memory**
 - Example of random-access blocks with four files.
 - Data blocks for each file can be accessed by going back and forth the disk, because they are **contiguous**.

A1	A2	B1	B2	C1	C2	D1	D2
----	----	----	----	----	----	----	----

- File b and d is deleted.


A1	A2			C1	C2		
----	----	--	--	----	----	--	--

- File E is created with free blocks. (**spread across** the disk!)

A1	A2	E1	E2	C1	C2	E3	E4
----	----	----	----	----	----	----	----

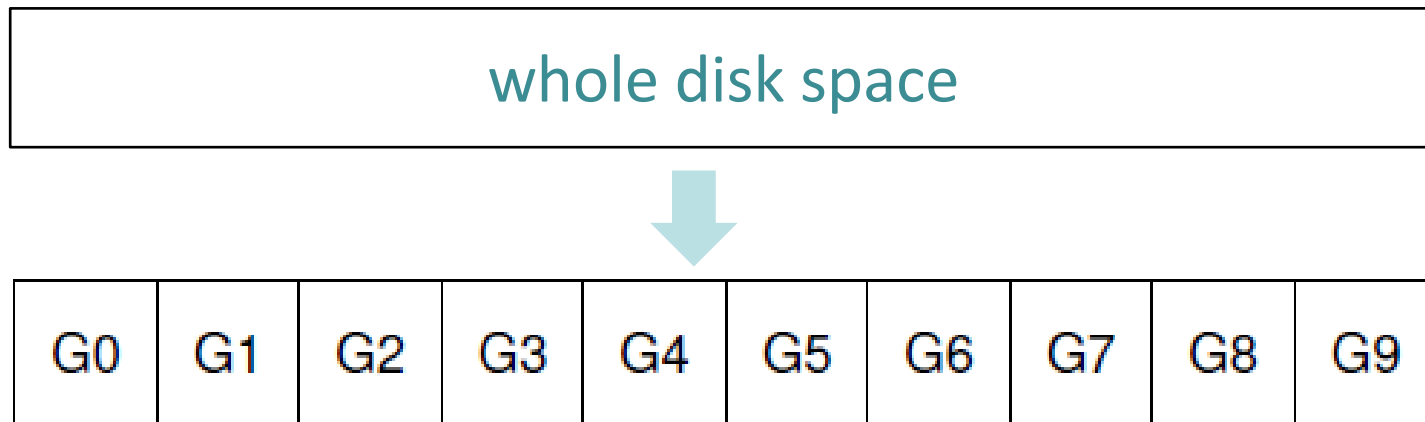
The 1st Unix File System (~1974)

- Why the performance is terrible
 - major issue: treated the disk as a **random-access memory**
 - The actual disk is different from DRAM (recap)
 - access latency varies much at different locations
 - e.g., seek time varies




The Fast File System (FFS, ~1984)

- Key insight: disk awareness
 - FS structures and allocation policies match the internals of disks
 - Divide the disk into cylinder groups (block groups)
 - place related stuff in the same group, avoid long seek



The Fast File System (FFS, ~1984)

- Each cylinder group has a structure similar to the Unix FS



The Fast File System (FFS, ~1984)

- Heuristics to Allocate Related Stuff Together
 - For directory
 - find the cylinder group with:
 - a low number of allocated directories (b/c we want to balance directories across groups)
 - a high number of free inodes (b/c we want to be able to allocate a bunch of files later)
 - put the directory data and inode in that same group
 - For file
 - allocate the data and inode of the file in the same group
 - places all files that are in the same directory in the cylinder group of the directory they are in

The Fast File System (FFS, ~1984)

- Much Better Performance
 - e.g., 14-47% of raw disk bandwidth
- Main Lesson of FFS
 - treat disk like it's a disk
- A watershed moment in FS history
- Many FSes today take cues from FFS

Agenda

- ~~Recap~~
- ~~File System Implementation II~~
 - ~~Access Paths~~
 - ~~Caching & Buffering~~
 - ~~Fast File System (FFS)~~

Questions?



*acknowledgement: slides include content from “Modern Operating Systems” by A. Tanenbaum, “Operating Systems Concepts” by A. Silberschatz etc., “Operating Systems: Three Easy Pieces” by R. Arpaci-Dusseau etc., and anonymous pictures from internet.