

# Relational Database

# The dominate data model

- IBM's DB2 family
- Oracle
- Microsoft's Access and SQLServer
- A multi-billion dollar industry

# Relational Database

- **Relational database**: a set of **relations**
- **Relation**: made up of 2 parts:
  - **Schema** : specifies the name of relations, plus name and type of each column.
    - E.G. Students (sid: string, name: string, login: string, age: integer, gpa: real).
  - **Instance** : a *table*, with rows and columns
- Can think of a relation as a **set** of **rows** or **tuples or records** (i.e., all rows are distinct). \*
- \*commercial systems allow duplicate rows, but

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Cardinality  
(#rows)= 3  
degree  
(#columns)= 5  
all rows are  
distinct

# Does the order Matter?

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

sid	name	login	gpa	age
53666	Jones	jones@cs	3.4	18
53650	Smith	smith@math	3.8	19
53688	Smith	smith@eecs	3.2	18

- Record (row): no
  - Why: set
- Columns (field): depends

# Integrity Constraints (ICs)

- **IC:** condition that must be true for *any* instance of the database
  - ICs are specified when schemas are defined
  - ICs are checked when relations are modified
- A *legal* instance of a relation is one that satisfies all specified ICs
  - DBMS should not allow illegal instances
- If the DBMS checks ICs, stored data is more faithful to real-world meaning
  - Avoids data entry errors, too

## Two Common Types of ICs

- Primary Key Constraints (key)
- Referential Integrity Constraint (Foreign key)

# Primary Key

- **Key**: *minimal* set of the fields of a relation that can uniquely identify a tuple
  - {SSN} is a key
  - If there's >1 key for a relation, one of the keys is chosen (by DBA) to be the **primary key**. The other keys are called **candidate keys**.
- **Superkey**: set of the fields of a relation that can uniquely identify a tuple
  - E.g., {SSN, Name, Age} is a superkey.
  - A key must be a super key, but not vice versa

# Excise


- T/F: every relation is guaranteed to have a key
  - Yes. Because relation is a set of tuples, the set of all fields is always a superkey
- Facebook. What is/are the key/keys?

# Foreign Key

**Foreign key (FK)** : Set of fields in one relation that is used to 'refer' to a tuple in another relation. (Must correspond to primary key of the second relation.)

- FK is Like a 'logical pointer'.

- Referential Integrity



Employee: referencing relation  
Department: referenced relation

- DNO is a foreign key of EMPLOYEE.



# The SQL Query Language

- Data Definition Language (DDL)
  - enables creation, deletion, and modification of definitions for tables and views and integrity constraint specification.
- Data Manipulation Language (DML)
  - allows users to query, to insert, or to delete rows.

# Creating Relations in SQL

- Creates the Students relation. Observe that the type (domain) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.
- As another example, the Enrolled table holds information about courses that students take.

```
CREATE TABLE Students
(sid CHAR(20),
 name CHAR(20),
 login CHAR(10),
 age INTEGER,
 gpa REAL
primary key (sid))
```

```
CREATE TABLE Enrolled
(sid CHAR(20),
 cid CHAR(20),
 grade CHAR(2))
```

# Primary and Candidate Keys in SQL

- Possibly many candidate keys (specified using **UNIQUE**), one of which is chosen as the *primary key*.
  - ❖ “A student can take a same course only once.”

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid) )
```
  - ❖ “Students can take only one course, and no two students in a course receive the same grade.”

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid),
UNIQUE (cid, grade) )
```

# Foreign Keys in SQL

- Only students listed in the Students relation should be allowed to enroll for courses.

```
CREATE TABLE Enrolled  
  (sid CHAR(20), cid CHAR(20), grade CHAR(2),  
   PRIMARY KEY (sid,cid),  
   FOREIGN KEY (sid) REFERENCES Students )
```

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Enrolled

sid	cid	grade
53666	Carnatic101	C
53666	Reggae203	B
53650	Topology112	A
53666	History105	B

# Enforcing Referential Integrity

- Consider Students and Enrolled; *sid* in Enrolled is a foreign key that references Students.
- What should be done if an Enrolled tuple with a non-existent student id is inserted? (*Reject it!*)
- What should be done if a Students tuple is deleted?
  - Not allowed if it appears in Enrolled
  - Delete all Enrolled tuples that refer to it.
  - Set *sid* in Enrolled tuples that refer to it to a *default sid*.
  - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting '*unknown*' or '*inapplicable*'.)

# Referential Integrity in SQL/92

- SQL/92 supports all 4 options on deletes and updates.
  - Default is **NO ACTION** (*delete/update is rejected*)
  - **CASCADE** (also delete all tuples that refer to the deleted tuple)
  - **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE)
```

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE SET NULL)
```

# Destroying and Altering Relations

`DROP TABLE` Students

- Destroys the relation Students. The schema information *and* the tuples are deleted.

`ALTER TABLE` Students  
`ADD` firstYear integer

- The schema of Students is altered by adding a new field; every tuple in the current instance is extended with a *NULL* value in the new field.

`ALTER TABLE` Students  
`drop column` firstYear

# Adding and Deleting Tuples

- Insert a single tuple:

```
INSERT INTO Students (sid, name, login, age, gpa)
VALUES ('53688', 'Smith', 'smith@ee', 18, 3.2)
```


- Can delete all tuples satisfying some conditions (e.g., name = Smith):

```
DELETE
FROM Students S
WHERE S.name = 'Smith'
```



# Logical DB Design: ER to Relational

- Entity sets to tables.




```
CREATE TABLE Employees  
(ssn CHAR(11),  
name CHAR(20),  
lot INTEGER,  
PRIMARY KEY (ssn))
```

For each regular entity set  $E$  in the ER diagram, create a relation  $R$  that includes all the simple attributes of  $E$ . Include only the simple component attributes of a composite attribute.

Choose one of the key attributes of  $E$  as the primary key of  $R$ .

If the chosen key is composite, the simple attributes that form the key are taken together as the primary key of  $R$ .


# Relationship Sets to Tables



- In translating a relationship set to a relation, attributes of the relation must include:
  - Keys for each participating entity set (as foreign keys).
    - This set of attributes forms a *key* for the relation.
  - All descriptive attributes.

```
CREATE TABLE Works_In(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (ssn, did),  
  FOREIGN KEY (ssn)  
    REFERENCES Employees,  
  FOREIGN KEY (did)  
    REFERENCES Departments)
```

# Uni-participation Constraints




- Map relationship to a table:
  - Note that **did** is the key now!
  - Separate tables for Employees and Departments.


```
CREATE TABLE Manages(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  FOREIGN KEY (did) REFERENCES Departments)
```

Since each department has a unique manager, we could instead combine Manages and Departments.

```
CREATE TABLE Dept (  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  mgrssn CHAR(11),  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees)
```



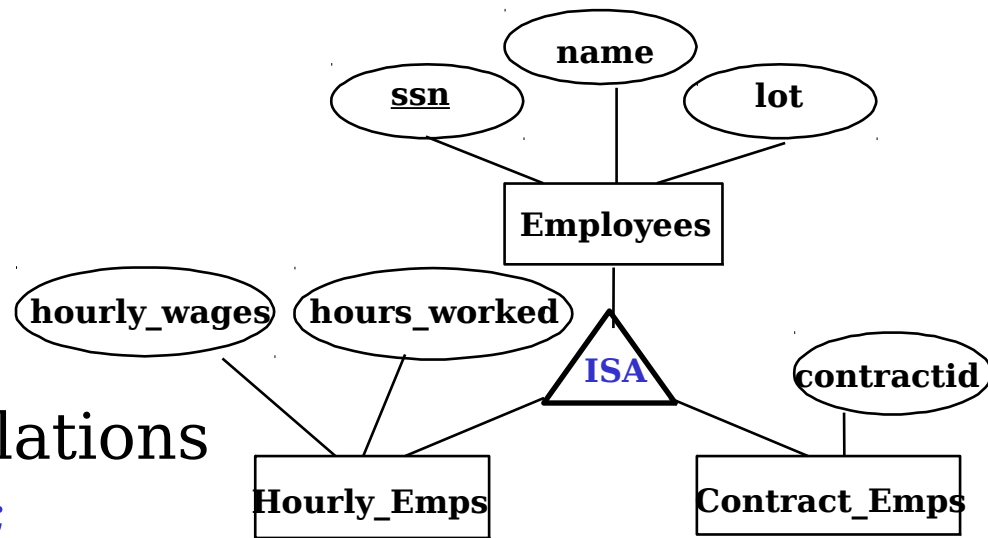
# Total and Uni-Participation Constraints



```
CREATE TABLE Dept (  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  mgrssn CHAR(11), NOT NULL  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees)
```

By specifying “NOT NULL”, mgrssn must have a value, thus satisfying total participation constraint

# Translating ISA Hierarchies to Relations



- General approach: 3 relations
  - `Employee(ssn, name, lot);`
  - `Hourly_Emps (ssn, hourly_wages, hours_worked)`
  - `Contract_Emps(ssn, contractid)`
    - Must delete Hourly\_Emps tuple if referenced Employees tuple is deleted
    - Queries involving all employees are easy
    - Queries involving just Hourly\_Emps require a join to get some attributes
- Alternative: Just Hourly\_Emps and Contract\_Emps
  - `Hourly_Emps (ssn, name, lot, hourly_wages, hours_worked)`
  - `Contract_Emps (ssn, name, lot, contractid)`
    - Each employee must be in one of these two

User database  
Requirements in  
natural languages




ER-diagrams

- Entity sets
- Relationship sets
- Constrains (uni and total)

Database

- relations
- constrains (key, foreign key)

1. Which one is easier, less likely to make mistake?
  - From requirements to ER
  - From ER to database
2. Why and what can we do?



1. There are a number of Employees, each of which has a unique SSN, a name, a dob
2. There are a number of Departments, each of which has a unique did, a name, an address
3. Some employees works in some departments
  - 1) An employee can work in at most one department
  - 2) Every department must have one



User database  
Requirements in  
natural languages

Convert the  
requirements in  
COMS 561  
language

ER-diagrams

- Entity sets
- Relationship sets
- Constrains (uni and total)

Database

- relations
- constrains (key, foreign key)

## Entity Sets

- There are a number of  $E_1$ , each of which has  $a_1, a_2, \dots, a_{n1}$ , where  $a_x$  is unique,  $a_y$  is unique, and  $\dots$ , so on
- There are a number of  $E_2$ , each of which has  $a_1, a_2, \dots, a_{n2}$ , where  $attr1_x$  is unique
- .....

## Relationship Sets

- $E_x$  has to do with  $E_y$
- .....

## Constraints

- Each one in  $E_x$  can do with someone in  $E_y$ , at most one time, and/or at least one time
- .....