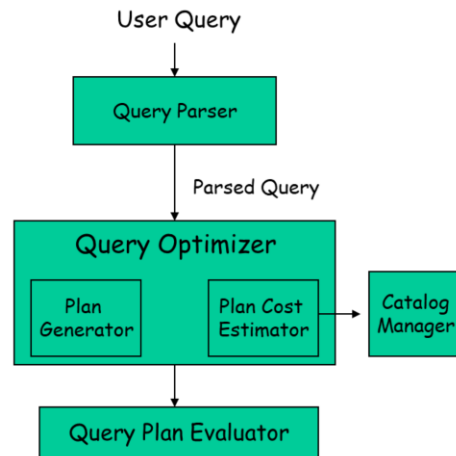


Overview of Query Evaluation

- A user query is expressed using SQL or other languages, ideally natural languages
- A parsed query is essentially treated as a **relational algebra** expression
 - Selection (σ)
 - Join (\bowtie)
 - Project(π)
 - Union, intersection and difference, cross product
- Query optimizer
 - Enumerates the possible plans to evaluate expression,
 - Select a small subset of these plans and estimate their cost



This lecture will introduce Relational Algebra

Relational Algebra

- What is it?
- Why we need it?

Think about algebra.

For example $(a^2+b)*c/d$.

Algebra has symbols defining for different operations.

These operations has different level of priority. So $(a^2+b)*c/d$ is different from a^2+b*c/d

We can rewrite the same formula for faster calculation. $(a^2+b)*c/d = a^2*c/d+b*c/d$, but the left hand side have less operations, so faster to calculate

Union, Intersection and Difference Operators

- These are defined if r and s have the same schema
 - $r \cup s = \{x: x \in r \vee x \in s\}$; its schema is same as that of r or s
 - $r - s = \{x: x \in r \wedge x \notin s\}$; its schema is same as that of r or s
 - $r \cap s = \{x: x \in r \wedge x \in s\}$; its schema is same as that of r or s
- Example

| r | | s | | $r \cup s$ | | $r - s$ | | $r \cap s$ | |
|---|---|---|---|------------|---|---------|---|------------|---|
| A | B | A | B | A | B | A | B | A | B |
| 1 | 2 | 1 | 3 | 1 | 2 | 1 | 2 | 1 | 2 |
| 3 | 4 | 3 | 4 | 3 | 4 | | | 3 | 4 |
| 5 | 6 | 5 | 6 | 5 | 6 | | | 5 | 6 |
| 7 | 8 | | | 7 | 8 | 7 | 8 | | |
| | | | | 1 | 3 | | | | |

Start with the basic

Relations are sets of tuples

So all set operators can be applied here (but there is a requirement: the schema of the two relations has to be the same)


Union \cup

Intersection \cap

Difference $-$

Question

- Is $r \cap s = s \cap r$
- Is $r \cup s = s \cup r$
- Is $r - s = s - r$



Select Operator

- $\sigma_c(r)$ selects those tuples of r that satisfy condition c
- Examples of conditions
 - $A \leq B$,
 - $A > 2 \wedge A < B$
- $\sigma_c(r) = \{x: x \in r \wedge x \text{ satisfies condition } c\}$
- Example

r

| A | B |
|---|---|
| 1 | 2 |
| 3 | 4 |
| 5 | 6 |
| 7 | 6 |

$\sigma_{A > 2 \wedge A < B}(r)$

| A | B |
|---|---|
| 3 | 4 |
| 5 | 6 |

Sigma σ

Question

- Do you remember what we learned in the Chapter on Select?

Project Operator

- $\Pi_X(r)$ retains only X columns
 - In $\Pi_X(r)$, X must consist of some attributes of r
- The schema of $\Pi_X(r)$ is X
- $\Pi_X(r) = \{x[X] : x \in r\}$
 - Here $x[X]$ includes only attributes in X, other attributes are excluded
- Example

| r | | | $\Pi_{AC}(r)$ | |
|---|---|---|---------------|---|
| A | B | C | A | C |
| 1 | 2 | 4 | 1 | 4 |
| 5 | 4 | 2 | 5 | 2 |
| 5 | 6 | 2 | 7 | 3 |
| 7 | 8 | 3 | | |

- $\Pi_{AC}(r)$ is a set from which a “duplicate tuple was removed”

Pi π

Row VS Column

- In relational algebra, “Selection” is on rows (i.e., which rows satisfy the conditions), while “Projection” is on columns.
- Don’t get confused with the Select statement in SQL.
 - Select name from Students where id=“001”
 - This sql is $\pi_{name}(\sigma_{id="001"}(Students))$

$\pi_{name}(\sigma_{id="001"}(Students))$ chooses the row first and then choose the column.

Can we switch these two operations?

Natural Join Operator

- $r \bowtie s$ combines tuples of r and s agreeing on common attributes
- $r \bowtie s$ is always defined
- Suppose schema of r is R and schema of s is S , Then schema of $r \bowtie s$ is $R \cup S$
- $r \bowtie s = \{x: x[R] \in r \wedge x[S] \in s\}$
 - For a tuple to be in $r \bowtie s$, its R portion should be in r and S portion in S
 - Implicitly, R and S portions match on the common attributes $R \cap S$

• Example

| r | | s | | r \bowtie s | | |
|----|---|---|---|---------------|---|---|
| A | B | B | C | A | B | C |
| 1 | 2 | 2 | 3 | 1 | 2 | 3 |
| 3 | 4 | 4 | 4 | 3 | 4 | 4 |
| 5 | 6 | 4 | 6 | 3 | 4 | 6 |
| 9 | 4 | 9 | 4 | 9 | 4 | 4 |
| 10 | 5 | | | 9 | 4 | 6 |

Join \bowtie

Question

- Do you remember what we learned in Chapter on Join?

Renaming Operator

- Sometimes it becomes necessary to rename attributes or relations before or after an operator is applied
- Renaming does not alter the information content, i.e. tuples
- We use the syntactic form $\alpha \rightarrow \beta$ to express that the identifier α is changed to β .
- Example
 - Consider $r \cup (s: C \rightarrow B): \rightarrow q$
 - Here, attribute C in s is renamed to B
 - After computing the union the result is named as q
 - “.” has the lowest priority

r

| A | B |
|---|---|
| 1 | 2 |
| 3 | 4 |
| 5 | 6 |

s

| A | C |
|---|---|
| 1 | 3 |
| 3 | 4 |
| 5 | 6 |

q

| A | B |
|---|---|
| 1 | 2 |
| 3 | 4 |
| 5 | 6 |
| 1 | 3 |

Cross Product Operator

- $r \times s$ concatenates tuples of r and s
- Every common attribute in r and s renamed
- In renaming prefix “ $r.$ ” and “ $s.$ ” are used to make attributes unique across r and s
- $r \times s = \{x \circ y : x \in \text{renamed } r \wedge y \in \text{renamed } s\}$
 - Here, \circ denotes concatenation of two tuples to form a larger tuple
 - The schema of $r \times s$ consists of all attributes of r and s after renaming
- Example

| r | | s | | r × s | | | |
|---|---|---|---|-------|-----|-----|---|
| A | B | B | C | A | r.B | s.B | C |
| 1 | 2 | 4 | 2 | 1 | 2 | 4 | 2 |
| 3 | 4 | 7 | 3 | 1 | 2 | 7 | 3 |
| 5 | 6 | | | 3 | 4 | 4 | 2 |
| | | | | 3 | 4 | 7 | 3 |
| | | | | 5 | 6 | 4 | 2 |
| | | | | 5 | 6 | 7 | 3 |

Cross product (aka Cartesian product)

Additional Join Operators

- $r \bowtie_c s = \{x \bowtie y : x \in \text{renamed } r \wedge y \in \text{renamed } s \wedge x \bowtie y \text{ satisfies } c\}$
 - The schema of $r \bowtie_c s$ consists of all attributes of r and s after renaming
- Example

| A | B |
|---|---|
| 1 | 2 |
| 3 | 4 |
| 5 | 6 |

| B | C |
|---|---|
| 4 | 2 |
| 4 | 3 |

| A | r.B | s.B | C |
|---|-----|-----|---|
| 1 | 2 | 4 | 2 |
| 1 | 2 | 4 | 3 |
| 5 | 6 | 4 | 2 |
| 5 | 6 | 4 | 3 |

Question

- Can you link natural join with cross product?
- $s \times r$
- $s \bowtie r$

Relational Algebra as a Query Language

Emp

| Name | DName | Salary |
|------|-------|--------|
| John | Toys | 50K |
| Mary | Toys | 60K |
| Leu | Shoes | 55K |

Dept

| DName | MName |
|-------|-------|
| Toys | Mary |

- Give names and salaries of employees in Toys or Credit.

$\Pi_{\text{Name, Salary}}(\sigma_{\text{DName} = \text{'Toys'} \vee \text{DName} = \text{'Credit'}}(\text{Emp}))$

| Name | Salary |
|------|--------|
| John | 50K |
| Mary | 60K |

- Another algebraic expression for the query is as follows:

$\Pi_{\text{Name, Salary}}(\sigma_{\text{DName} = \text{'Toys'}}(\text{Emp}))$
 $\cup \Pi_{\text{Name, Salary}}(\sigma_{\text{DName} = \text{'Credit'}}(\text{Emp}))$

- Different users may express the query in different ways

Relational Algebra as a Query Language

Emp

| Name | DName | Salary |
|------|-------|--------|
| John | Toys | 50K |
| Mary | Toys | 60K |
| Leu | Shoes | 55K |

Dept

| DName | MName |
|-------|-------|
| Toys | Mary |

- Who is John's manager?

$\Pi_{MName}(\sigma_{emp.DName=Dept.DName \wedge Name='John'}(Emp \times Dept))$

- Observe how operators are composed
- Retrieves the following relation

| MName |
|-------|
| Mary |

Quiz on Canvas

- Access code: sailors
- 2 attempts allowed
- Open book, open notes, no communication

Exercise

Sailors(sid, sname, rating, age)

Boats(bid, bname, color)

Reserve(sid, bid, day)

- Find the names of sailors who have reserved boat 103

$$\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$$

Exercise

Sailors(sid, sname, rating, age)

Boats(bid, bname, color)

Reserve(sid, bid, day)

- Find the names of sailors who have reserved a red boat

$$\pi_{sname} \left(\left(\sigma_{color="red"}(Reserve \bowtie Boats) \right) \bowtie sailors \right)$$

Exercise

Sailors(sid, sname, rating, age)

Boats(bid, bname, color)

Reserve(sid, bid, day)

- Find the colors of boats reserved by Lubber

$$\pi_{color} \left(\left(\sigma_{sname="Lubber"}(Reserve \bowtie Sailors) \right) \bowtie Boats \right)$$

Sailors(sid, sname, rating, age)
Boats(bid, bname, color)
Reserve(sid, bid, day)

- Find the names of sailors who have reserved a red or a green boat

$$\begin{aligned} & \pi_{sname} \left(\left(\sigma_{color="red" \vee color="green"} (Reserve \bowtie Boats) \right) \bowtie sailors \right) \\ & \pi_{sname} \left(\left(\sigma_{color="red"} (Reserve \bowtie Boats) \right) \bowtie sailors \right) \\ & \cup \pi_{sname} \left(\left(\sigma_{color="green"} (Reserve \bowtie Boats) \right) \bowtie sailors \right) \end{aligned}$$

Sailors(sid, sname, rating, age)
 Boats(bid, bname, color)
 Reserve(sid, bid, day)

- Find the names of sailors who have reserved a red and a green boat

$\pi_{sname} \left(\left(\sigma_{color="red" \wedge color="green"} (Reserve \bowtie Boats) \right) \bowtie sailors \right)$ is wrong.

Because this means select the boats that are both red and green.

$\pi_{sname} \left(\left(\sigma_{color="red"} (Reserve \bowtie Boats) \right) \bowtie sailors \right) \cap$

$\pi_{sname} \left(\left(\sigma_{color="red"} (Reserve \bowtie Boats) \right) \bowtie sailors \right)$ can also run into problem, if two sailors have the same name.

It's better to use sid

$\pi_{sname} \left(\pi_{sid} \left(\left(\sigma_{color="red"} (Reserve \bowtie Boats) \right) \bowtie sailors \right) \right)$

Sailors(sid, sname, rating, age)
Boats(bid, bname, color)
Reserve(sid, bid, day)

- Find the names of sailors who have reserved at least one boats

Sailors(sid, sname, rating, age)
Boats(bid, bname, color)
Reserve(sid, bid, day)

- Find the sids of sailors with age over 20 who have not reserved a red boat