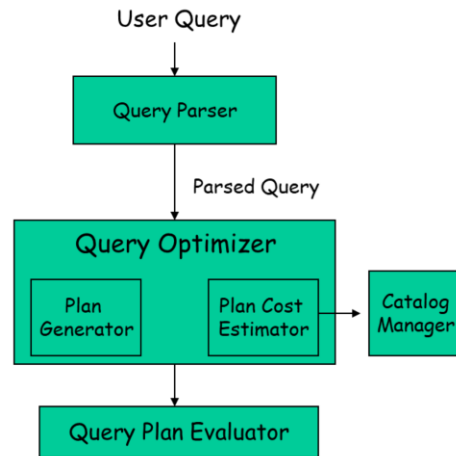


Overview of Query Evaluation

- A user query is expressed using SQL or other languages, ideally natural languages
- A parsed query is essentially treated as a **relational algebra** expression
 - Selection (σ)
 - Join (\bowtie)
 - Project(Π)
 - Union, intersection and difference, cross product
- Query optimizer
 - Enumerates the possible plans to evaluate expression,
 - Select a small subset of these plans and estimate their cost



Query Blocks: Units of Optimization

- An SQL query is parsed into a collection of *query blocks*, and these are optimized one block at a time.
 - A query block is an SQL query with no nesting and exactly one SELECT clause and one FROM clause and at most one WHERE clause, etc.
- Nested blocks are usually treated as calls to a subroutine, made once per outer tuple.
 - This is an over-simplification, but serves for now.


```
SELECT S.sname
FROM Sailors S
WHERE S.age IN
  (SELECT MAX (S2.age)
   FROM Sailors S2
   GROUP BY S2.rating)
```

Outer block

Nested block

Relational Algebra Tree and Evaluation Plan

- A plan consists of an extended relational algebra tree
 - Similar to a parse tree for an arithmetic expression
- Additional annotations are used to indicate the access and/or implementation method
 - Plan with/without indexes, etc.
- Problems
 - Which access method should be used for an operation?
 - Which plan should be used?




- 1) query block
- 2) each block is translated into a relational algebraic expression
- 3) each expression is represented by a tree
- 4) each tree can have a number of execution plans because of relational algebra equivalence
- 5) The number of plans is large
- 6) Estimating each plan needs to solve these issues: estimate the cost each operator and the output size
- 7) General strategy is 1) do selection as early as possible; 2) do projection as early as possible; 3) do join last
- 8) when do join, consider only left-deep plans

Sailors(sid, sname, rating, age)
 Boats(bid, bname, color)
 Reserve(sid, bid, day)

Find the names of sailors who
 reserve a boat with bid = 100 and
 whose rating is greater than 5

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

$\Pi_{sname} (\sigma_{bid=100 \wedge rating > 5} (Reserve \bowtie_{sid=sid} Sailors))$



A plan is a tree with annotations
 that specify the access method
 for each operator


Sailors(sid, sname, rating, age)
 Boats(bid, bname, color)
 Reserve(sid, bid, day)

Find the names of sailors who
 reserve a boat with bid = 100 and
 whose rating is greater than 5


```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

$\Pi_{sname} (\sigma_{bid=100 \wedge rating > 5} (Reserve \bowtie_{sid=sid} Sailors))$

RA Tree:

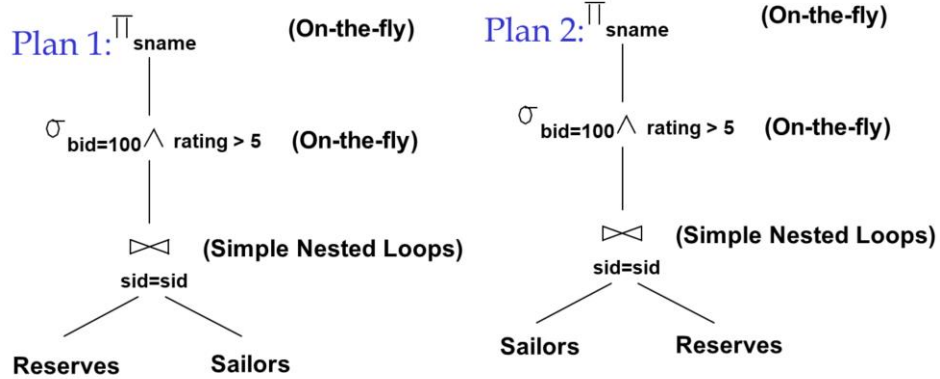


Plan:



There are usually many plans. Which one is the best? Optimization issue!

$$\Pi_{sname}(\sigma_{bid=100 \wedge rating > 5}(Reserve \bowtie_{sid=sid} Sailors))$$

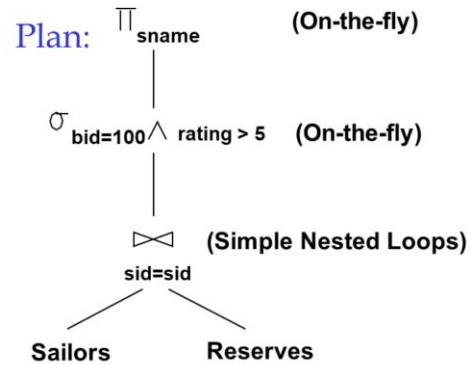
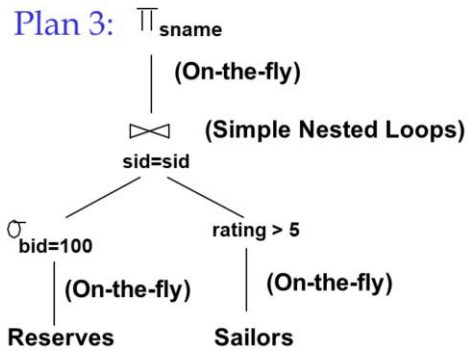


Question 1: are they the same?

Question 2: which one might be faster?

There are usually many plans. Which one is the best? Optimization issue!

$$\Pi_{sname}(\sigma_{bid=100 \wedge rating > 5}(Reserve \bowtie_{sid=sid} Sailors))$$



Relational Algebra Equivalences

- Allow us to choose different join orders and to 'push' selections and projections ahead of joins
- **Selections** $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots \sigma_{cn}(R))$ (*Cascade*)
 $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$ (*Commute*)
- **Projections:** $\pi_{a1}(R) \equiv \pi_{a1}(\dots (\pi_{an}(R)))$ (*Cascade*)
- **Joins:** $R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T$ (*Associative*)
 $(R \bowtie S) \equiv (S \bowtie R)$ (*Commute*)
- Show that: $R \bowtie (S \bowtie T) \equiv (T \bowtie R) \bowtie S$

Question 1: think about what those symbols mean

Question 2: think about why they are the same

More Equivalences To Explore

$$\sigma_{C \text{ AND } C'}(R) = \sigma_C(\sigma_{C'}(R)) = \sigma_C(R) \cap \sigma_{C'}(R)$$

$$\sigma_{C \text{ OR } C'}(R) = \sigma_C(R) \cup \sigma_{C'}(R)$$

$$\sigma_C(R \bowtie S) = \sigma_C(R) \bowtie S$$

Example: R(A, B, C, D), S(E, F, G)

$$\sigma_{F=3}(R \bowtie_{D=E} S) = \quad ?$$

$$\sigma_{A=5 \text{ AND } G=9}(R \bowtie_{D=E} S) = \quad ?$$

More Equivalences

- A projection commutes with a selection that only uses attributes retained by the projection
- Selection between attributes of the two arguments of a cross-product converts cross-product to a join
- A selection on just attributes of R commutes with $R \bowtie S$ (i.e., $\sigma(R \bowtie S) \equiv \sigma(R) \bowtie S$)
- Similarly, if a projection follows a join $R \bowtie S$, we can 'push' it by retaining only attributes of R (and S) that are needed for the join or are kept by the projection

Challenges of Query Optimization

- Two main issues:
 - For a given query, what plans are considered?
 - Search space is huge
 - How to estimate the cost of a plan?
- Solutions
 - Ideally: find best plan
 - Practically: avoid worst plans

Highlights of IBM System R Optimizer (1979)

- Impact
 - Most widely used currently; works well for < 10 joins
- Cost estimation
 - Statistics, maintained in system catalogs, used to estimate cost of operations and result sizes
 - Considers combination of CPU and I/O costs
- Plan Space
 - Only the space of *left-deep plans* is considered
 - Left-deep plans allow output of each operator to be pipelined into the next operator without storing it in a temporary relation
 - Cartesian products avoided

Cost Estimation

- Given a plan, we need to
 - Estimate *cost* of each operation in plan tree
 - Use the information recorded in statistics and system catalogs
 - Depends on input cardinalities
 - We've already discussed how to estimate the cost of operations (sequential scan, index scan, joins, etc.)
 - Estimate *size of result* for each operation in tree
 - Use information about the input relations
 - For selections and joins, assume independence of predicates

System Catalogs

- For each relation:
 - name, file name, file structure (e.g., Heap file)
 - attribute name and type, for each attribute
 - index name, for each index
 - integrity constraints
- For each index:
 - structure (e.g., B+ tree) and search key fields
- For each view:
 - view name and definition
- Plus statistics, authorization, buffer pool size, etc.

Catalogs are themselves stored as relations!

Statistics stored in Catalog

- Cardinality
 - Number of tuples/rows $NTuples(R)$ for each relation R
- Size
 - Number of pages $NPages(R)$ for each relation R
- Index Cardinality
 - Number of distinct key values $NKeys(I)$ for each index I
- Index Size
 - Number of pages $INPages(I)$ for each index I
 - For example, for a B+ tree index, we take the number of leaf pages to be the index size
- Index Height
 - Number of nonleaf levels $IHeight(I)$ for each tree index I
- Index Range
 - The minimum present key value $ILow(I)$ and the maximum present key value $IHigh(I)$ for each index I

These statistics are updated periodically

Size Estimation and Reduction Factors

- Consider a query block:

```
SELECT attribute list
FROM relation list
WHERE term1 AND ... AND termk
```

- Maximum # tuples in result is the product of the cardinalities of relations in the FROM clause
- Reduction factor (RF) associated with each term reflects the impact of the term in reducing result size
- Result cardinality = Max # tuples * product of all RF's.
 - Implicit assumption that terms are independent!
 - Term col=value has RF $1/NKeys(I)$, given index I on col
 - Term col1=col2 has RF $1/MAX(NKeys(I1), NKeys(I2))$
 - Term col>value has RF $(High(I)-value)/(High(I)-Low(I))$

Optimization Strategies

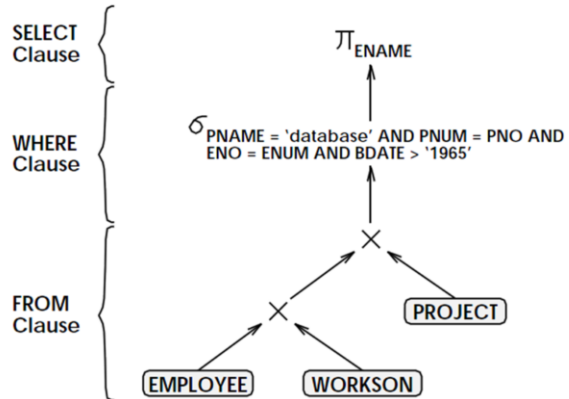
- Move SELECTs and PROJECTs as far down as possible
- Among the SELECTs, order them such that the lowest selectivity factor is performed first
- Among Joins, order them such that the join with the lowest join selectivity factor is performed first

EMPLOYEE(ENUM, ENAME, BDATE)
 PROJECT(PNUM, PNAME)
 WORKSON(ENO, PNO) where ENO is a FK to Employee(ENUM) and PNO is a FK to PROJECT(PNUM)


SELECT ENAME
 FROM EMPLOYEE, PROJECT, WORKSON
 WHERE PNAME='Database'
 AND PNUM=PNO
 AND ENUM=ENO
 AND BDATE>'1965'

Canonical Query Tree

IBM's System R




Canonical Query Tree




EMPLOYEE(ENUM, ENAME, BDATE)
 PROJECT(PNUM, PNAME)
 WORKSON(ENO, PNO)

Move SELECTs down

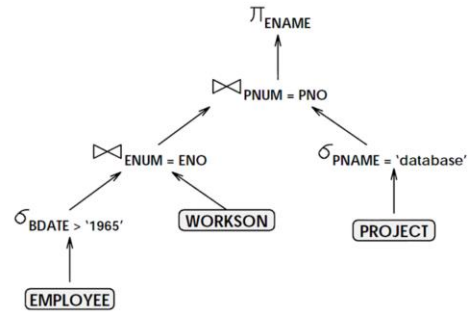
After Optimization

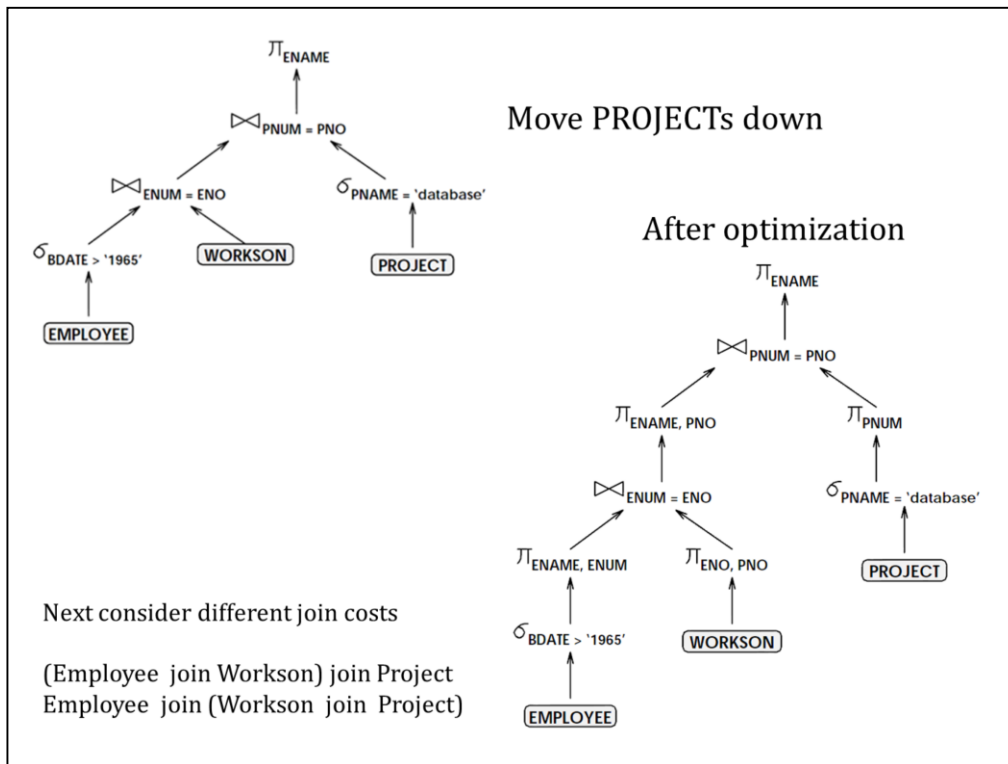


Replace “ $\sigma - \times$ ” by “ \bowtie ”



After Optimization





Question: Can you recall the cost of join?

Join Optimization

- Consider just three relations R, S, and T

Form 1:

$(R \bowtie S) \bowtie T$	$(S \bowtie R) \bowtie T$
$(R \bowtie T) \bowtie S$	$(T \bowtie R) \bowtie S$
$(S \bowtie T) \bowtie R$	$(T \bowtie S) \bowtie R$

Form 2:

$R \bowtie (S \bowtie T)$	$R \bowtie (T \bowtie S)$
$S \bowtie (R \bowtie T)$	$S \bowtie (T \bowtie R)$
$T \bowtie (R \bowtie S)$	$T \bowtie (S \bowtie R)$

Form 1: $(\bullet \bowtie \bullet) \bowtie \bullet$

Form 2: $\bullet \bowtie (\bullet \bowtie \bullet)$

The **number** of *different join orderings* of *n relations* is **exponentially large** !!!



(Because the **number** of **permutations** is **exponentially large**)

- The **number** of *possible join trees* to consider is just *too large*


- We **need** to **reduce** the **search space**....

Plans to consider


- Fundamental decision in System R: *only left-deep join trees* are considered.
 - As the number of joins increases, the number of alternative plans grows rapidly; we need to restrict the search space.
 - Left-deep trees allow us to generate all *fully pipelined plans*.
 - Intermediate results not written to temporary files.



Three-way Join



- Possible left-deep join trees:



Enumeration of Left-Deep Plans

- Left-deep plans differ only in the order of relations, the access method for each relation, and the join method for each join
- Enumerated using N passes (if N relations joined):
 - **Pass 1:** Find best 1-relation plan for each relation
 - **Pass 2:** Find best way to join result of each 1-relation plan (as outer) to another relation (*All 2-relation plans*)
 - **Pass N:** Find best way to join result of a (N-1)-relation plan (as outer) to the N'th relation (*All N-relation plans*)
- For each subset of relations, retain only:
 - Cheapest plan overall, plus
 - Cheapest plan for each *interesting order* of the tuples
- In spite of pruning plan space, this approach is **still exponential** in the # of tables. (works well for most queries with less than 15 tables)
- **ORDER BY, GROUP BY, aggregates** etc. handled as a final step, using either an 'interestingly ordered' plan or an additional sorting operator.

Pass 1 question: what access methods should be used for this selection? (B+ tree? Hash index? Sequential scan?)

Pass2 question: which two relations should join? How to join?

Summary

- Two parts to optimize a query
 - Consider a set of alternative plans, typically, left-deep plans only
 - Must estimate cost of each plan that is considered
 - Must estimate size of result and cost for each plan node
 - Key issues: Statistics, indexes, operator implementations
- Single-relation queries
 - All access paths considered, cheapest is chosen
 - Issues: Selections that *match* index, whether index key has all needed fields and/or provides tuples in a desired order
- Multiple-relation queries
 - All single-relation plans are first enumerated
 - Selections/projections considered as early as possible
 - For each 1-relation plan, all ways of joining another relation are considered
 - For each 2-relation plan that is 'retained', all ways of joining another relation are considered, etc.