


What Is Data Mining?

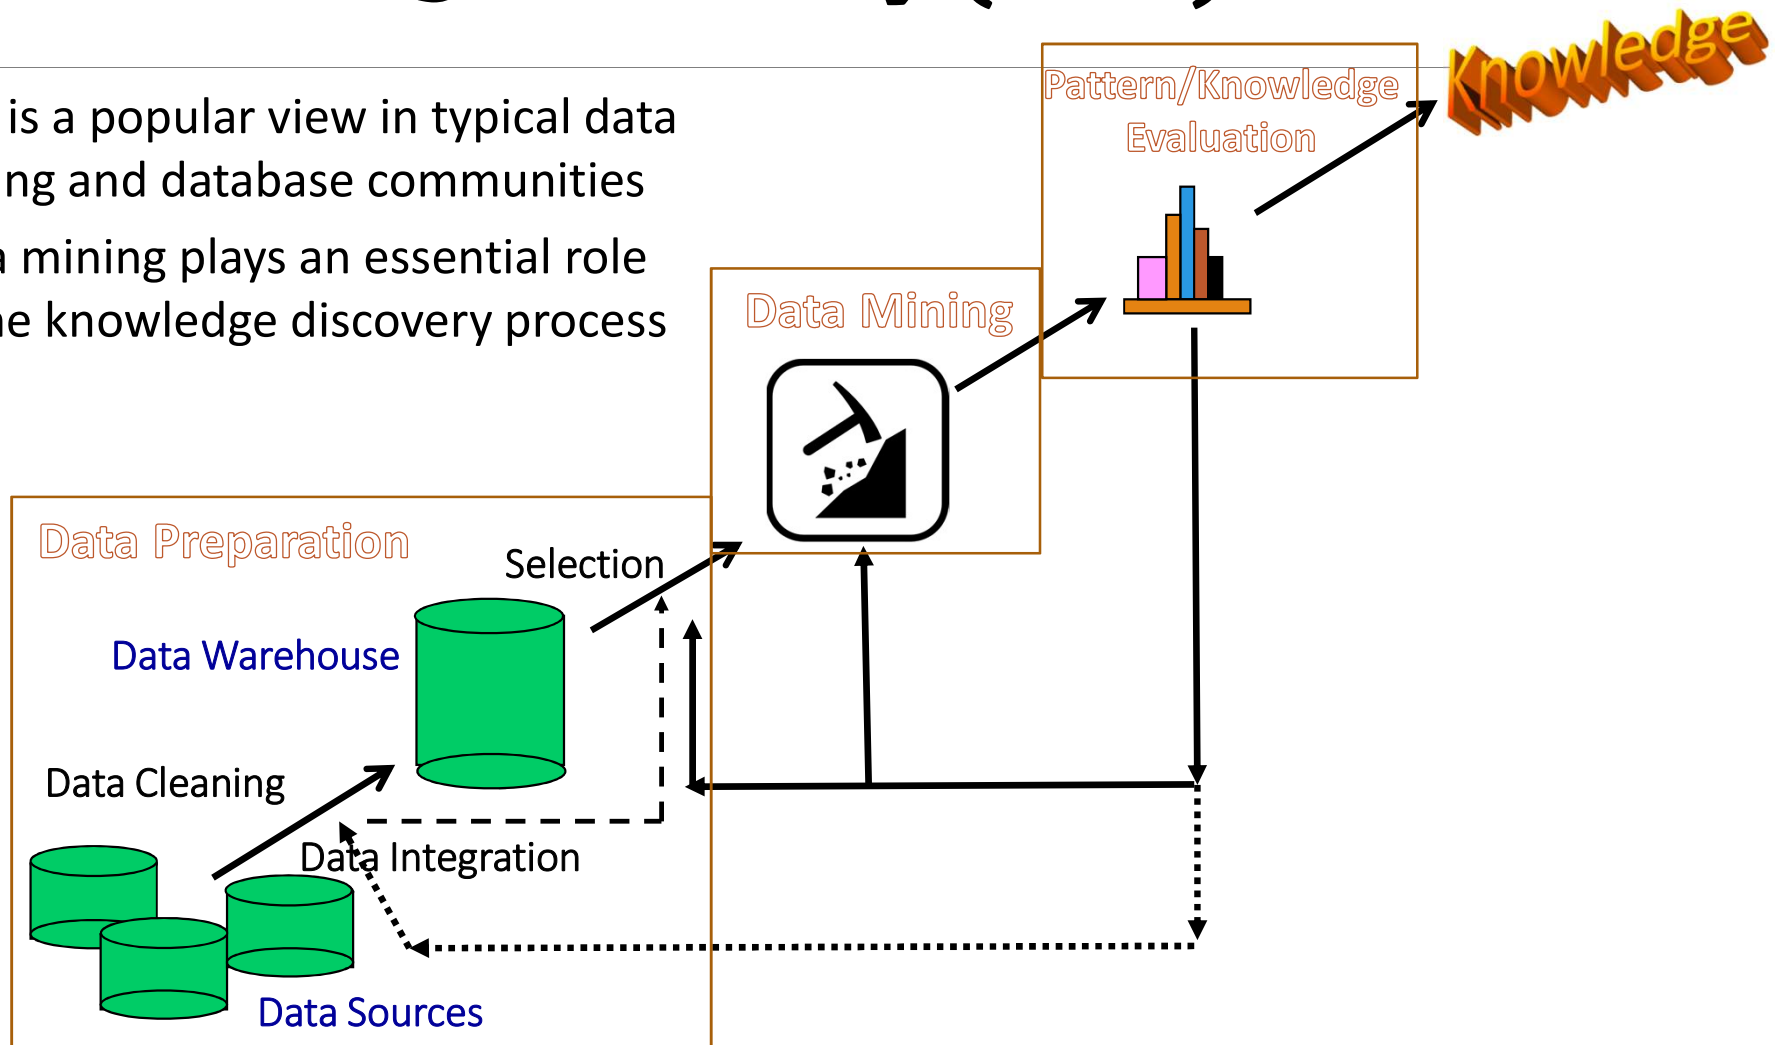


- ❑ Data mining (knowledge discovery from data)
 - ❑ Extraction of interesting (non-trivial, implicit, previously unknown and potentially useful) patterns or knowledge from huge amount of data
- ❑ Alternative names
 - ❑ Knowledge discovery (mining) in databases (KDD), knowledge extraction, data/pattern analysis, data archeology, data dredging, information harvesting, business intelligence, etc.

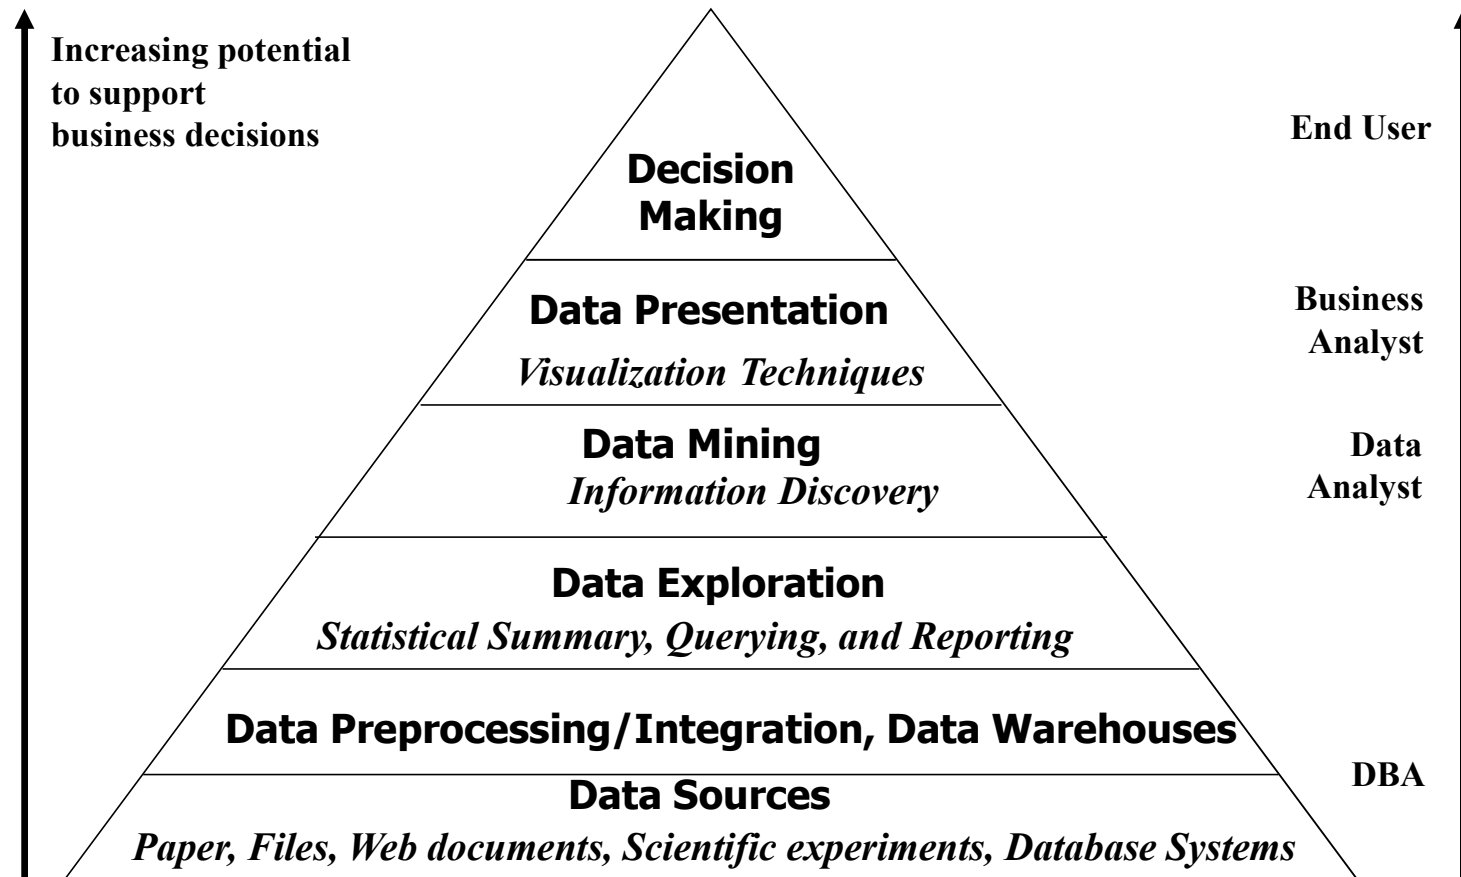


Knowledge Discovery (KDD) Process

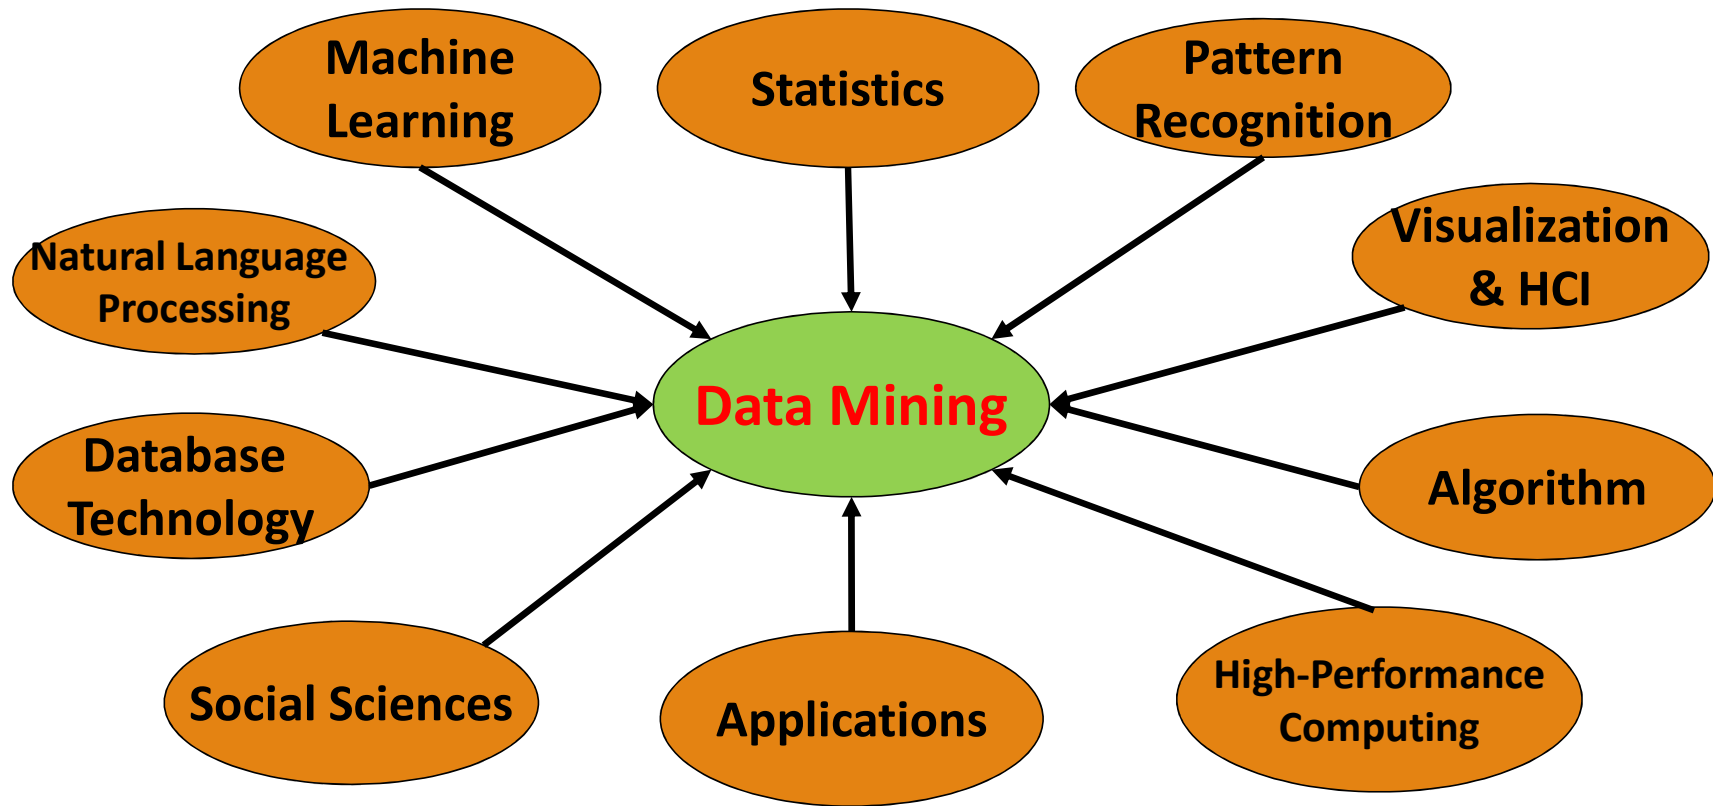
- ❑ This is a popular view in typical data mining and database communities
- ❑ Data mining plays an essential role in the knowledge discovery process



Data Mining in Business Intelligence



Data Mining: Confluence of Multiple Disciplines



Recommender System

- ❑ Product recommendation (Amazon, EBay)
 - ❑ Search recommendation (Google, Bing)
 - ❑ Video/music/post recommendation (Netflix, Pandora, Pinterst)
 - ❑ Friend recommendation (Facebook, twitter)
 - ❑ Job recommendation (linkIn)
-
- ❑ collaborative filtering
 - ❑ content-based filtering
 - ❑ hybrid

Commerce, Profiling and Finance

- Planning and Forecasting

 - Dynamic pricing

 - Ads bidding

- Profiling

 - User profiling

 - Churn Prediction: knowing which users are going to stop using your platform in the future.

 - Product profiling

- Fintech

 - Stock market

 - Sentiment analysis

Urban Planning

- ❑ Energy and power
- ❑ Traffic prediction and management
 - ❑ Parking detection
 - ❑ Traffic control
- ❑ Transportation sharing system
 - ❑ Uber
 - ❑ Bike-sharing
- ❑ Pollution
 - ❑ Air quality prediction

Medicine and Healthcare

- ❑ Disease prediction
 - ❑ Computer Aided Detection
 - ❑ EHR
 - ❑ Risk prediction
 - ❑ Disease progression prediction
- ❑ Healthcare
 - ❑ Epidemic and outbreak prediction
 - ❑ Food safety
- ❑ Medicine study
 - ❑ Drug discovery and prediction
- ❑ Bioinformatics

Other Sciences and Applications

- ❑ Education
 - ❑ MOOC (massive open online course)
- ❑ Political science and Social science
 - ❑ Fake news
 - ❑ Crime and terrorist detection
 - ❑ Disaster detection
 - ❑ Opinion mining
 - ❑ Social influence
- ❑ Environmental Science
 - ❑ Climate

Pattern Discovery: Basic Concepts

- ❑ What Is Pattern Discovery? Why Is It Important?
- ❑ Basic Concepts: Frequent Patterns and Association Rules
- ❑ Compressed Representation: Closed Patterns and Max-Patterns

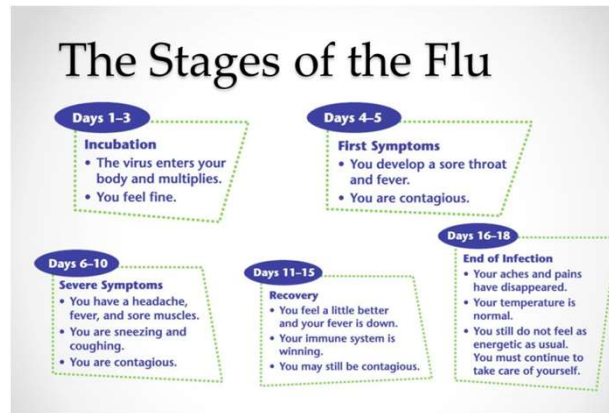
What are Patterns?

❑ What are patterns?


- ❑ **Patterns**: A set of items, subsequences, or substructures that occur frequently together (or strongly correlated) in a data set
- ❑ Patterns represent **intrinsic** and **important properties** of datasets



Frequent item set



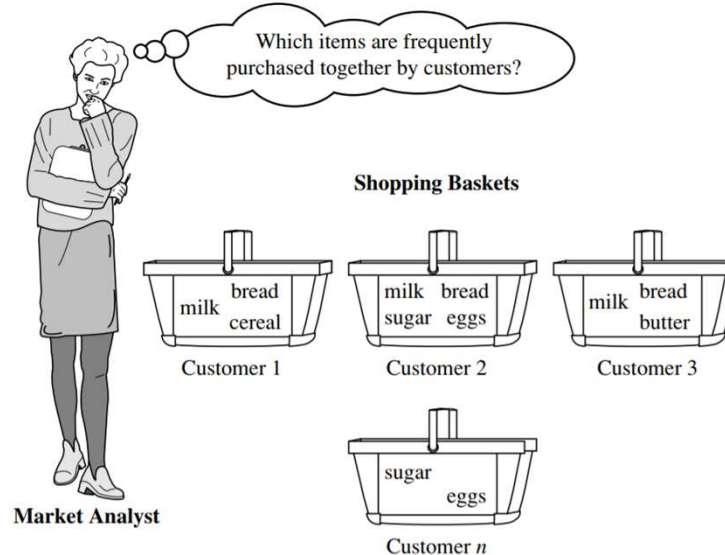
Frequent sequences



Frequent structures

What Is Pattern Discovery?

- ❑ **Pattern discovery:** Uncovering patterns from massive data sets
- ❑ It can answer questions such as:
 - ❑ What products were often purchased together?
 - ❑ What are the subsequent purchases after buying an iPad?



Pattern Discovery: Why Is It Important?

- ❑ **Foundation** for many essential data mining tasks
 - ❑ Association, correlation, and causality analysis
 - ❑ Mining **sequential**, structural (e.g., sub-graph) patterns
- ❑ Broad applications
 - ❑ Market basket analysis, cross-marketing, catalog design, sale campaign analysis, Web log analysis, biological sequence analysis
 - ❑ Many types of data: spatiotemporal, multimedia, time-series, and stream data

Basic Concepts: Transactional Database

□ Transactional Database (TDB)

- Each transaction is associated with an identifier, called a TID.
- May also have counts associated with each item sold

Tid	Items bought
1	Beer, Nuts, Diaper
2	Beer, Coffee, Diaper
3	Beer, Diaper, Eggs
4	Nuts, Eggs, Milk
5	Nuts, Coffee, Diaper, Eggs, Milk

Basic Concepts: k-Itemsets and Their Supports

- Itemset: A set of one or more items

$$I = \{I_1, I_2, \dots, I_m\}$$

- k-itemset: An itemset containing k items:

$$X = \{x_1, \dots, x_k\}$$

- Ex. {Beer, Nuts, Diaper} is a 3-itemset

- Absolute support (count)

- sup{X} = occurrences of an itemset X

- Ex. sup{Beer} = 3
- Ex. sup{Diaper} = 4
- Ex. sup{Beer, Diaper} = 3
- Ex. sup{Beer, Eggs} = 1

Tid	Items bought
1	Beer, Nuts, Diaper
2	Beer, Coffee, Diaper
3	Beer, Diaper, Eggs
4	Nuts, Eggs, Milk
5	Nuts, Coffee, Diaper, Eggs, Milk

- Relative support

- s{X} = The fraction of transactions that contains X (i.e., the probability that a transaction contains X)
- Ex. s{Beer} = 3/5 = 60%
- Ex. s{Diaper} = 4/5 = 80%
- Ex. s{Beer, Eggs} = 1/5 = 20%

Basic Concepts: Frequent Itemsets (Patterns)

- An itemset (or a pattern) X is *frequent* if the support of X is no less than a *minsup* threshold σ
- Let $\sigma = 50\%$ (σ : *minsup* threshold) for the given 5-transaction dataset

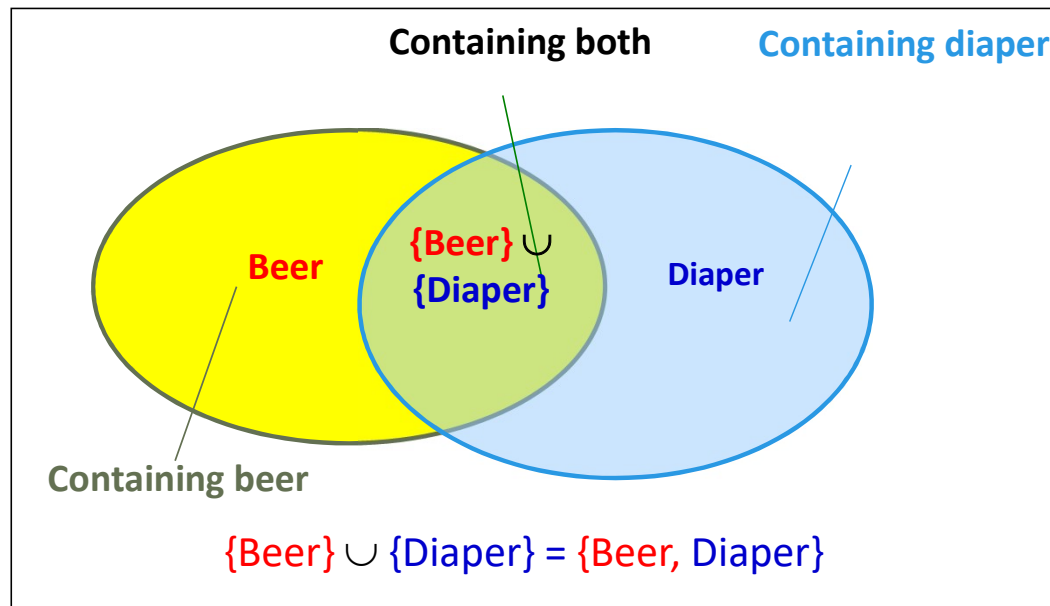


Tid	Items bought
1	Beer, Nuts, Diaper
2	Beer, Coffee, Diaper
3	Beer, Diaper, Eggs
4	Nuts, Eggs, Milk
5	Nuts, Coffee, Diaper, Eggs, Milk

- All the frequent 1-itemsets:
 - Beer: 3/5 (60%); Nuts: 3/5 (60%); Diaper: 4/5 (80%); Eggs: 3/5 (60%)
 - All the frequent 2-itemsets:
 - {Beer, Diaper}: 3/5 (60%)
 - All the frequent 3-itemsets?
 - None
- Why do these itemsets (shown on the left) form the complete set of frequent k -itemsets (patterns) for any k ?
 - **Observation:** We may need an efficient method to mine a complete set of frequent patterns

From Frequent Itemsets to Association Rules

- Compared with itemsets, association rules can be more telling
 - Ex. *Diaper* \rightarrow *Beer*
 - *Buying diapers may likely lead to buying beers*



Note: $X \cup Y$: the union of two itemsets

■ The set contains both X and Y

Association Rules

- How do we compute the strength of an association rule $X \rightarrow Y$ (Both X and Y are itemsets)?

- We first compute the following two metrics, s and c .

- **Support of $X \cup Y$**

- Ex. $s\{\text{Diaper, Beer}\} = 3/5 = 0.6$ (i.e., 60%)

- **Confidence of $X \rightarrow Y$**

- The *conditional probability* that a transaction containing X also contains Y :

$$c = \text{sup}(X, Y) / \text{sup}(X)$$

- Ex. $c = \text{sup}\{\text{Diaper, Beer}\} / \text{sup}\{\text{Diaper}\} = 3/4 = 0.75$

- In pattern analysis, we are often interested in those rules that dominate the database, and these two metrics ensure the popularity and correlation of X and Y .

Tid	Items bought
1	Beer, Nuts, Diaper
2	Beer, Coffee, Diaper
3	Beer, Diaper, Eggs
4	Nuts, Eggs, Milk
5	Nuts, Coffee, Diaper, Eggs, Milk

Mining Frequent Itemsets and Association Rules

Association rule mining

- Given two thresholds: $minsup$, $minconf$
- Find **all** of the rules, $X \rightarrow Y (s, c)$ such that $s \geq minsup$ and $c \geq minconf$

Let $minsup = 50\%$

- Freq. 1-itemsets: Beer: 3, Nuts: 3, Diaper: 4, Eggs: 3
- Freq. 2-itemsets: {Beer, Diaper}: 3

Let $minconf = 50\%$

- $Beer \rightarrow Diaper$ (60%, 100%)
- $Diaper \rightarrow Beer$ (60%, 75%)

(Q: Are these all rules?)



Tid	Items bought
1	Beer, Nuts, Diaper
2	Beer, Coffee, Diaper
3	Beer, Diaper, Eggs
4	Nuts, Eggs, Milk
5	Nuts, Coffee, Diaper, Eggs, Milk

Observations:

- Mining association rules and mining frequent patterns are very close problems
- Scalable methods are needed for mining large datasets

Challenge: There Are Too Many Frequent Patterns!

- ❑ A long pattern contains a combinatorial number of sub-patterns
- ❑ How many frequent itemsets does the following TDB₁ contain (minsup = 2)?

❑ TDB₁: T₁: {a₁, ..., a₅₀}; T₂: {a₁, ..., a₁₀₀}

❑ Let's have a try

1-itemsets: {a₁}: 2, {a₂}: 2, ..., {a₅₀}: 2

2-itemsets: {a₁, a₂}: 2, ..., {a₁, a₅₀}: 2, {a₂, a₃}: 2 ..., ..., {a₄₉, a₅₀}: 2,

..., ..., ..., ...

49-itemsets: {a₁, a₂, ..., a₄₉}: 2, ..., {a₂, a₃, ..., a₅₀}: 2

50-itemset: {a₁, a₂, ..., a₅₀}: 2

- ❑ The total number of frequent itemsets:

$$\binom{50}{1} + \binom{50}{2} + \dots + \binom{50}{50} = 2^{50} - 1$$

A too huge set for any one to compute or store!



Expressing Patterns in Compressed Form

- How to reduce the redundancy of the list of all the frequent itemsets?
- Solution 1: **Closed patterns**: A pattern (itemset) X is **closed** if X is *frequent*, and there exists *no super-pattern* $Y \supset X$, with the same support as X
 - Ex. TDB_1 : $T_1: \{a_1, \dots, a_{50}\}$; $T_2: \{a_1, \dots, a_{100}\}$; $T_3: \{a_1, \dots, a_{10}\}$
 - Suppose $minsup = 2$. How many closed patterns does TDB_1 contain?
 - Two: $P_1: \{\{a_1, \dots, a_{50}\}: 2\}$; $P_2: \{\{a_1, \dots, a_{10}\}: 3\}$

Expressing Patterns in Compressed Form: Closed Patterns

- Closed pattern is a lossless compression of frequent patterns
 - Lossless: no information loss
 - Reduces the # of patterns but does not lose the support information!
 - Given $P_1: \{a_1, \dots, a_{50}\}: 2$; $P_2: \{a_1, \dots, a_{10}\}: 3$;
 - You will still be able to say: $\{a_2, \dots, a_{40}\}: 2$, $\{a_5, a_1\}: 3$

Expressing Patterns in Compressed Form: Max-Patterns


- ❑ Solution 2: **Max-patterns**: A pattern X is a **max-pattern** if X is frequent and there exists no frequent super-pattern $Y \supset X$
- ❑ Difference from close-patterns?
 - ❑ Do not care the real support of the sub-patterns of a max-pattern
 - ❑ Let Transaction DB TDB_1 : $T_1: \{a_1, \dots, a_{50}\}; T_2: \{a_1, \dots, a_{100}\}; T_3: \{a_1, \dots, a_{10}\}$
 - ❑ Suppose $minsup = 2$. How many max-patterns does TDB_1 contain?
 - ❑ One: $P: \{\{a_1, \dots, a_{50}\}: 2\}$

Expressing Patterns in Compressed Form: Max-Patterns

- **Max-pattern** is a **lossy compression**!
 - We only any subset of the max-pattern $P:\{a_1, \dots, a_{50}\}$ is frequent
 - But we do not know the real support of $\{a_1, \dots, a_{10}\}, \dots$, any more!
 - More compressed than closed pattern (that is smaller in size)

The Downward Closure Property of Frequent Patterns

- ❑ **Frequent** itemset: $\{a_1, \dots, a_{50}\}$
 - ❑ Subsets are all **frequent**: $\{a_1\}, \{a_2\}, \dots, \{a_{50}\}, \{a_1, a_2\}, \dots, \{a_1, \dots, a_{49}\}, \dots$
- ❑ Downward closure (Apriori): Any subset of a frequent itemset must be frequent
 - ❑ If **{beer, diaper, nuts}** is frequent, so is **{beer, diaper}**
 - ❑ If **any subset of an itemset S** is **infrequent**, then there is no chance for S to be frequent.

 A sharp knife for pruning!

Apriori: A Candidate Generation & Test Approach

- ❑ Outline of Apriori (level-wise, candidate generation and test)
 - ❑ **Scan** DB once to get frequent 1-itemset
 - ❑ **Repeat**
 - ❑ Generate length-(k+1) candidate itemsets from length-k frequent itemsets
 - ❑ Test the candidates against DB to find frequent (k+1)-itemsets
 - ❑ Set $k := k + 1$
 - ❑ **Until** no frequent or candidate set can be generated
 - ❑ Return all the frequent itemsets derived

The Apriori Algorithm—An Example

Database TDB

Tid	Items
10	A, C, D
20	B, C, E
30	A, B, C, E
40	B, E

minsup = 2

1st scan

C_1

Itemset	sup
{A}	2
{B}	3
{C}	3
{D}	1
{E}	3

F_1

Itemset	sup
{A}	2
{B}	3
{C}	3
{E}	3

C_2

Itemset	sup
{A, B}	1
{A, C}	2
{A, E}	1
{B, C}	2
{B, E}	3
{C, E}	2

2nd scan

C_2

Itemset
{A, B}
{A, C}
{A, E}
{B, C}
{B, E}
{C, E}

F_2

Itemset	sup
{A, C}	2
{B, C}	2
{B, E}	3
{C, E}	2


C_3

Itemset
{B, C, E}

3rd scan

F_3

Itemset	sup
{B, C, E}	2



The Apriori Algorithm (Pseudo-Code)

C_k : Candidate itemset of size k

F_k : Frequent itemset of size k

$K := 1$;

$F_k := \{\text{frequent items}\}$; // frequent 1-itemset

While ($F_k \neq \emptyset$) **do** { // when F_k is non-empty

$C_{k+1} := \text{candidates generated from } F_k$; // candidate generation

 Derive F_{k+1} by counting candidates in C_{k+1} with respect to TDB at minsup;

$k := k + 1$

}

return $\cup_k F_k$ // return F_k generated at each level

Candidate Generation (Pseudo-Code)

□ Suppose the items in F_{k-1} are listed in an order

□ // Step 1: Joining

for each p in F_{k-1}

for each q in F_{k-1}

if $p.item_1 = q.item_1, \dots, p.item_{k-2} = q.item_{k-2}, p.item_{k-1} < q.item_{k-1}$ {

$c = \text{join}(p, q)$


□ // Step 2: pruning

if $\text{has_infrequent_subset}(c, F_{k-1})$

continue // prune


else add c to C_k

}



Partitioning: Scan Database Only Twice

- Theorem: *Any itemset that is potentially frequent in TDB must be frequent in at least one of the partitions of TDB*



Partitioning: Scan Database Only Twice

- ❑ Method: Scan DB **twice** (A. Savasere, E. Omiecinski and S. Navathe, *VLDB'95*)
 - ❑ Scan 1: Partition database so that each partition can fit in main memory (why?)
 - ❑ Mine local frequent patterns in this partition
 - ❑ Scan 2: Consolidate global frequent patterns
 - ❑ Find global frequent itemset candidates (those frequent in at least one partition)
 - ❑ Find the true frequency of those candidates, by scanning TDB_i one more time

Why Mining Frequent Patterns by Pattern Growth?

- ❑ Apriori: A *breadth-first search* mining algorithm
 - ❑ First find the complete set of frequent k -itemsets
 - ❑ Then derive frequent $(k+1)$ -itemset candidates
 - ❑ Scan DB again to find true frequent $(k+1)$ -itemsets
- ❑ Are there depth-first search algorithm?
 - ❑ Yes