

Sockets Programming

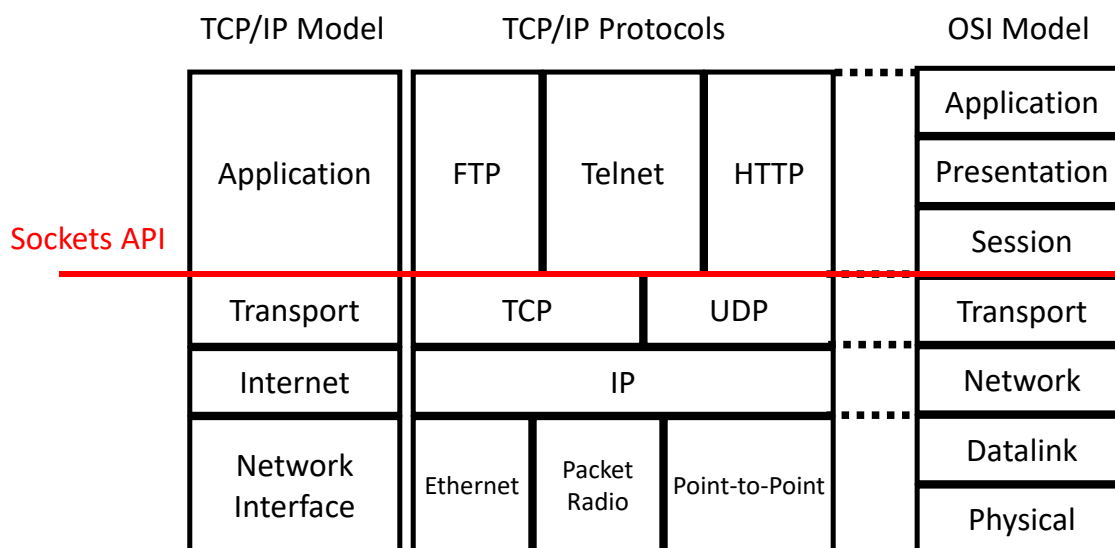
What is a socket?

- ⊕ A socket is a **communication endpoint** in the computer network
- ⊕ Sockets work with UNIX I/O services just like files
- ⊕ Sockets have special needs:
 - Specifying addresses for communication endpoints
 - Establishing a connection
- ⊕ For applications to interact with communication protocol suites via sockets, we need sockets API (Application Programming Interface)

Sockets API

- ✦ `sd = socket(protofamily, type, protocol)`
- ✦ `connect(sd, remoteaddr, addrlen)`
- ✦ `bind(sd, localaddr, addrlen)`
- ✦ `listen(sd, qlen)`
- ✦ `td = accept(sd, &clientaddr, &addrlen)`
- ✦ `close(sd)`
- ✦ `send(td, data, length, flags)`
- ✦ `recv(td, &buffer, length, flags)`
- ✦ `select(maxfd, readset, writeset, exceptset, timeout)`
- ✦ `get/setsockopt(sd, ...)`
- ✦ Others:
 - `sendto, recvfrom, poll, read, write`

Where is Sockets API?



What is a client-server paradigm?

⊕ Server:

- Server is a process with resources
- A computer can host multiple servers
- Server waits for connection requests from clients
- Server specifies **well-known port #** when creating socket

⊕ Client:

- Client is a process that needs resources
- Client is assigned **ephemeral port #**
- Client initiates connection with server
- Client needs to know server's IP address & port #
- Server learns client's address & port #

Types of Servers

⊕ Iterative Server

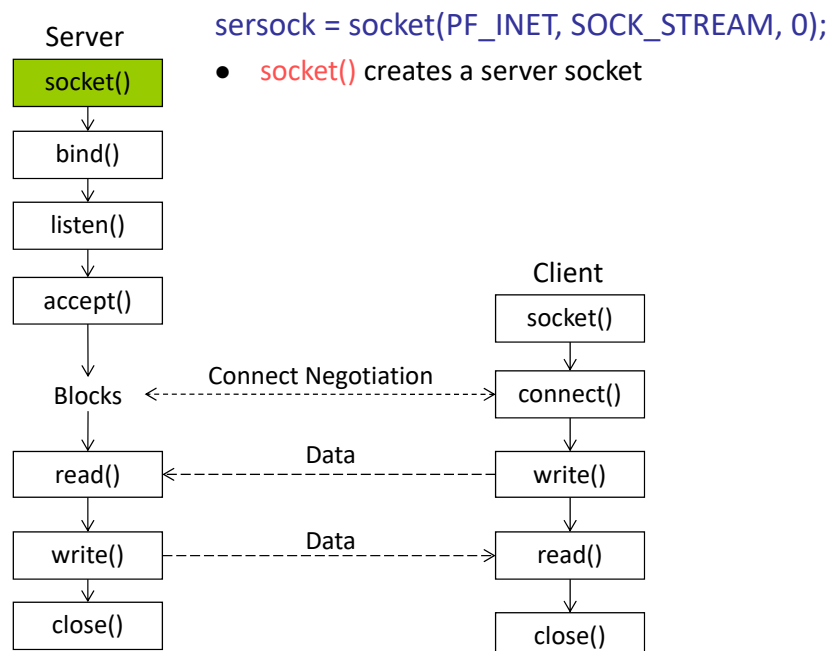
- Accepts only one connection at any time

⊕ Concurrent Server

- May accept more than one connection and serve all connections concurrently
- When the concurrent server receives a request, it forks a child server to serve that request; the server itself waits for a new request
- When the child server finishes the service, it exits

TCP Sockets Programming

[Server Side] Step 1: Create a TCP Socket



Create a TCP Socket

socket()

```
% man -s 2 socket
cc [flag ...] file ... -lsocket -lnsl [library ...]
#include <sys/types.h>
#include <sys/socket.h>
int socket(int domain, int type, int protocol);
```

+ domain (family)

- PF_UNIX: Unix system internal protocols
- PF_INET: IPv4 Internet protocols

+ type

- SOCK_STREAM: stream socket, for connection-oriented
- SOCK_DGRAM: datagram socket, for connectionless
- SOCK_RAW: raw socket, bypassing the transport layer

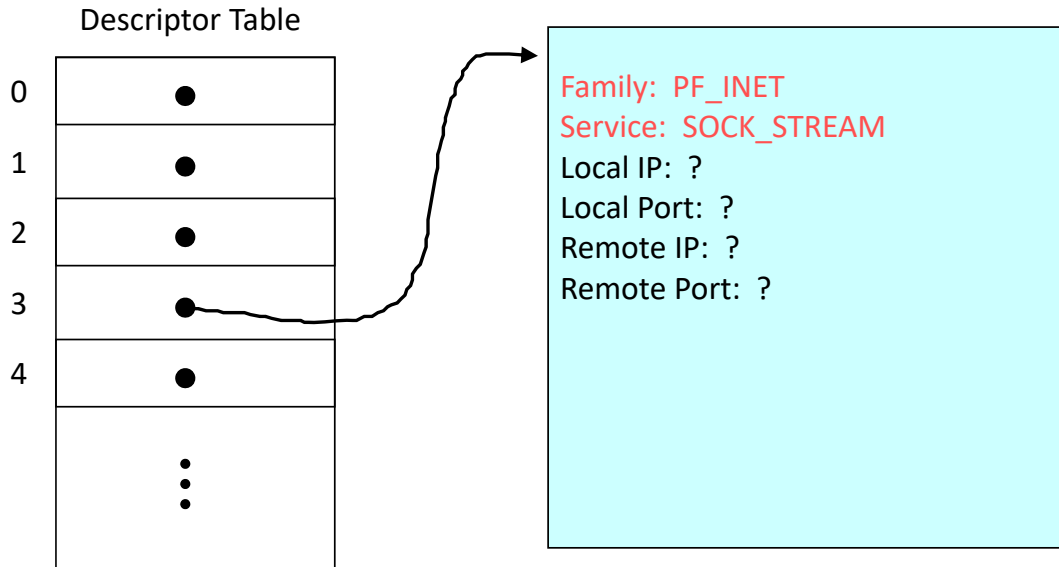
+ protocol

- IPPROTO_TCP: used with SOCK_STREAM, TCP Protocol
- IPPROTO_UDP: used with SOCK_DGRAM, UDP Protocol
- IPPROTO_IP: used with SOCK_RAW, IP protocol
- IPPROTO_ICMP: used with SOCK_RAW, ICMP protocol
- IPPROTO_RAW: used with SOCK_RAW, bypassing the IP layer
- 0: the default protocol

+ socket() system call returns

- a non-negative integer (socket descriptor/handle)
- or -1 on error

Socket Descriptor Data Structure



Error Checking and Reporting

- ✚ Whenever an error occurs during a system call
 - ➡ System call sets the value of a global variable `errno`
 - ➡ To print out the error message, you should use the special "`void perror(const char * string)`" function right after the error was produced; otherwise, `errno` may be overwritten in calls to other functions
 - ✓ `perror("The following error occurred:");`
 - ✗ `printf("The error code is %d\n", errno);`

Error Checking and Reporting

- ⊕ Do error checking after every system call

- ⊕ Example:

```
if ( socket(PF_INET, SOCK_STREAM, 0) < 0 )
{
    perror("The following error occurred:");
    exit(1);
}
```

- ⊕ Please don't forget to do error checking in your labs!

[Server Side] Step 2: Specify an Endpoint Address for Server

- ⊕ TCP/IP uses an IP address and a port number to specify an endpoint address

- Other protocol suites (families) may use other addressing schemes

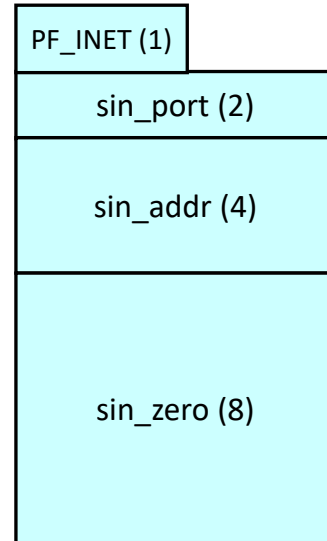
- ⊕ Ports

- Each user process associates its connection with the transport layer protocol with a unique port # (16 bits)
- The same port number can be used for different protocols
 - E.g., you may use port 2000 with TCP and port 2000 with UDP, but you cannot have two TCP services both with port 2000
- Some well-known ports are reserved for well-known services
 - TCP port 21 for [ftp](#)
 - TCP port 80 for [http](#)
- Ports 0 ~ 1023 are reserved, so user processes must use ports ≥ 1024

IPv4 Socket Address Structure

```
struct sockaddr_in {
    sa_family_t sin_family; // PF_INET
    in_port_t sin_port;     // 16-bit port #
    struct in_addr sin_addr; // 32-bit IPv4 address
    char sin_zero[8];       // unused
};
```

```
struct in_addr {
    // 32-bit IPv4 address in network byte order
    in_addr_t s_addr;
};
```



Byte Order

- ⊕ Little-Endian Byte Order
 - Little end of a multi-byte number is stored at the starting address
- ⊕ Big-Endian Byte Order
 - Big end of a multi-byte number is stored at the starting address
- ⊕ Example: `c = 0x 0001`

`c[0]` `c[1]`

 - little-endian:
 - big-endian:
- ⊕ Different systems have different byte orders (host byte order)
- ⊕ Internet protocols use big-endian byte ordering!

Network Byte Order

- ⊕ All values stored in a `sockaddr_in` must be in network byte order
 - `sin_port`: a 16-bit TCP or UDP port number
 - `sin_addr`: a 32-bit IP address

Common Mistake: Ignoring Network Byte Order!

Network Byte Order Functions

'h' : host byte order

'n' : network byte order

's' : short (16 bit)

'l' : long (32 bit)

⊕ `uint16_t htons(uint16_t);`

⊕ `uint16_t ntohs(uint16_t);`

⊕ `uint32_t htonl(uint32_t);`

⊕ `uint32_t ntohl(uint32_t);`

IPv4 Address Conversion Functions

⊕ `in_addr_t inet_addr(const char *)`

- ➡ convert an ASCII dotted-decimal IP address to a 32-bit network byte ordered IP address

⊕ `char *inet_ntoa(struct in_addr)`

- ➡ convert a network byte ordered value to an ASCII dotted-decimal string

⊕ Example:

```
char IP_ADDRESS[16] = "129.186.23.86";  
serveraddr.sin_addr.s_addr = inet_addr(IP_ADDRESS);
```

Specify an Endpoint Address for Server

⊕ Specifying an endpoint address using the socket address structure:

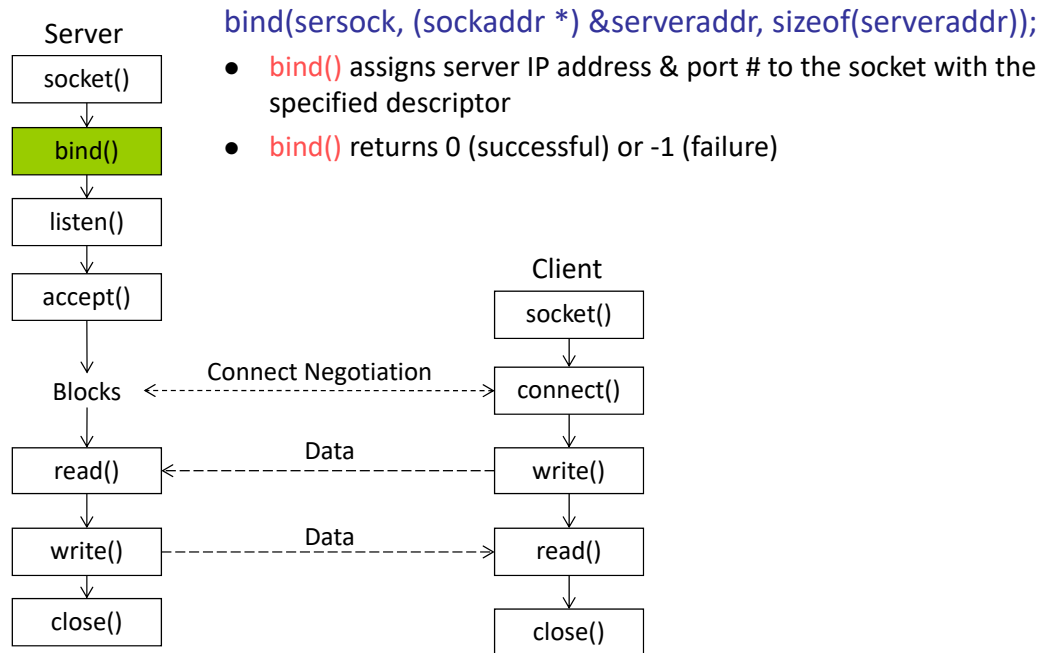
```
struct sockaddr_in serveraddr;
```

```
serveraddr.sin_family = PF_INET;
```

```
serveraddr.sin_port = htons( 80 );
```

```
serveraddr.sin_addr.s_addr = inet_addr( IP_ADDRESS );
```

[Server Side] Step 3: Assign Address to Socket



Cpr E 489 -- D.Q.

21

Assign Address to Socket

bind()

```
% man -s 2 bind
cc [flag ...] file ... -lsocket -lnsl [library ...]
#include <sys/types.h>
#include <sys/socket.h>
int bind(int sockfd, const struct sockaddr *myaddr, socklen_t addrlen);
```

+ sockfd

- ➡ Socket descriptor returned by `socket()`

+ myaddr

- ➡ Pointer to a `sockaddr` structure

+ addrlen

- ➡ Length of the `sockaddr` structure

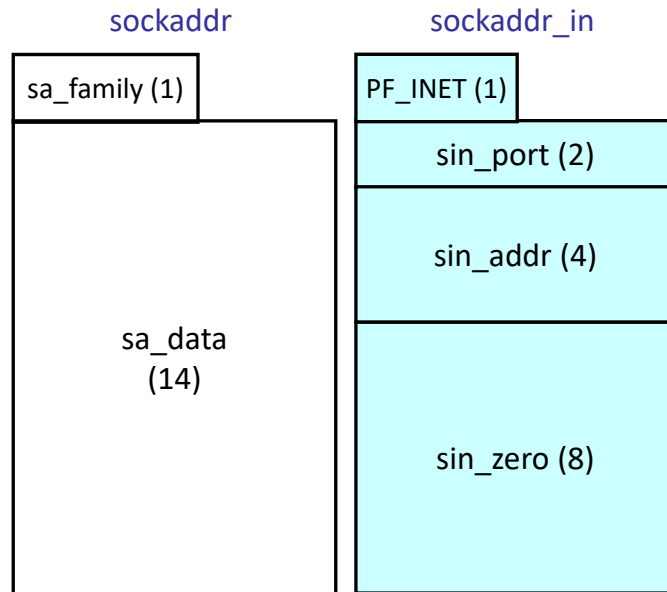
Cpr E 489 -- D.Q.

22

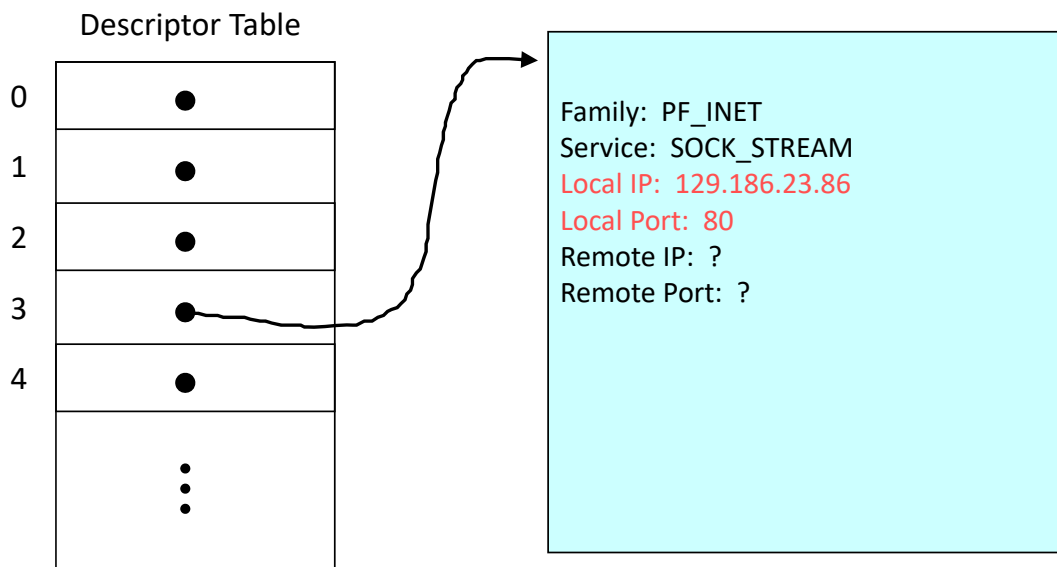
Generic Socket Address Structure

The C functions that make up the sockets API expect structures of generic `sockaddr` type

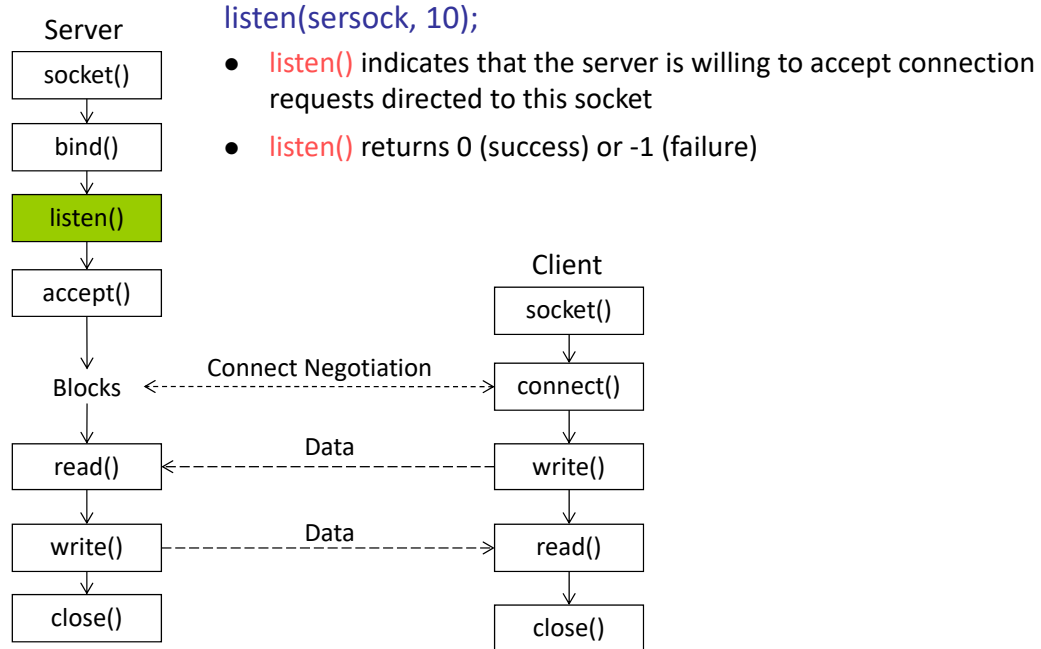
- so if omitting the cast `(struct sockaddr *)` in front of pointer to a `sockaddr_in` structure, you will get a warning when compiling



Socket Descriptor Data Structure



[Server Side] Step 4: Make a Passive-Mode Socket



Cpr E 489 -- D.Q.

25

listen()

```
% man -s 2 listen
cc [flag ...] file ... -lsocket -lnsl [library ...]
#include <sys/types.h>
#include <sys/socket.h>
int listen(int sockfd, int backlog);
```

+ sockfd

- ➡ Socket descriptor returned by `socket()` and already bound to server address via `bind()`

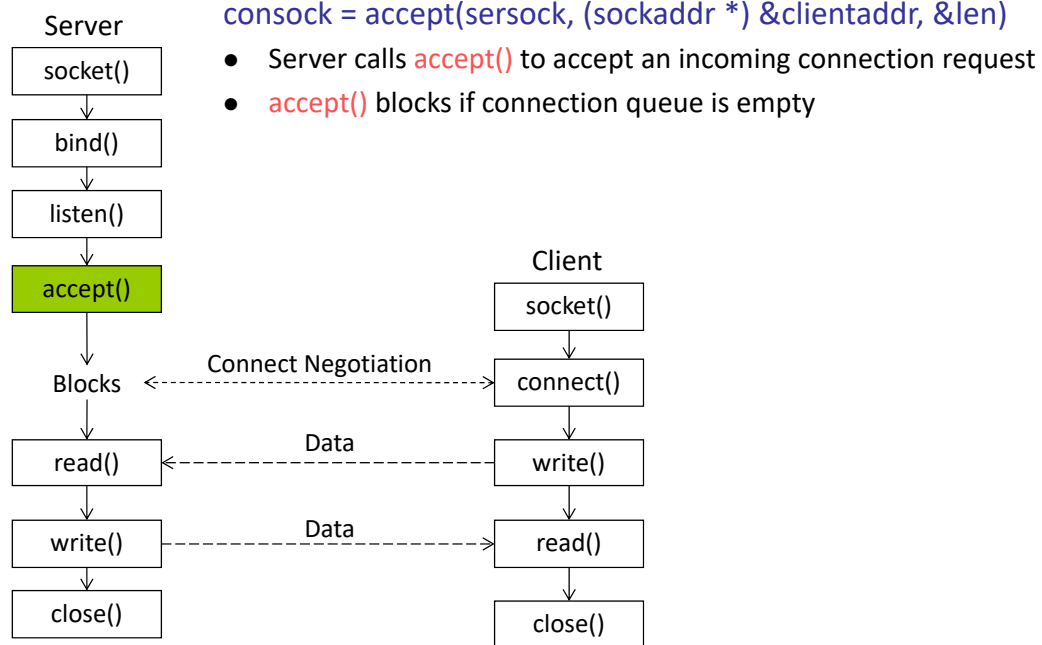
+ backlog

- ➡ Maximum number of connections that may be queued while waiting for the server to accept them

Cpr E 489 -- D.Q.

26

[Server Side] Step 5: Accept a Connection



Cpr E 489 -- D.Q.

27

`accept()`

```
% man -s 2 accept
cc [flag ...] file ... -lsocket -lnsl [library ...]
#include <sys/types.h>
#include <sys/socket.h>
int accept(int sockfd, struct sockaddr *cliaddr, int *addrlen);
```

✚ `sockfd`

- Socket descriptor of the passive-mode TCP socket

✚ `cliaddr`

- Pointer to `sockaddr` structure in which the client's address will be stored

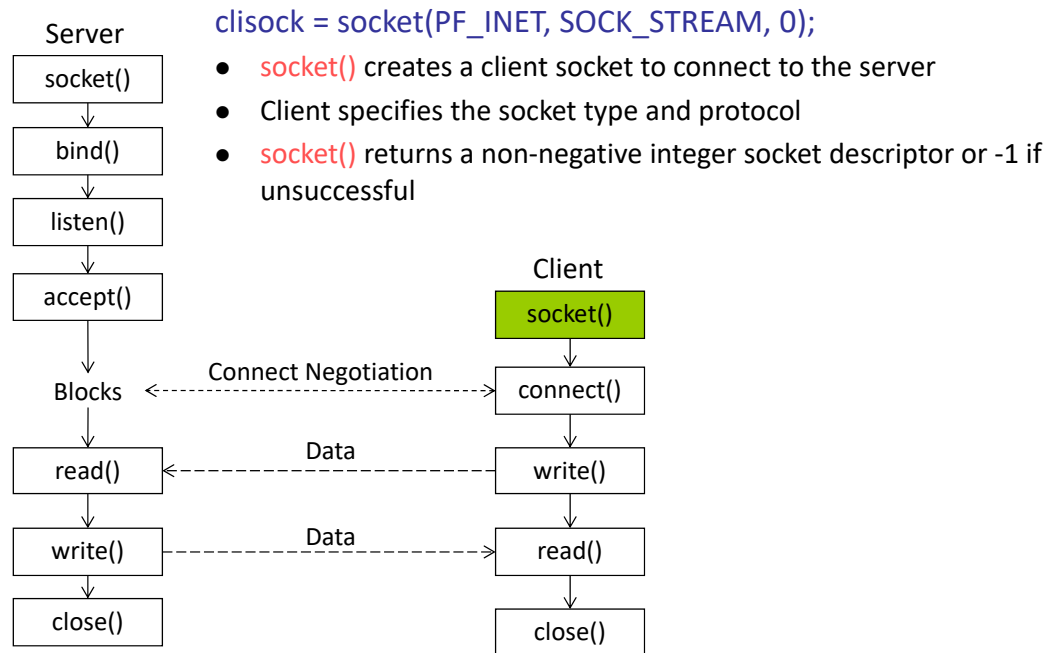
✚ `addrlen`

- Pointer to an integer in which the length of the `sockaddr` structure is returned

Cpr E 489 -- D.Q.

28

[Client Side] Step 1



Cpr E 489 -- D.Q.

29

[Client Side] Step 2, 3

✚ Skipped ...

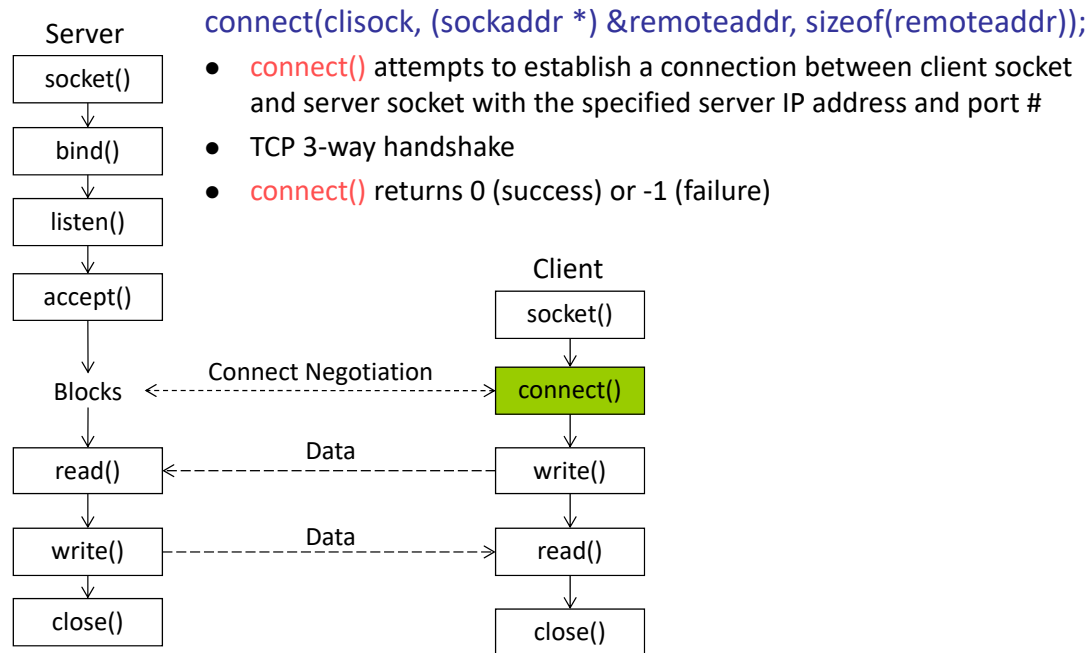
✚ When using a connection-oriented service, the client does not have to `bind` its socket:

- ➡ Kernel will take care of assigning the client address (Port #, IP address)

Cpr E 489 -- D.Q.

30

[Client Side] Step 4: Connect to Server



Cpr E 489 -- D.Q.

31

`connect()`

```
% man -s 2 connect
cc [flag ...] file ... -lsocket -lnsl [library ...]
#include <sys/types.h>
#include <sys/socket.h>
int connect(int sockfd, struct sockaddr *seraddr, socklen_t addrlen);
```

✚ `sockfd`

➡ Socket descriptor returned by `socket()` and already implicitly bound

✚ `seraddr`

➡ Pointer to `sockaddr` structure that contains the server's endpoint address

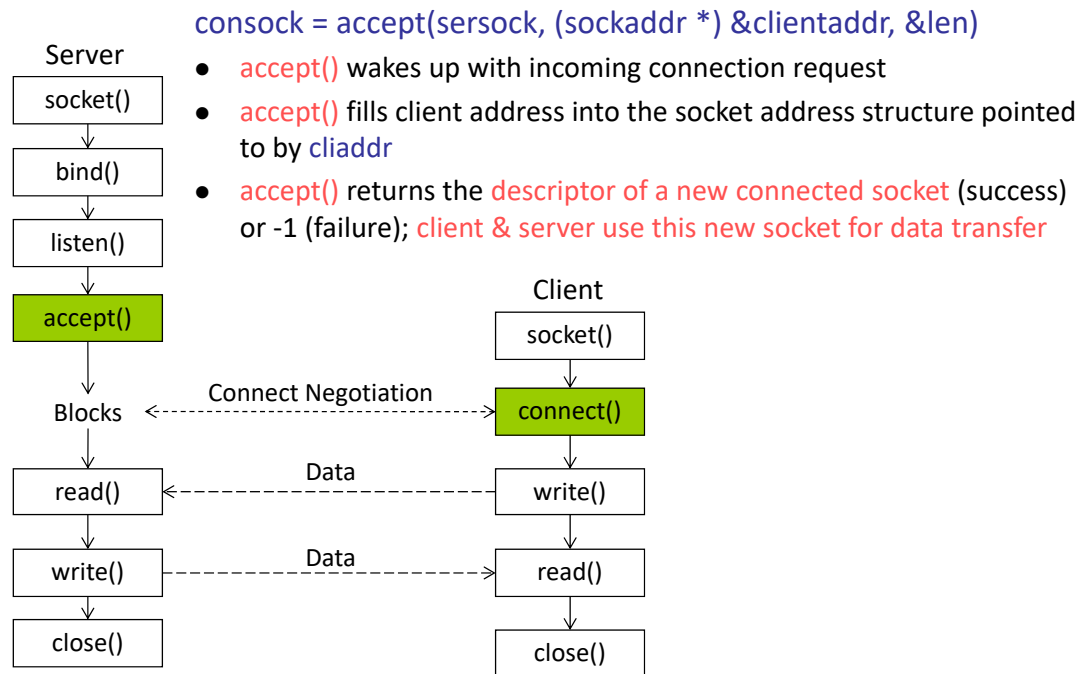
✚ `addrlen`

➡ Size of the `sockaddr` structure

Cpr E 489 -- D.Q.

32

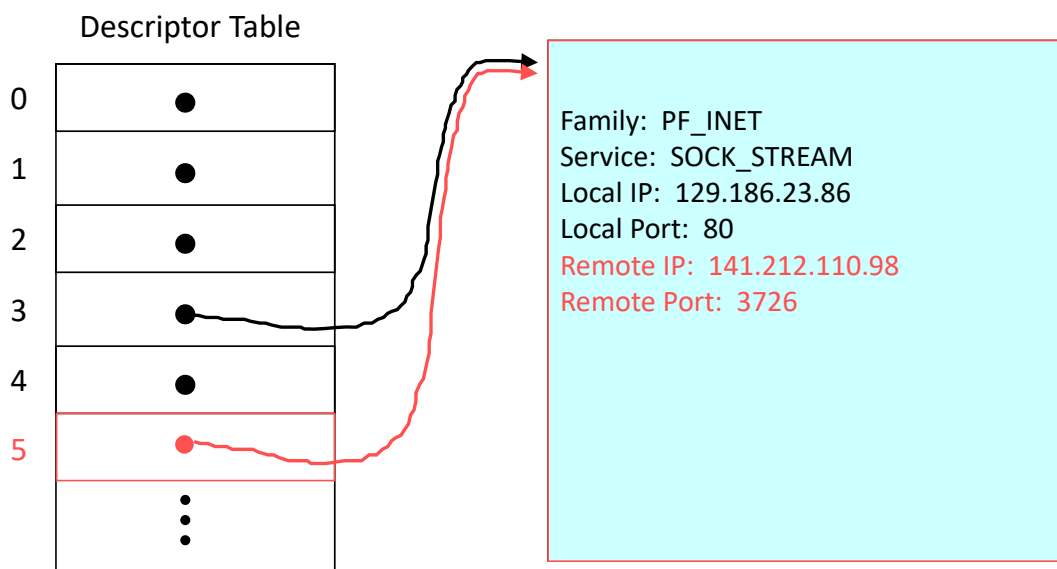
Back to [Server Side] Step 5



Cpr E 489 -- D.Q.

33

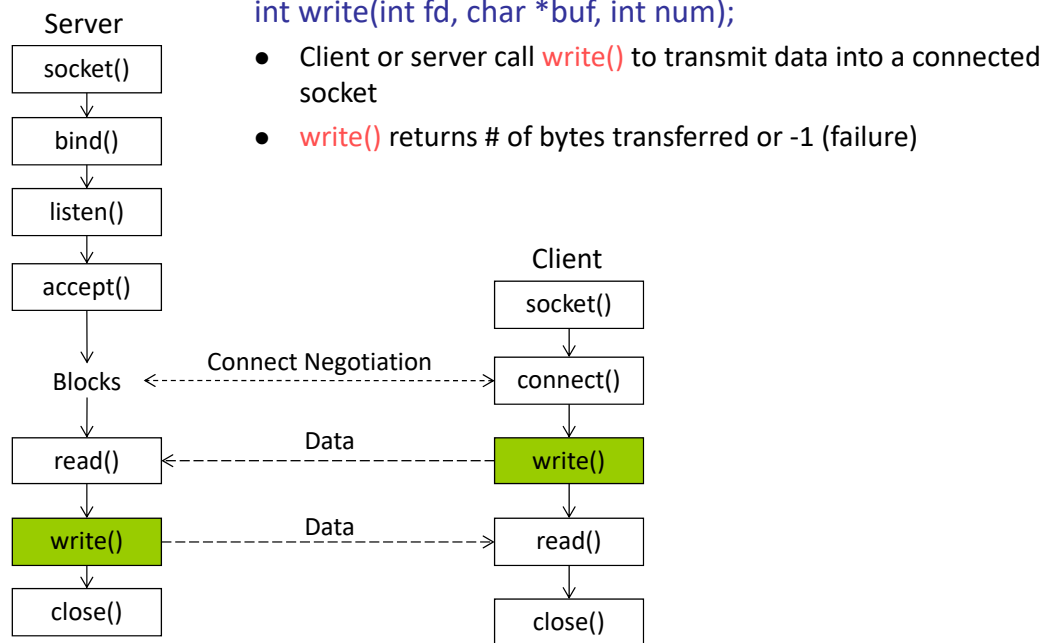
Socket Descriptor Data Structure



Cpr E 489 -- D.Q.

34

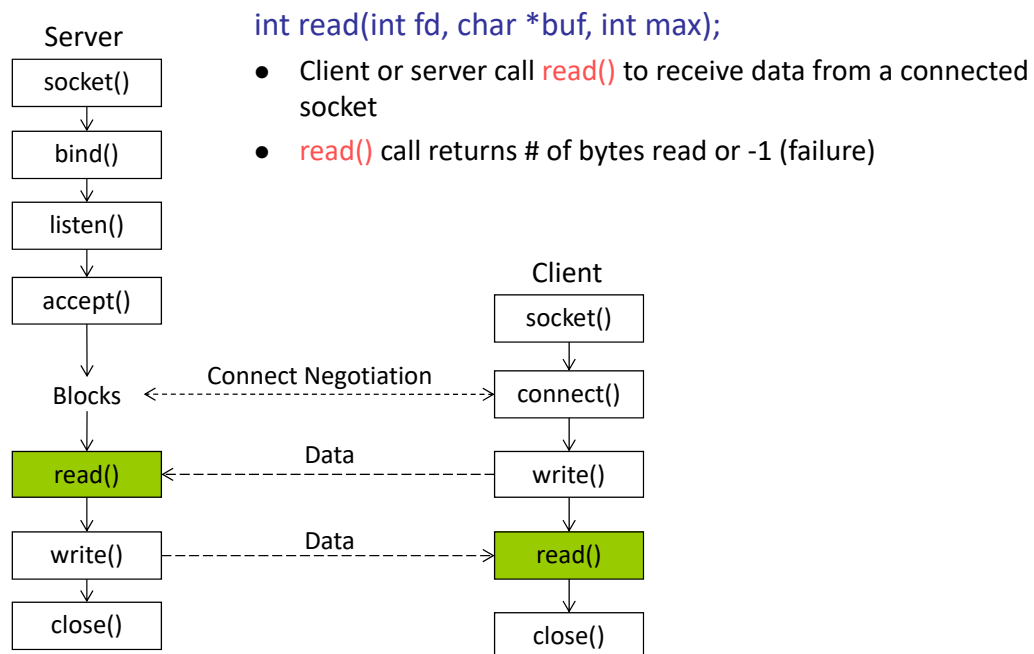
[Server-Client Interaction]



Cpr E 489 -- D.Q.

35

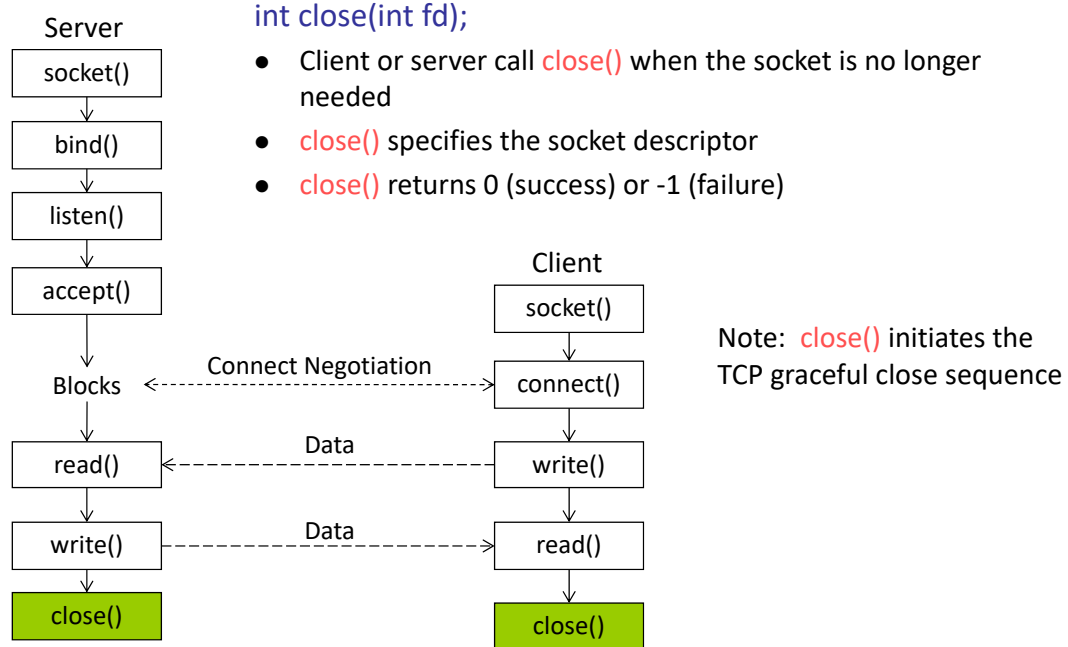
[Server-Client Interaction]



Cpr E 489 -- D.Q.

36

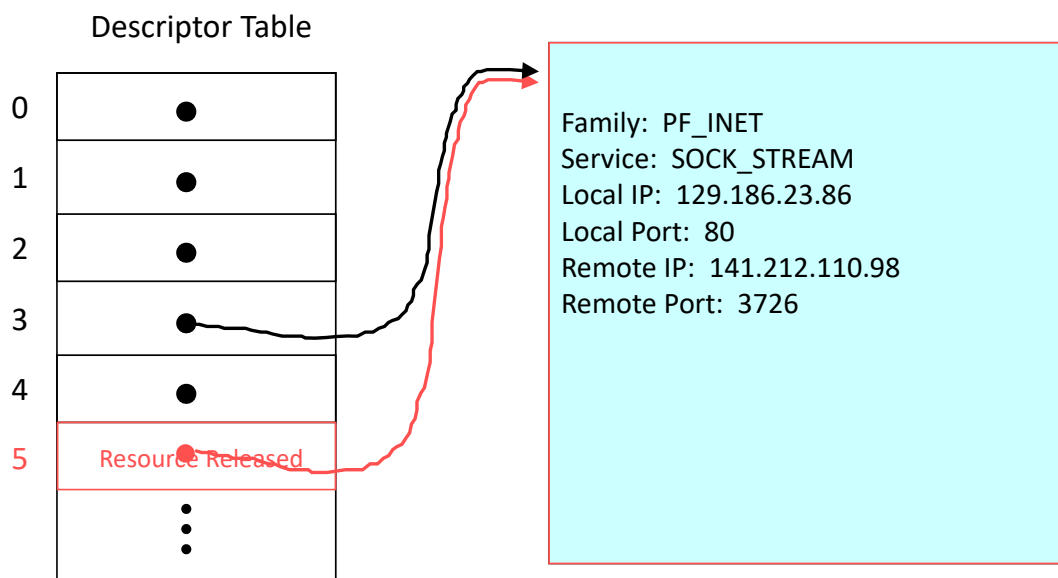
[Server or Client Side] Final Step: Connection Termination



Cpr E 489 -- D.Q.

37

Socket Descriptor Data Structure

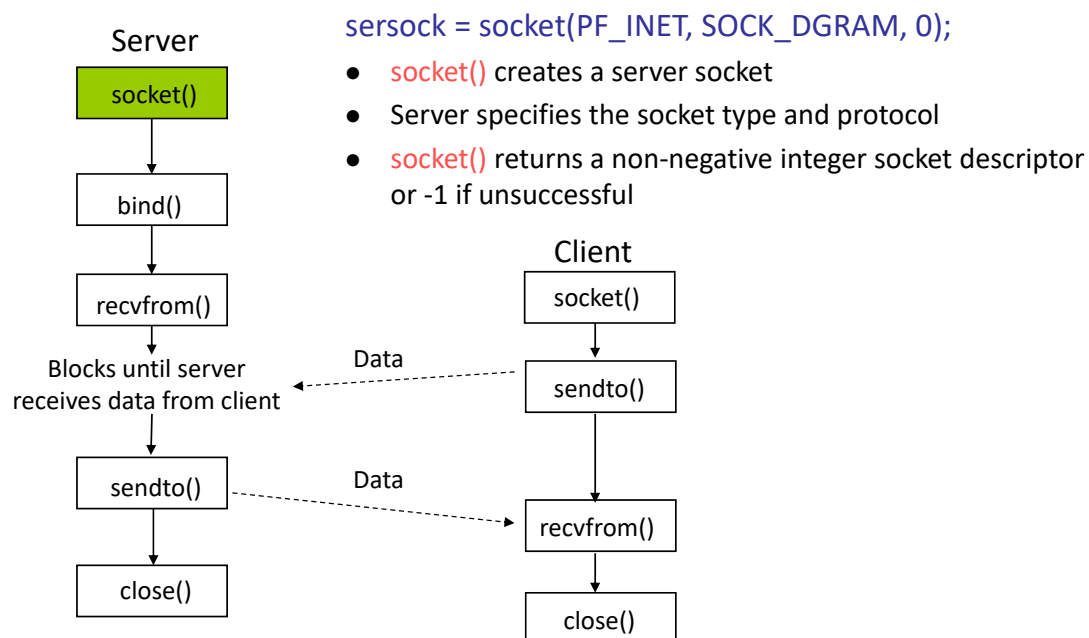


Cpr E 489 -- D.Q.

38

UDP Sockets Programming

[Server Side] Step 1



[Server Side] Step 2

- ✦ Specifying the server address:

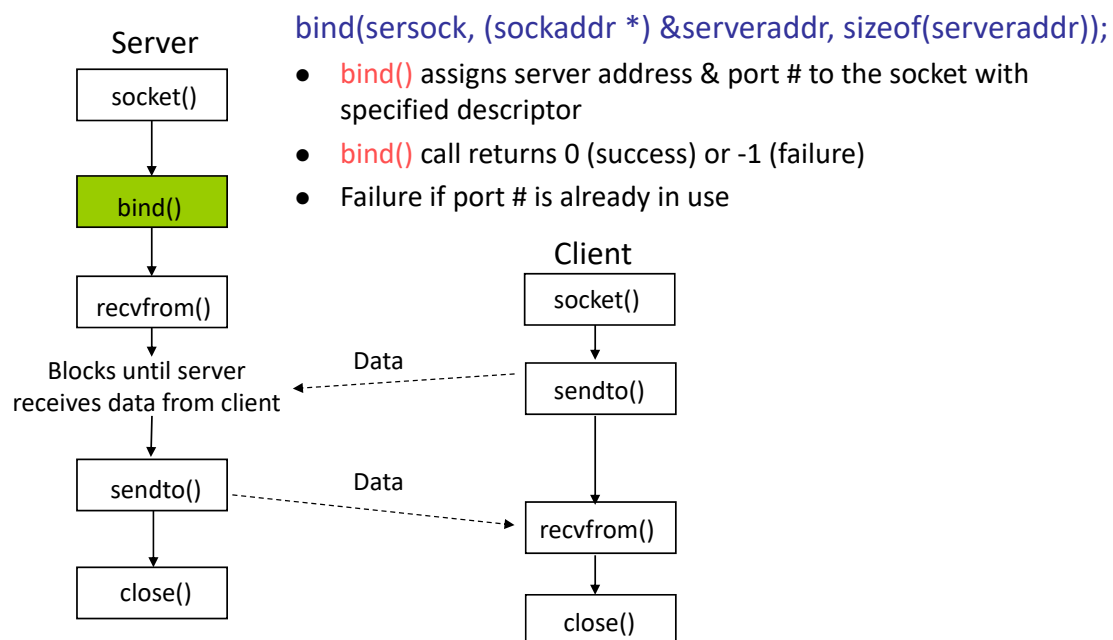
```
struct sockaddr_in serveraddr;
```

```
serveraddr.sin_family = PF_INET;
```

```
serveraddr.sin_port = htons( 37 );
```

```
serveraddr.sin_addr.s_addr = htonl( INADDR_ANY );
```

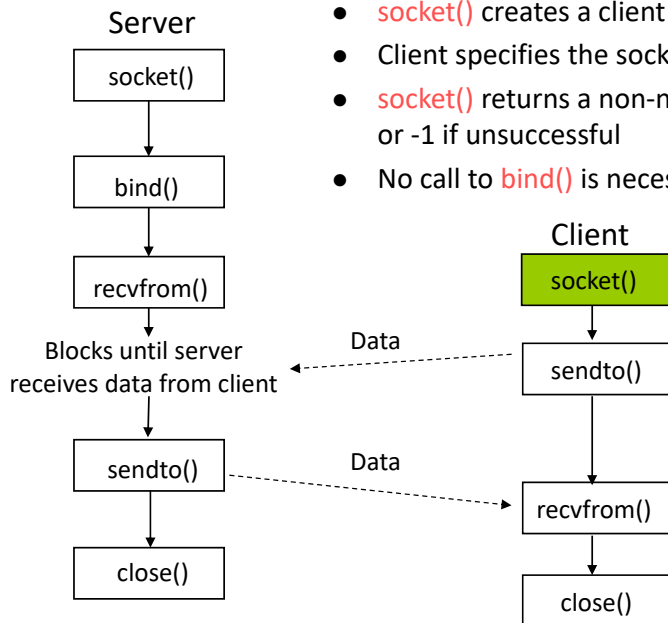
[Server Side] Step 3



[Client Side] Step 1, 2, 3

```
clisock = socket(PF_INET, SOCK_DGRAM, 0);
```

- `socket()` creates a client socket
- Client specifies the socket type and protocol
- `socket()` returns a non-negative integer socket descriptor or -1 if unsuccessful
- No call to `bind()` is necessary!



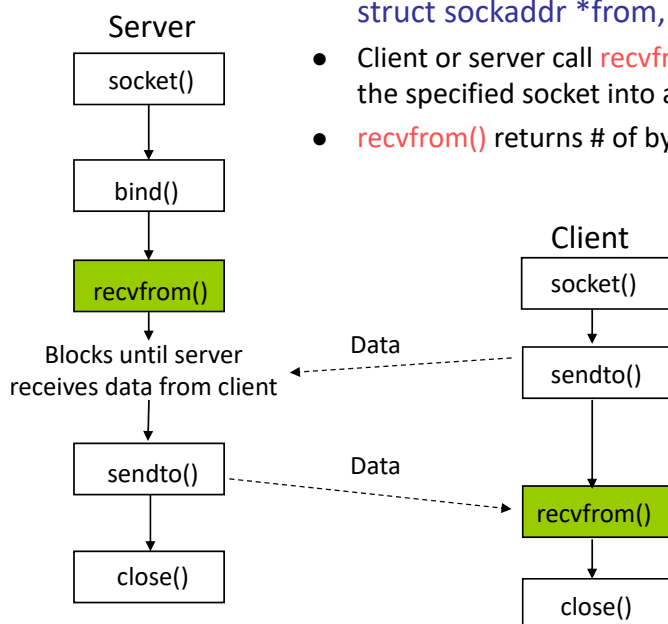
Cpr E 489 -- D.Q.

43

[Server-Client Interaction]

```
int recvfrom(int fd, char *buf, int max, int flags,  
struct sockaddr *from, int *fromlen);
```

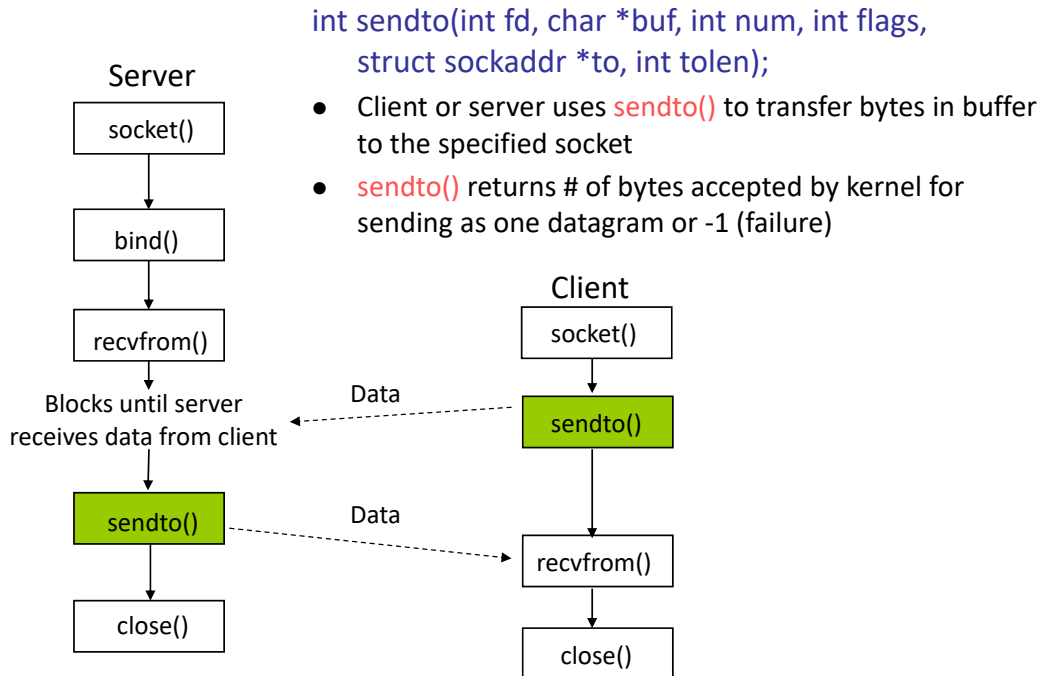
- Client or server call `recvfrom()` to copy bytes received in the specified socket into a specified location
- `recvfrom()` returns # of bytes received or -1 (failure)



Cpr E 489 -- D.Q.

44

[Server-Client Interaction]



Cpr E 489 -- D.Q.

45

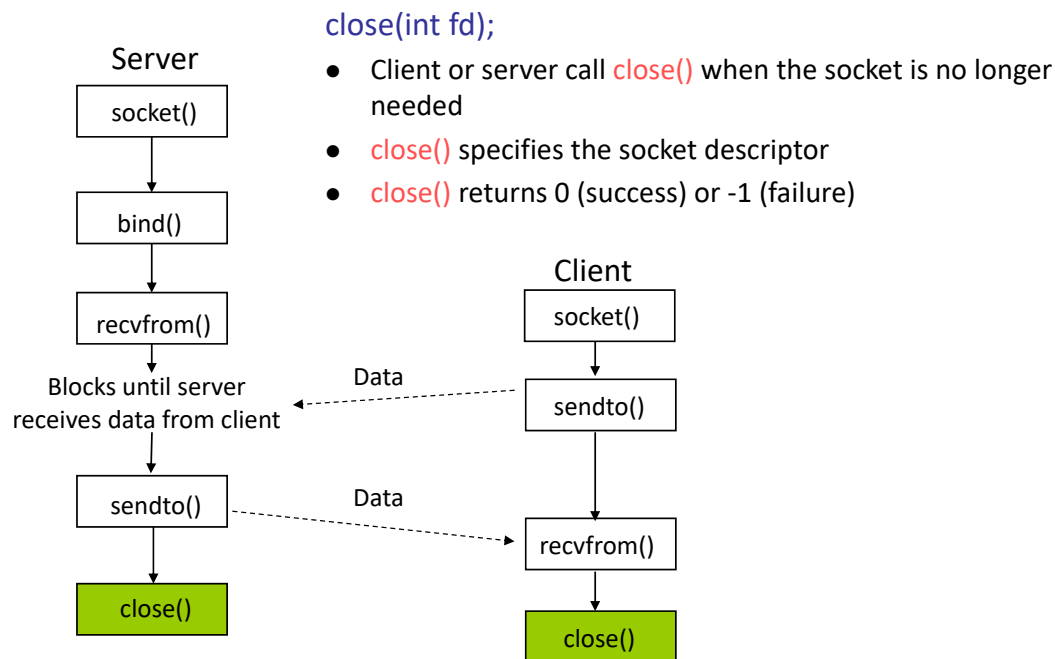
You can also `connect()` to a UDP socket!

- ✦ A UDP socket can be used in a call to `connect()`
- ✦ But:
 - This simply tells the kernel the address of the server
 - No handshake is performed
 - No data of any kind is sent on the network as a result of calling `connect()` on a UDP socket
- ✦ `connect()` is typically used with UDP when communication is with a single peer only
- ✦ Once a UDP socket is connected:
 - can use `write()` and `read()`

Cpr E 489 -- D.Q.

46

[Server or Client Side] Final Step: Close the Socket



Cpr E 489 -- D.Q.

47

Other Useful Functions

1. `int getsockopt(int s, int level, int optname, void *optval, int *optlen);`
2. `int setsockopt(int s, int level, int optname, const void *optval, int optlen);`

✚ Get and set socket options

✚ Some useful options

➡ `SO_REUSEADDR`:

- Allows server to restart
- Allows multiple servers to bind to the same port with different IP addresses
- All TCP servers should specify this option to allow the server to be restarted

Cpr E 489 -- D.Q.

48

⊕ Some useful options

➡ `SO_KEEPALIVE`

- Server sends a probe which client must respond to indicate that it is still alive

➡ `SO_LINGER`

- Indicates whether `close()` should return immediately or wait for connection termination

➡ `SO_RCVBUF` and `SO_SNDBUF`

- Set receive and send buffer sizes

3. `int gethostname(char *name, int namelen);`

- ⊕ Fills `name` with the standard host name of the current host

4. `int getsockname(int sockfd, struct sockaddr *localaddr, int *addrlen);`

- ⊕ Fills `localaddr` with the address of socket `sockfd`

5. `int getpeername(int sockfd, struct sockaddr *peeraddr, int *addrlen);`

- ⊕ Fills `peeraddr` with the address of the peer connected to socket `sockfd`

6. struct hostent *gethostbyname(const char *name);

✦ name: string that holds the host name

✦ struct hostent {
 char *h_name; // canonical name of host
 char **h_alias; // alias list
 int h_addrtype; // host address type
 int h_length; // length of address
 char **h_addr_list; // list of addresses
}

✦ Example

```
struct hostent *Address;  
Address = gethostbyname("www.ece.iastate.edu");
```

7. struct hostent *gethostbyaddr(const char *addr, int len, int type);

✦ addr: a pointer to in_addr structure containing the IP address

✦ Example

```
struct hostent *Address;  
unsigned long IP;  
  
IP = inet_addr("129.186.23.86");  
Address = gethostbyaddr((const char *) &IP, sizeof(IP), AF_INET);
```