

Back to Inference

Topics:

- 1.) Estimation of parameters
- 2.) confidence intervals
- 3.) Hypothesis Testing
- 4.) Prediction

Estimation

Start with $X_1, \dots, X_n \stackrel{iid}{\sim} f_X(x)$, where there is some parameter, say θ , associated with f_X . In statistics, θ is unknown to us, so we must estimate it from the data.

Def:

A statistic, $T(X_1, \dots, X_n)$, that is used to learn about an unknown parameter θ is called an Estimator.

Notes: 1) The term Estimator is used when referring to a statistic as a function of random variables

2) Typical notation is to put a "hat" over the parameter being estimated to denote an Estimator.

— $\hat{\theta}$ is an Estimator of θ

Def:

The observed value of a statistic used to learn about an unknown parameter is called an Estimate.

- Notes: 1) An Estimate is a function of the observed data values x_1, \dots, x_n (16)
- 2) It is an actual numeric value

EX

I have a sample $X_1, \dots, X_n \sim f_X(x)$ and want an estimator for $E(X) = \mu$. My Estimator will be $\hat{\mu} = \bar{X}$. I observe the values 6, 7, 7, 8, 9, 10. My estimate of μ is

$$\bar{x} = \frac{6+7+7+8+9+10}{6} = 7.83$$

Since an Estimator, $\hat{\theta}$, is a function of R.V.'s, it is a R.V. also. Thus it has its own distribution called the Sampling Distribution of $\hat{\theta}$.

- $E(\hat{\theta})$ is the mean of the sampling distribution
- The standard deviation of the sampling distribution is called the "Standard Error" and is denoted as $SE(\hat{\theta})$
- $SE(\hat{\theta}) = \sqrt{\text{Var}(\hat{\theta})}$

We make use of the sampling distribution in confidence intervals and hypothesis testing

A natural question now is:

IS my Estimator any good?

There are some properties we can look at.

- unbiasedness
- consistency
- mean square Error (MSE)
- Asymptotic Normality

in many cases, $\hat{\theta} \approx N(\theta, [SE(\hat{\theta})]^2)$

Def

An estimator, $\hat{\theta}$, is an unbiased estimator for θ if $E(\hat{\theta}) = \theta$

[on Average, we hit the target]

Def

An estimator, $\hat{\theta}$, is a consistent estimator for θ if $IP(|\hat{\theta} - \theta| > \epsilon) \rightarrow 0$ as $n \rightarrow \infty$

[as the sample size gets bigger, there is a high probability $\hat{\theta}$ will be close to θ]

Note: unbiasedness used to be a very desirable property. Still is important. Consistency is a good thing to have in an Estimator

Earlier in Notes, we said that we should use

\bar{X} as an Estimator for $E(X) = \mu$ and

S^2 as an Estimator for $Var(X) = \sigma^2$

Thm

\bar{X} and S^2 are both unbiased and consistent Estimators for their Respective Parameters

 \bar{X} :

$$E(\bar{X}) = \cancel{\frac{1}{n} \sum_{i=1}^n E(X_i)} \cdot E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) = \frac{1}{n} \sum_{i=1}^n E(X_i) = \frac{1}{n} n \mu = \mu$$

\Rightarrow unbiased

$$IP(|\bar{X} - \mu| > \epsilon) \leq \frac{Var(\bar{X})}{\epsilon^2} = \frac{\sigma^2}{n\epsilon^2} \rightarrow 0 \text{ as } n \rightarrow \infty$$

\Rightarrow Consistent (using Chebychev's Inequality on the R.V. \bar{X})

A popular metric to compare Estimators is the mean Square Error (MSE)

(18)

Def

The MSE of an Estimator $\hat{\theta}$ is $MSE(\hat{\theta}) = E((\hat{\theta} - \theta)^2)$

It can be shown that $MSE(\hat{\theta}) = [\text{Bias}(\hat{\theta})]^2 + \text{Var}(\hat{\theta})$

We would like an Estimator with a good mix of Small bias and Small Variance.

Ex

Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ and we want an Estimator for μ . Two choices we will compare:

- 1.) $\hat{\mu}_1 = X_1$
2.) $\hat{\mu}_2 = \bar{X}$ } compare MSEs

First, both Estimators have sampling distributions that are the Normal Distribution.

Second, $E(X_1) = \mu$ and $E(\bar{X}) = \mu$ so both are unbiased Estimators

Thus, to compare MSEs we just have to compare Variances.


$$\text{Var}(X_1) = \sigma^2 \quad \text{and} \quad \text{Var}(\bar{X}) = \sigma^2/n$$

Thus, ~~the~~ $MSE(\bar{X}) < MSE(X_1)$


preferred

and should be

sampling distribution of X_1



sampling distribution of \bar{X}



Statistical models

we want a model for our entire sample that we can use for Inference on parameters etc.

Def

A Statistical model is the joint distribution of our sample.

Recall: we saw joint distributions for two discrete R.V.'s earlier in the semester $\Rightarrow P_{XY}(x, y) = P(X=x, Y=y)$

Also, if X & Y were Independent, then the joint distribution could be written as $P_{XY}(x, y) = P(X=x, Y=y) = P_X(x)P_Y(y)$

Let X_1, \dots, X_n iid $f_X(x)$. The joint distribution of our sample is

$$f(x_1, x_2, \dots, x_n) = \prod_{i=1}^n f_X(x_i)$$

The joint distribution will involve unknown parameter(s) that we need to estimate to have a "working" model.

The first thing we can do as in the previous pages is use our Statistical model to come up with a single estimate (point Estimate) of the parameter(s) which we can then plug in so the model is usable

- In statistics this is called "fitting the model"

- In machine Learning this is called "Learning the model" with training data

EX I have a sample $X_1, \dots, X_n \stackrel{iid}{\sim} f_X(x)$

where $X_i = \# \text{ goals scored by ISU women's Soccer team in game } i$.

- 1.) come up with a model for the sample
- 2.) Estimate the parameter
- 3.) use fitted model to estimate probability the score more than 2 goals in next game.

Each X_i is a discrete R.V. that is the # of occurrences (goals) in some time frame (1 game) {maybe poisson Distribution?}

Our Assumption will be:

$X_1, \dots, X_n \stackrel{iid}{\sim} \text{Pois}(\lambda)$

(Recall: $P_X(x) = \frac{e^{-\lambda} \lambda^{x_i}}{x_i!}$ for $X_i \sim \text{Pois}(\lambda)$)

So our joint model for our sample will be:

$$f(x_1, \dots, x_n) = \prod_{i=1}^n f_X(x_i) =$$

distribution where the X_i 's are coming from

$$\frac{e^{-n\lambda} \lambda^{\sum_{i=1}^n x_i}}{\prod_{i=1}^n x_i!}$$

There is one parameter, λ , that needs to be estimated from the sample

Since $\lambda = E(X)$ I will use the plug in Estimator \bar{X} . Suppose the observed values were

0, 0, 1, 0, 1, 2, 2, 0, 1, 1

my estimate of λ is: $\hat{\lambda} = \bar{x} = \boxed{.8}$. Thus I assume that a poisson distribution with $\lambda = .8$ ~~dis~~ is the ~~gen~~ distribution that generated my data.

So, for the next game I can use my fitted model to answer

$$P(X > 2) = 1 - P(X \leq 2) \text{ where } X \sim \text{Pois}(.8)$$

$$= \boxed{.047}$$