

472 Recitation

Week 5

Sensorless Problem

Solution: a sequence of actions from search in the space B of **belief states** (b-states).

States: 2^N possible subsets of the set Σ of N physical states. $C_n^0 + C_n^1 + C_n^2 + \dots + C_n^n = 2^n$

Initial state: S .

Actions: At the belief state b , $\text{ACTIONS}(b) = \bigcup_{s \in b} \text{ACTIONS}_P(s)$
↑
Action on a physical state

Goal Test: possibly if one of the states $s \in b$ passes the test;
necessarily if every states $s \in b$ passes the test.

Cost: Could be one of several values if the same action as different costs in different states.

- Use of compact description
- Incremental belief-state search to build up solution one physical state at a time

Partially Observable Environments

Transition
Model:

- **Prediction**: computes the b-state from an action a .

estimate $\rightarrow \hat{b} = \text{PREDICT}(b, a)$

- **Possible percepts**: computes the set of percepts that could be observed in the predicted b-state.

$$\text{POSSIBLE-PERCEPTS}(\hat{b}) = \{o \mid o = \text{PERCEPTS}(s) \text{ and } s \in \hat{b}\}$$

- **Update**: computes the set of states in \hat{b} that could have produced the percept.

Algorithm: **AND-OR search algorithm**

Game


- s_0 : initial state – game setup.

At a state s :

- $\text{TO-MOVE}(s)$: the player to move in the state.
- $\text{ACTIONS}(s)$: the set of legal moves in the state.
- $\text{RESULT}(s, a)$: the transition model defining the next state from taking action a .
- $\text{IS-TERMINAL}(s)$: to test if the game is over, i.e., if s is a terminal state.
- $\text{UTILITY}(s, p)$: a **utility function** to return a value to the player p if the game ends in terminal state s .

Minimax

Optimal against an optimal opponent



Minimax

function MINIMAX-SEARCH(*game*, *state*) **returns** *an action*

$\text{player} \leftarrow \text{game}.\text{TO-MOVE}(\text{state})$

$\text{value}, \text{move} \leftarrow \text{MAX-VALUE}(\text{game}, \text{state})$

return *move*

function MAX-VALUE(*game*, *state*) **returns** *a (utility, move) pair*

if *game.IS-TERMINAL*(*state*) **then return** *game.UTILITY*(*state*, *player*), *null*

$v \leftarrow -\infty$

for each *a* **in** *game.ACTIONS*(*state*) **do**

$v2, a2 \leftarrow \text{MIN-VALUE}(\text{game}, \text{game.RESULT}(\text{state}, a))$

if $v2 > v$ **then**

$v, \text{move} \leftarrow v2, a$

return *v*, *move*

function MIN-VALUE(*game*, *state*) **returns** *a (utility, move) pair*

if *game.IS-TERMINAL*(*state*) **then return** *game.UTILITY*(*state*, *player*), *null*

$v \leftarrow +\infty$

for each *a* **in** *game.ACTIONS*(*state*) **do**

$v2, a2 \leftarrow \text{MAX-VALUE}(\text{game}, \text{game.RESULT}(\text{state}, a))$

if $v2 < v$ **then**

$v, \text{move} \leftarrow v2, a$

return *v*, *move*

Minimax

Minimax Algorithm

function MINIMAX-DECISION(*state*) *returns an action*
inputs: *state*, current state in game
return *argmax*_{*a* in ACTIONS(*state*)} MIN-VALUE(RESULT(*state*, *a*))

function MAX-VALUE(*state*) *returns a utility value*
if TERMINAL-TEST(*state*) **then** **return** UTILITY(*state*)
v ← −∞
for each *a* in ACTIONS(*state*) **do**
 v ← MAX(*v*, MIN-VALUE(RESULT(*state*, *a*)))
return *v*

function MIN-VALUE(*state*) *returns a utility value*
if TERMINAL-TEST(*state*) **then** **return** UTILITY(*state*)
v ← ∞
for each *a* in ACTIONS(*state*) **do**
 v ← MIN(*v*, MAX-VALUE(RESULT(*state*, *a*)))
return *v*



| Criteria | Performance |
|------------------|-------------------------|
| Completeness | Yes (if tree is finite) |
| Optimality | Yes |
| Time Complexity | $O(b^m)$ |
| Space Complexity | $O(bm)$ |

Alpha-beta pruning

- Improve the efficiency of Minimax search
- Alpha is the value of the choice for MAX
- Beta is the value of the choice for MIN

The α - β algorithm

function ALPHA-BETA-SEARCH(*state*) returns an action
inputs: *state*, current state in game
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$
return the action in ACTIONS(*state*) with value v

function MAX-VALUE(*state*, α , β) returns a utility value
if TERMINAL-TEST(*state*) then return UTILITY(*state*)
 $v \leftarrow -\infty$
for each *a* in ACTIONS(*state*) do
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{Result}(s,a), \alpha, \beta))$
 if $v \geq \beta$ then return v ← Beta pruning
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
return v

function MIN-VALUE(*state*, α , β) returns a utility value
if TERMINAL-TEST(*state*) then return UTILITY(*state*)
 $v \leftarrow +\infty$
for each *a* in ACTIONS(*state*) do
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{Result}(s,a), \alpha, \beta))$
 if $v \leq \alpha$ then return v ← Alpha pruning
 $\beta \leftarrow \text{MIN}(\beta, v)$
return v

