

CS 474/574 Machine Learning

1. Introduction

Prof. Dr. Forrest Sheng Bao
Dept. of Computer Science
Iowa State University
Ames, IA, USA

September 4, 2020

Why Machine Learning (ML)

- ▶ How computers know how to do things?
- ▶ Two ways:
 1. programming: steps detailed by human programmer
 2. learning: without being specifically told
- ▶ Example 1: machine translation
 1. programming: writing many rules to replace and reposition words, e.g., *Do you speak Julia? Sprechen Sie Julia?*
 2. learning: feeding the computer many bilingual documents
- ▶ Example 2: sorting
 1. programming: Quicksort, etc.
 2. learning: feeding the computer many pairs of unsorted and sorted list of numbers.
- ▶ The first approach in the context of AI is also called rule-based system or expert system, e.g. MyCin, Grammarly.

Why ML is attractive

- ▶ We are lazy. We want to shift the heavy lifting to the computers.
- ▶ We are incompetent. No kidding! Sometimes it is very difficult to come up with step-by-step instructions.
- ▶ Examples: Self-driving, AlphaGo, Automated circuit routing, Machine translation, Commonsense reasoning, text entailment, Document generation, auto-reply of messages/emails, fly a helicopter inversely, van-Gogh-lize paints.
- ▶ It is a dream. “Creating an artificial being has been the dream since the beginning of science.” – Movie A.I., Spielberg et al., 2001

Three types of MLs

ML (in current approaches) is about finding/approximating functions.

- ▶ Supervised, finding $\hat{f}(x) \approx f(x)$ with ground truth provided by human.
 - ▶ Let x and y be two (vectors of) variables, and a function connecting them $y = f(x)$ But only god knows f .
 - ▶ We construct another function \hat{f} to approximate f such that $\hat{y} = \hat{f}(x) \approx y = f(x)$ for a(ny) given x .
 - ▶ **Supervised** because we provide many pairs of x 's and y 's for the computer to know the difference between \hat{y} and y on a large pool of samples.
 - ▶ Examples: object detection from images, Flavia, CPU branch prediction, COVID-19 diagnosis from blood profile, Epileptic EEG recognition, depression treatment from brain shapes.
 - ▶ Beyond categorization/classification: Mflux, Review helpfulness prediction, Document summarization, predict house prices
- ▶ Unsupervised, finding $\hat{f}(x)$ without ground truth
- ▶ Reinforcement, let the machine find ground truth itself

Representation of x

- ▶ x is usually not a simple (vector of) number(s). How to tell it to a computer?
- ▶ Example: bananas vs. apples
- ▶ **Feature engineering**: manually craft functions to **extract** features from raw data, e.g., SIFT, bag-of-words.
- ▶ Automated feature extraction in deep learning: E.g., filters in CNNs.
- ▶ If x involves categorical values (e.g., gender), there are usually two approaches: **One-hot encoding** and **embedding** (in DL context, to be discussed later).

Supervised ML

- ▶ Given many pairs of inputs and outputs:
 $\{(\mathbf{X}_1, \mathbf{y}_1), (\mathbf{X}_2, \mathbf{y}_2), \dots, (\mathbf{X}_N, \mathbf{y}_N)\},$
- ▶ that underline a “black-box” function $f : \mathbb{R}^n \mapsto \mathbb{R}^m$ such that $\forall i \in [1..n], f(\mathbf{X}_i) = \mathbf{y}_i,$
- ▶ construct a function \hat{f} that approximates the function f .
- ▶ “approximate”: usually $\min ||\hat{f}(x) - f(x)||^p$ where p is usually 1 or 2. See ℓ_p -norm .
- ▶ The process of finding the approximation function \hat{f} is called **training** or **learning**.
- ▶ \hat{f} is called a **model** or an **estimator**.
- ▶ \mathbf{X}_i : an **input** (especially when raw data is used as the input) or **feature vector** (if using feature engineering).
- ▶ \mathbf{y}_i , often $\in \mathbb{R}^1$ a **label** (in classification) or **target** (used more generally and lately).
- ▶ Classification vs. Regression: When y is continuous or discrete. In modern DL context, such division is usually no mentioned, especially in generative tasks.